

Bayesian Adaptive Estimation and Inference in General Recognition
Theory Models With Psychometric Functions

Joni Pääkkö

Maisterintutkielma
Jyväskylän Yliopisto
MUTKU
Syksy 2020

Abstract

General Recognition Theory (GRT) is a multidimensional generalization of Signal Detection Theory. It is used to model the detection of signals with multiple dimensions, e.g. auditory signals which can vary not only in their pitch but also in their timbre independently. Main focus is in if the detection of these dimensions is coupled.

Bayesian adaptive estimation has been successfully applied to many different tasks and models in psychophysics. The main goal of this thesis is to study the application of it to GRT models. To achieve this, I will introduce a GRT models for a Yes/No and 2I-4AFC procedure that include psychometric functions to model the relationship between physical signal strengths and d' explicitly. Also, methods for Bayesian estimation of these models are introduced.

The performance of the models are evaluated in simulations ($N = 772$) and in a small scale psychophysical experiment ($N = 2$).

Simulations indicate that the adaptive algorithm is more efficient on average, but offers little practical improvement over random sampling in this context.

Data from the psychophysical experiments indicate that non-sensory factors play a big part, indicating that more work should be put into developing models/procedures that can identify between these factors. The data also suggests that the often default choice of coupling means does not necessarily capture all relevant properties of it.

Contents

1	Introduction	1
2	Theories of Detection	5
2.1	Signal Detection Theory	5
2.1.1	Relationship between stimulus and d'	6
2.1.2	Modelling responses: relationship between d' and response	8
2.2	General Recognition Theory	12
2.2.1	Calculating d' in two dimensions	13
2.2.2	Modelling responses in two dimensions	14
2.2.3	Interactions between the dimensions	18
2.3	Lapsing behaviour	20
2.4	Critical discussion	21
2.4.1	Relationship to other models	21
2.4.2	Nonidentifiabilities and other limitations of the model	24
3	Bayesian statistics	27
3.1	Basics of Bayesian statistics	27
3.2	Hierarchical models	30
3.3	Approximating the Posterior Probability Densities	32
3.3.1	Resample-Move algorithm	34
3.3.2	Laplace Approximation	36
3.4	Bayesian inference in the General Recognition Theory	37
4	Adaptive estimation	40
5	Simulations	43
5.1	Methods	44
5.2	Results	46
5.3	Discussion	52
6	Psychophysical experiments	54
6.1	Methods	54
6.2	Results	55
6.2.1	Posterior predictive plots	56
6.2.2	Participant OK, Models 1 to 3	58
6.2.3	Participant OK, Model with both interactions	63
6.2.4	Participant JP, Models 1 to 3	65
6.2.5	Participant JP, Model with both interactions	69
6.2.6	Participant JP, Model with both interactions, $\kappa_\sigma > 0$	71
6.3	Some introspective remarks	72
6.4	General discussion	73
7	Conclusions	75

Appendices

A Algorithm for calculating ϕ_2	82
B Code for models for the Yes/No task	84
B.1 Model 1	84
B.2 Model 2	85
B.3 Model 3	86
B.4 Both	87
B.5 Both (Truncated)	88
C Code for models for the 2I-4AFC task	91
C.1 Model 1	91
C.2 Model 2	92
C.3 Model 3	93
C.4 Both	94

1 Introduction

Multidimensionality of perception Perception can be thought of as being multidimensional. For example, a single tone played with a guitar could be said to consist of pitch, length and timbre, which are some of the dimensions of an auditory signal, in this case, a musical tone. Another way in which perception is multidimensional is when it is *multimodal*: in many situations we combine information from different sensory modalities, for example haptic and auditory information when pressing a key on the keyboard, or visual and auditory information when watching someone's lips as they speak.

How we perceive such multidimensional signals is a fundamental psychological question: during perception, are the dimensions processed independently, or do they interfere with each other? Is the perception of auditory signals influenced by visual or other signals, and if so, in what situations? Do the physical dimensions, such as amplitude, correspond to the psychological (in the case of amplitude loudness) or do we perceive them as combinations (do e.g. higher pitches sound louder)?

In the context of the empirical part of this work (see Section 6 *Psychophysical experiments*) the coupling of pitch and timbre is the most relevant. These present a special case in which the features belong to the same modality: both are features of audition.

What does coupling or interaction mean? This is an important question, especially since pitch and timbre belong to the same modality.

Concretely, for example in an experiment in which the participants have to detect changes in pitch and timbre, there could be a loud noise that occurs on random trials. This loud noise, of course, affects the detection of the auditory stimulus as a whole: detection thresholds for both pitch and timbre would be affected by the noise. This would make it seem like those dimensions are connected, and of course they are, since they are both features of an auditory stimulus, but usually the focus is on perceptual or cognitive interactions.

As a visual example consider the optical illusion presented in Figure 1. The sides of the box should appear ever-so-slightly concave. The circles affect how the straight lines that constitute the box are *experience* on a perceptual/cognitive level.

It is this kind of interaction that is of interest; an interaction which affects the *appearance* of the stimuli, or their constituent dimensions.

One possible approach for modelling multidimensional perception is the General Recognition Theory (GRT), which is a multidimensional generalization of Signal Detection Theory (SDT) (Ashby & Soto, 2015; Ashby & Townsend, 1986; Wickens, 2002). SDT is a much-studied mathematical framework on one hand describing the detection of signals in the presence of noise and on the other hand for dissociating perceptual and decisional processes. In Signal Detection Theory—as will be discussed in more detail later—participant's tendency to report that they detected a signal is modelled separately from their sensory processes.

Adaptive estimation and Bayesian statistics Most experiments employing GRT use pre-selected stimuli, i.e. use a paradigm called *method of constant stimuli* (MOCS). However, it is known that using methods that adapt the stimulus values to the observed behaviour of the subject can make the testing more efficient (Kontsevich & Tyler, 1999). These kinds of methods are usually simply referred to as *adaptive* methods.

There are many adaptive methods that employ either *ad hoc* or common sense principles such

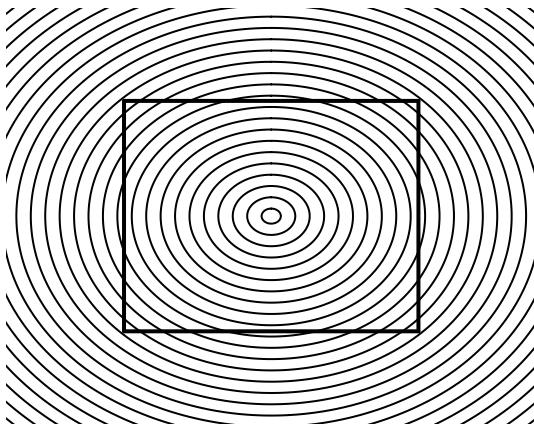


Figure 1: The sides of the box should appear slightly concave. Note that there is individual variability in whether this illusion works or not.

as "*If the subject correctly identifies the signal twice in a row, reduce its level, otherwise increase its level*", which could be a useful heuristic for finding e.g. the hearing threshold. In Bayesian Adaptive estimation the heuristic for choosing the stimuli is based on, as the name implies, Bayesian statistics.

In Bayesian statistics uncertainty about the parameters of a statistical model is represented as a joint probability distribution over them. For example when modelling hearing threshold in broadband noise, one could use a normal distribution for quantifying their prior uncertainty by saying that they expect the hearing threshold to be around 40 decibels with a standard deviation of 20 decibels.

When data is collected, it informs the experimenter through the statistical model which parameter values are more likely; or conversely what parameter values are *less* likely: if the subject is not able to detect 40 decibel signals correctly, this makes it less likely that values below are their hearing threshold, those values can be ruled out. This "ruling out" of parts of the parameter space can be said to reduce our uncertainty about plausible values for hearing threshold.

It should also be clear that other data points reduce information more than others: if the subject is always presented with signals over 60 decibels and they always detect them, this does not inform the experimenter if 30 decibels is more likely than 40 decibels as the hearing threshold. In Bayesian adaptive estimation one simply tries to pick the test signals in such a way that the uncertainty about the parameters is minimized on each trial, or in other words, as many possibilities as possible can be ruled out.

Structure of this thesis The main focus of this thesis is in applying Bayesian adaptive estimation to General Recognition Theory (GRT). However, this is a multifaceted problem, and it's useful to break it down into a few separate topics—which on the other hand are interrelated.

First, the "classic" GRT models could be described as non-parametric in the sense that they don't parameterize the relationship between the signal strength and the probability of a response; response probabilities are estimated only for stimulus categories. However, in order to meaningfully implement the adaptive method, one should be able to calculate response probabilities for arbitrary stimuli. In order to achieve just this, I will be integrating psychophysical functions—functions that relate continuous stimulus levels to response probabilities (Kingdom & Prins, 2010, Chapter 4)—into the theory of GRT. This constitutes the main part of Section 2 *Theories of Detection*.

Second, when such functional relationships are implemented, the question of *how* or what kind of functional relationships between stimuli and responses should be modelled becomes an issue. There are many different kinds of functions one can choose from, and many kinds of interactions that one could incorporate into the model. This issue will be touched upon in Section 2 when different kinds of interactions are discussed. Practical ramifications of this are discussed when analysing the data from the psychophysical experiments.

Related to this, I will be criticizing the implementation of the widely used 2X2 identification experiment in the context of auditory psychophysics and will instead argue for modelling discrimination. The core of the criticism is that in a psychophysical experiment, in which the stimulus levels are extremely low, it is likely that participants base their decisions on perceived differences between the stimuli rather than their absolute levels (Section 2.4.1 *Relationship to other models*).

Third, Bayesian statistical inference is different from the usual frequentist methods used when analysing GRT models. This is most apparent in how the concept of *interaction* is not dependent on statistical significance, but rather the posterior probability densities for the coefficients that model interactions. This is, of course, closely related to the second point, since the choice of the model determines how those coefficients are to be interpreted. (Section 3.4 *Bayesian inference in the General Recognition Theory*)

Fourth, Bayesian adaptive estimation requires, as its name already implies, the calculation of prior and posterior probability distributions. These distributions lack analytical solutions, so approximations based on Monte Carlo methods (random sampling) are used. This issue is discussed in Section 3.3 *Approximating the Posterior Probability Densities*. This is mainly a computational problem: since the aim of adaptive estimation is to select optimal stimuli in "real time", the algorithms used have to be relatively lightweight.

Related to all of the above angles are multiple practical problems, such as how to maximize the function that describes the optimality of stimuli, how many random samples one should use when approximating the posterior probability densities, what kind of parameterization of the likelihood function (the statistical model) to use, what is the most comfortable way of inputting responses for participants and so on. These questions will be addressed as they rise.

The reader should be aware that what is being evaluated here is the intersection of all the factors just mentioned, and changing any of these would change the results. This limits the generalizability of the results, which is a common problem with simulation studies. The main contribution of this thesis, then, is to explore the application of Bayesian (adaptive) methods to GRT and map some of the problems associated with this.

This thesis is divided into three main parts: The first part describes GRT, beginning from Signal Detection Theory and gradually exposing the reader to the models used in this thesis (Section 2 *Theories of Detection*); Second part describes Bayesian statistics and the adaptive algorithm (Sections 3 *Bayesian statistics* and 4 *Adaptive estimation*); the final part consists of simulations and psychophysical experiments which are used to evaluate the models (Sections 5 *Simulations* and 6 *Psychophysical experiments*).

2 Theories of Detection

The central theoretical concept in this thesis is General Recognition Theory (GRT), which is essentially a generalization of Signal Detection Theory (SDT) to multiple dimensions, with a focus on modelling interactions between them. Due to this hierarchy, I will first introduce the reader to SDT, before discussing GRT.

Both SDT and GRT are discussed in the context of a *discrimination task*. In such a task a single stimulus consists of two components: the *reference tone* and the *test tone*, as shown schematically in Figure 2. If for example the dimension of interest is pitch, the reference tone would always be the same, e.g. 150 Hz, and the test tone would either be the same or a higher pitch. The participant would then have to determine if there was a difference between the test and reference tones.

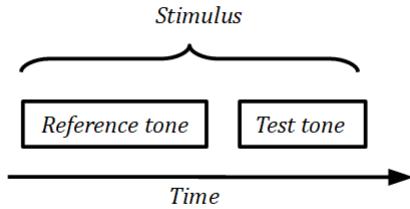


Figure 2: Schematic illustration of a stimulus in a discrimination task.

A stimulus in which the difference between reference and test tones is zero is called a *noise* stimulus while one in which the difference is greater than zero is called a *signal* stimulus; alternatively these can be called just noise and signal.

2.1 Signal Detection Theory

Continuing the example of a pitch discrimination task, during such experiment it is not rare to observe the participant sometimes give a negative and sometimes a positive response to the same stimulus on different presentations, that is, sometimes the participant will report hearing a difference and sometimes not even though the difference between the reference and test tone is the same.

The view that SDT takes is that the responses are not based directly on the physical signals, but rather some latent (unobserved) internal quantity, sometimes referred to as *evidence* (Wickens, 2002) or *judgments* Stigler (2003), I will be using the term *evidence* from here on. It is thought that the latent amount of evidence for e.g. the test tone being higher than the reference tone is subject to random variation, which in turn leads to variation in responses (Kingdom and Prins 2010, p. 154; Wickens 2002, p. 11). These random perturbations to the amount of evidence are commonly thought to arise primarily from sensory sources, but I will discuss the issue of how to interpret these latent quantities later in greater detail.

Some assumptions about the nature of the perturbations have to be made, for example that smaller ones are more common than large, and that the average amount of error is zero. A popular choice for modelling the distribution of the perturbations is the normal distribution with mean 0 and unknown variance. By making this assumption about the distribution of the perturbations, one can

calculate the predicted response probabilities for different stimuli, and consequently make statistical inferences about the sensory processes of interest.

2.1.1 Relationship between stimulus and d'

Since the amount of evidence is assumed to be random, it can be represented as probability distributions as in Figure 3. In the figure, S stands for physical signal level, for example the difference in frequencies between the reference and test tones. As already stated, normal distribution is used to model the distribution of evidence, however, in the context of lapses (accidental responses) I will discuss mixtures of distributions.

If the subject is presented with a noise stimulus (in which $S = 0$), the amount of evidence on any single trial is assumed to be a random sample from the zero-centred distribution in Figure 3; in the context of the pitch discrimination task this would correspond with stimulus in which the test and reference pitches are identical. As S is increased, the distribution from which the evidence is assumed to be sampled from shifts rightward, farther from zero, as is shown for $S = 2$ and $S = 4$.

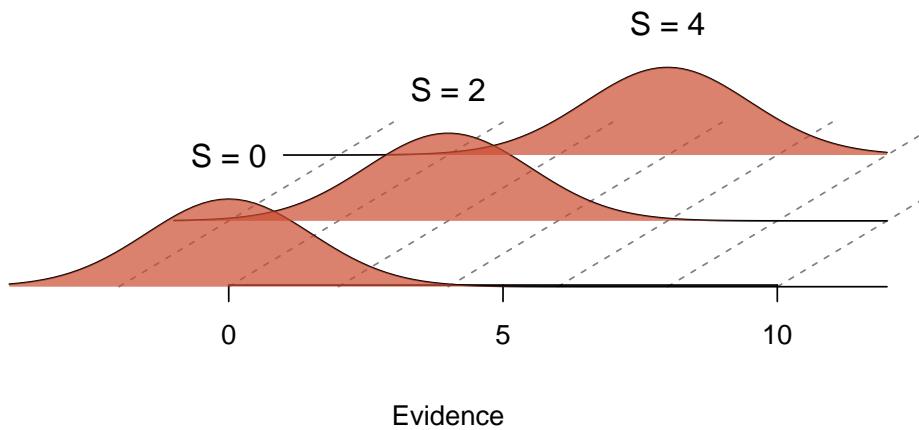


Figure 3: Evidence, according to SDT, is a random variable. It can be assumed to be normally distributed, with monotonically increasing μ as a function of signal strength, S , and constant variance, σ .

Signal level divided by the standard deviation of evidence is called d' , usually interpreted as *discriminability* (Kingdom & Prins, 2010, Chapter 6) or *signal-to-noise ratio* (Kontsevich & Tyler, 1999). Standardized in this way, the standard deviation of evidence distribution is always one, and the mean is d' , which is most conveniently interpreted indeed as signal-to-noise ratio.

The relationship between d' and S is not necessarily linear. Theoretically this kind of non-linearity could arise e.g. from changes in the variance of the evidence distribution, however, the

non-linearity is often interpreted as being the result of non-linearity in signal transduction, i.e. the change of physical signal into internal.¹

The functional relationship between physical signal strength and d' has been quantified in many different ways in the psychophysical literature (for a selection of functions, see Lesmes et al. (2015, Appendix A)). The functional form used here is based on the widely used *power law* model of signal transduction (see e.g. Dai and Micheyl (2011); Kontsevich and Tyler (1999); Lesmes et al. (2015)):

$$d' = \left(\frac{S}{\sigma}\right)^\beta \quad (1)$$

Here S is again the physical signal strength, σ corresponds to the standard deviation of the evidence distribution and β models non-linearity between signal level and signal-to-noise ratio. Note that often the letter α is used for the standard deviation of evidence (e.g. in Dai and Micheyl (2011); Kingdom and Prins (2010); Kontsevich and Tyler (1999)), but since this parameter is directly related to the standard deviation of the evidence distribution, I believe σ is more appropriate.

The effects of changing these parameters are shown in Figure 4. The smaller σ is, the faster d' increases. β parameter affects the linearity of the function: when $\beta < 1$ the function increases faster than the linear case before the point at which $d' = 1$ and slower after that; the opposite happens when $\beta > 1$.

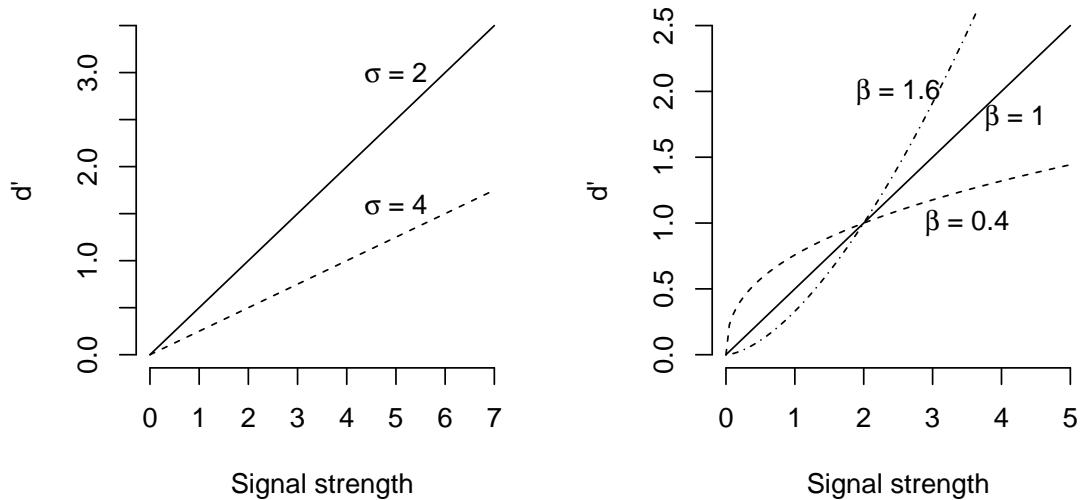


Figure 4: The effect of σ and β parameters on the d' function. Left panel shows the effect of changing σ parameter while keeping β constant ($\beta = 1$) and the right panel shows the effect of changing β parameter while keeping σ constant ($\sigma = 2$).

¹Familiar example of non-linear transduction is how pitch is experienced in large scale: in order to move one octave up on the psychological scale of pitch, one has to move increasingly fast on the physical frequency, see Fastl and Zwicker (2007, Chapter 5).

2.1.2 Modelling responses: relationship between d' and response

The preceding discussion about relating S -values to d' -values is just one half of SDT, there has to be a way of relating the d' -values to responses, R . Usually categorical decisions are required from the participant, and the specific response categories depend on what kind of task they are performing. I will discuss a Yes/No and a 2-interval 2-alternative forced choice (2I-2AFC) tasks. Structure of these tasks is shown schematically in Figure 5.

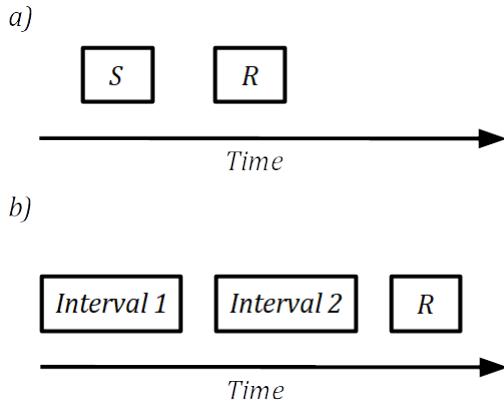


Figure 5: Schematic representations of the Yes/No (panel a) and 2-interval forced choice (panel b) tasks. S = signal, R = response.

In the Yes/No task (panel a of Figure 5) the participant's task is to indicate if they detected the signal. The name of the task is quite literal as the participants provide *Yes* and *No* answers. As shown in the figure, the participant is first presented with the stimulus after which they are expected to give their answer.

In contrast to this, in the 2I-2AFC task the participant is presented with two *observation intervals*, interval referring here to a temporal interval as can be seen from panel b of Figure 5. During each interval the participant is presented with a stimulus, one interval always contains only noise and the other contains the signal. The participant's task is then to indicate which interval they thought contained the signal. Again, the name is quite literal: the participant is presented with two *intervals* and they have to—they are *forced to*—choose their response among two *alternatives*.

Responses in the Yes/No task Decisional processing in the Yes/No task is modelled by assuming that the participant has an internal criterion for the evidence: when evidence is below this criterion the participant will respond negatively and when evidence exceeds the criterion they will respond positively.

When the participant is presented with a signal stimulus, positive responses are called *hits*; when they are presented with a noise stimulus, positive responses are called *false alarms*. (Kingdom & Prins, 2010; Wickens, 2002).

Criterion is represented by the red dashed vertical lines in Figure 6. As signal level increases, and as a consequence d' , the evidence distribution shifts to the right. Since the criterion is fixed, larger portion of the evidence distribution falls to the right of the criterion, indicating higher probability of a positive response.

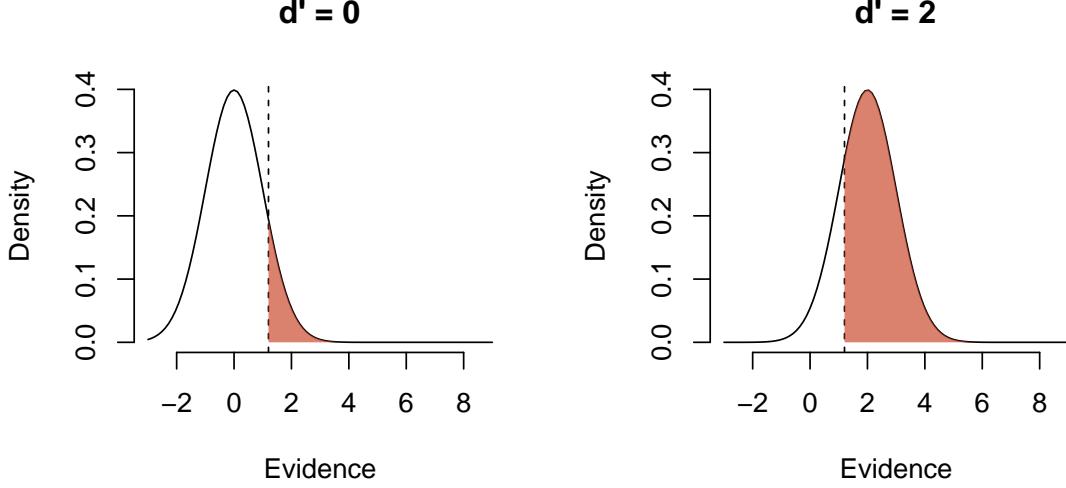


Figure 6: Binary decisions in the YesNo-paradigm. The distribution of evidence is divided into response regions.

The probability of evidence exceeding the decisional criterion on any trial can be found by finding the area (shaded part in Figure 6) under the normal distribution upwards from the criterion (see Kingdom and Prins (2010); Wickens (2002)):

$$P(R = 1) = \int_c^{\infty} N(d', 1) \quad (2)$$

where N is the normal distribution function with parameters μ (mean) and σ (standard deviation). The mean of the distribution is d' , which is the signal-to-noise ratio as calculated in Equation 1; c , the lower bound of the integral, is the decisional criterion.

The integral in this equation is usually written in terms of the cumulative standard normal distribution function, $\phi(u) = \int_{-\infty}^u N(0, 1)$. The resulting function, denoted with the Greek letter ψ , is called the *psychometric function*. (Kingdom & Prins, 2010, Chapter 4)

Usually only the positive response is considered, since the probability of a negative response is simply the complement of it, but it is more principled to show the probabilities for both the positive and negative responses:

$$\begin{aligned} \psi(S; \theta)_{\text{Yes}} &= \phi(-c + d') \\ \psi(S; \theta)_{\text{No}} &= 1 - \phi(-c + d') \end{aligned}$$

Here θ is a vector containing the parameters of the psychometric function (σ , c and β), and S is the physical signal level. Given vectors of observed signals and responses (S and R), one can find the best fitting values for the parameters, which then would be taken as the Signal Detection Theoretical interpretation of the observer's performance. In this interpretation the performance is divided into sensory (parameters σ and β) and decisional (parameter c) factors.

Since sensory and decisional processes are separate in the model, keeping the parameters that model sensory processing constant while varying the criterion will result in different predicted proportions of *Yes* and *No* responses and as a consequence hits and false alarms. That is to say that the perceptual processing capabilities can be exactly the same for two participants, but the observed amounts of hits and false alarms can be different, if one participant has different criterion. This is demonstrated in the left panel of Figure 7: in the figure, parameters relating to sensory factors are kept constant while the criterion is varied. The psychometric function drawn with a solid line represents response probabilities for an observer with more lax criterion in relation to the dashed function.

It is also important to notice that hits and false alarms are coupled. If the participant wants to avoid false alarms, they will have to adopt a stricter criterion for the evidence, and as a consequence they will inevitably also make less hits. This is because the aim of SDT is to model situation in which there is a non-zero probability of confusing signal with noise.

The coupling between hits and false alarms is demonstrated in the form the a *receiver operating characteristic* (ROC) curve (Kingdom & Prins, 2010, 161) in the right panel of Figure 7. The lines are drawn from two d' values, dashed line representing a stronger signal relative to the solid line. When the participant adopts stricter criterion, they move left on the ROC curves: indicating not only higher proportion of hits but also false alarms. Conversely stricter criterion leads to less false alarms but also less hits.

The result is that if the participant aims to avoid false alarms altogether, they will have to adopt a criterion that, theoretically, would also result in zero hits; the criterion has to be relaxed a bit in order for the participant to perform in the task in any meaningful way, resulting in some false alarms and hits, as just discussed. Generally, the participant is free to "set their own criterion", however it is possible to calculate the optimal value of the decisional criterion.

Responses in the 2I-2AFC task Since the participant is presented with two stimuli in the 2I-2AFC task, theoretically the participant is comparing two random numbers, one representing strength of evidence in the first interval and the other representing strength of evidence in the second interval. From which ever interval the larger of these two was from gets picked as the *signal* interval by the participant.

As a consequence the signal-to-noise ratio is decreased compared to the Yes/No task. Intuitively this can be understood by noticing that if one compares two things, both of which contain uncertainty, the individual uncertainties add up. Mathematically speaking, the participant can be thought of basing their decision on decision variable (here dv) that is the difference of the two random variables:

$$dv = [N(d', 1) - N(0, 1)]I \quad (3)$$

Here I is an indicator variable which is -1 when the signal is in the first interval and 1 when the signal is in the second interval. Consequently the decision rule is to respond *Interval 1* when the decision variable dv is negative and *Interval 2* when it's positive.

It is easier to think about the responses in terms of hits and false alarms, similar to the Yes/No task. This makes it possible to drop the the indicator variable from the equation, since we are only interested in the difference between the signal and noise distributions—this determines the

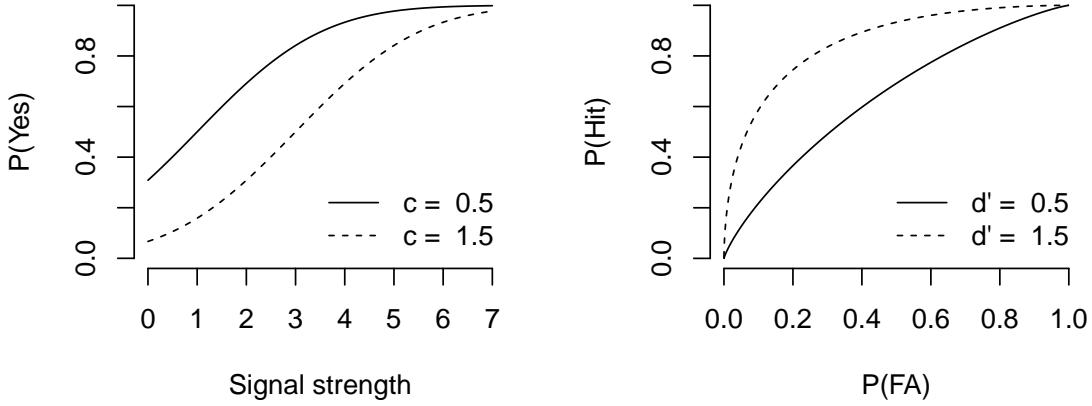


Figure 7: Decisional processing in Signal Detection Theory. Left panel: The effect of differing criteria on the predicted amounts of false alarms and hits while the parameters describing the sensory processing are kept constant. *Higher* criterion value implies stricter criterion—that the participant is less willing to respond *yes*—while a *lower* value implies more lax criterion. Right panel: Receiver operating characteristic (ROC) curve for two fixed d' -values, showing that the proportions of hits and false alarms are coupled.

probability of the participant choosing the correct interval. This basic structure of the 2AFC model is demonstrated in Figure 8.

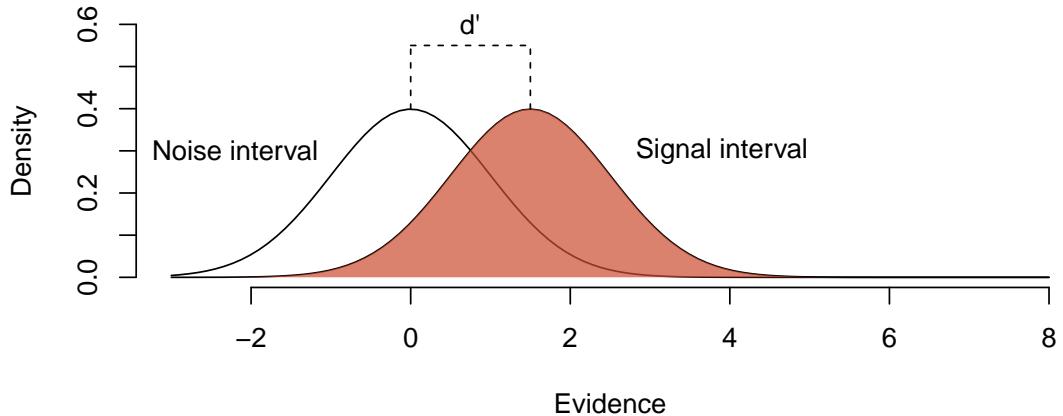
In the upper panel the zero-centred distribution represents evidence from the noise-only interval while the distribution centred on 1.5 represents evidence from the signal interval. Irregardless of whether the signal is in the first or second interval, the signal interval will be shifted d' amount. Since the probability of making the correct decision is the proportion of the decision variable distribution that exceeds 0, shifting the signal distribution to the right will increase this probability—as one would expect.

What is important to note is that the decision variable is the difference of two normally distributed random variables, both with standard deviation of one, since they represent the standardized quantity d' . It follows from the additivity of variances, and from the fact that standard deviation is the square root of variance, that the standard deviation of the decision variable is $\sqrt{1^2 + 1^2} = \sqrt{2}$. For this reason d' values in the psychometric function have to be divided by this constant:

$$\begin{aligned}\psi(S; \Theta)_{\text{Hit}} &= \phi\left(\frac{d'}{\sqrt{2}}\right) \\ \psi(S; \Theta)_{\text{False alarm}} &= 1 - \phi\left(\frac{d'}{\sqrt{2}}\right)\end{aligned}$$

The term c is dropped since it is assumed that, on average, the observer isn't biased towards either of the intervals. Some bias is probably not detrimental: any such bias would presumably

Distributions of evidence in the two intervals



Distribution of the decision variable

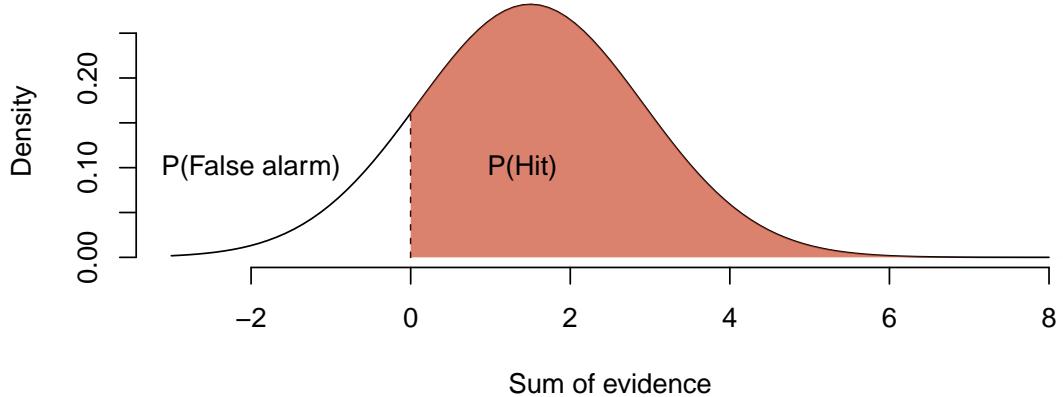


Figure 8: How responses are generated in the forced choice task. In the upper panel, the shaded distribution shows distribution of evidence from the signal interval while the non-shaded distribution shows evidence from the noise interval. The lower panel shows the distribution of the decision variable.

be most noticeable when the signal level is low, and in these cases performance is already close to chance.

2.2 General Recognition Theory

So far I have covered the case in which the participant has to make a single decision, e.g. if the pitch of the stimulus changed. In the two-dimensional case the participant has to observe two things at once. For example did pitch *and* did timbre change. The question of primary interest is if there are any *interactions* between the dimensions: does changing timbre also change response probabilities

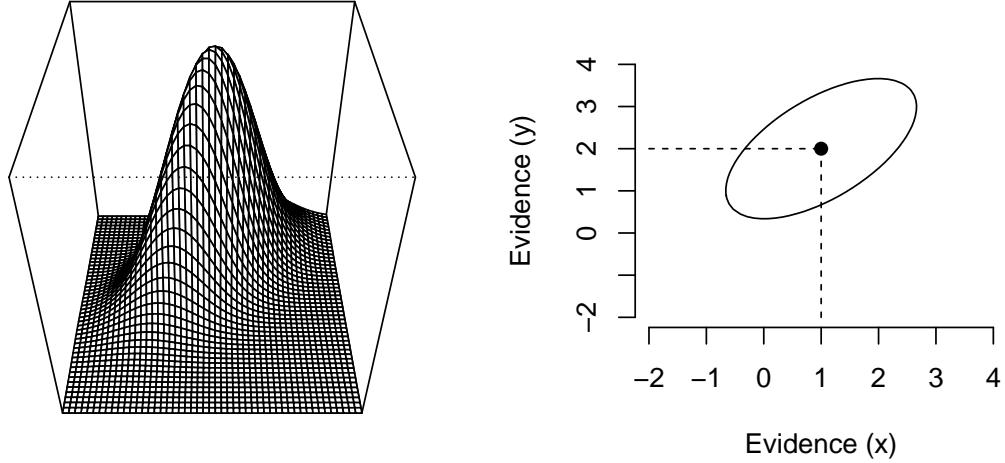


Figure 9: Two projections of a two-dimensional normal distribution with the parameters $\mu = [1.0, 2.0]$, $\rho = 0.6$.

to pitch changes.

As in SDT, also in GRT randomness in the responses is thought to arise from the evidence having a probabilistic nature (Ashby & Soto, 2015; Ashby & Townsend, 1986; Kadlec & Townsend, 1992).

In the two-dimensional case also the evidence distribution is two-dimensional. Following the assumption made earlier that evidence is distributed normally, it is possible to represent evidence in two dimensions by a bivariate normal distribution. A bivariate normal distribution can be defined by two means, which correspond to its location on each dimension, and a correlation coefficient, which corresponds to the shape of the distribution. A bivariate normal distribution with positive correlation is shown graphically in Figure 9. Left panel shows a projection of the 3-dimensional shape of the bivariate density; right panel shows the same distribution from above, as a contour. Rest of the figures in this thesis will represent bivariate distributions as contours similar to the one in the right panel.

2.2.1 Calculating d' in two dimensions

For two-dimensional stimuli, two d' values have to be calculated, one for each dimension. Similar to SDT, the d' values are the location of the evidence distribution, but now the location has to be determined in the two-dimensional space on the x and y axes separately: for example, how much evidence there is for a difference in pitch AND how much is there evidence for a difference in timbre. This implies that now there are two d' functions for which the parameters have to be estimated separately:

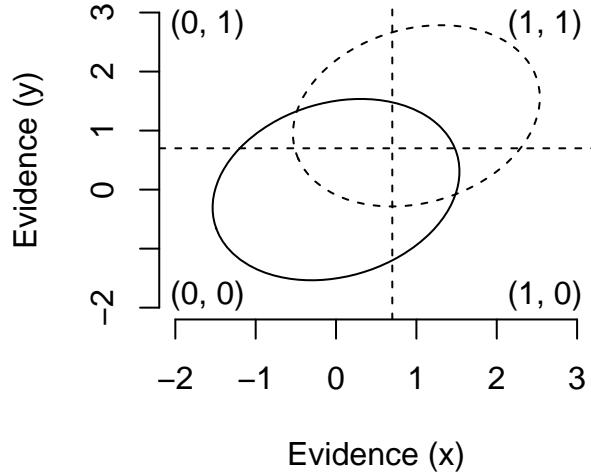


Figure 10: In two dimensions, the evidence distribution—when binary responses are required—is separated into four response regions.

$$\begin{aligned} d'_x &= \left(\frac{S_x}{\sigma_x}\right)^{\beta_x} \\ d'_y &= \left(\frac{S_y}{\sigma_y}\right)^{\beta_y} \end{aligned} \tag{4}$$

However, since the main motivation in GRT is to model interactions between the dimensions, this would be incomplete: these functions have to be coupled. I will first discuss how response probabilities are found in GRT, and then discuss the types of interference included in the models in this thesis.

2.2.2 Modelling responses in two dimensions

Whereas in SDT the participant is usually required to make a single decision on a single axis, in GRT the participant is required to make one decision per axis; in the two-dimensional case two decisions: did the signal level change on x axis and did the signal level change on y axis (a similar approach has been used also by Wickens (1992)).

As before, I will first discuss the Yes/No task before discussing the forced-choice task.

Yes/No task in two dimensions Figure 10 shows how the two-dimensional space is separated into response regions analogously to how it was shown from SDT Figure 3. The distribution represented by the solid line has $d' = [0.00, 0.00]$ and the distribution represented by the dashed lines $d' = [1.00, 1.25]$. The dashed straight lines are the criteria: the vertical line is criterion for the

dimension x and the horizontal for dimension y . These criteria divide the two-dimensional space into four response regions. Most of the solid lined distribution falls into the region labelled $(0,0)$ indicating high probability of a negative response on both dimensions. As the signal level on both dimensions is increased, as is done for the dashed line distribution, more of the distribution falls into the other regions, in this case the region $(1,1)$.

Again, similar to unidimensional SDT, the response probabilities are found by integrating over the evidence distribution, split into response regions at the decisional boundaries as shown in Figure 10. For example the probability of observing a positive response on both dimensions is found by integrating the two-dimensional normal distribution from the response criteria to positive infinity (region $(1,1)$). Whereas in the unidimensional case two psychometric functions were needed to model the two response categories, here four are needed, for each of the response possibilities:

$$\begin{aligned}\psi(\mathbf{S}; \theta)_{(No, No)} &= \int_{-\infty}^{C_x} \int_{-\infty}^{C_y} N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ \psi(\mathbf{S}; \theta)_{(Yes, No)} &= \int_{C_x}^{\infty} \int_{-\infty}^{C_y} N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ \psi(\mathbf{S}; \theta)_{(No, Yes)} &= \int_{-\infty}^{C_x} \int_{C_y}^{\infty} N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ \psi(\mathbf{S}; \theta)_{(Yes, Yes)} &= \int_{C_x}^{\infty} \int_{C_y}^{\infty} N(\boldsymbol{\mu}, \boldsymbol{\Sigma})\end{aligned}\tag{5}$$

Here \mathbf{S} is a vector containing the physical signal levels on the individual dimensions, $\boldsymbol{\mu}$ corresponds to the vector of d' -values. $\boldsymbol{\Sigma}$ is the correlation matrix, in which correlated evidence is modelled by treating the correlation coefficient as a free parameter:

$$\begin{bmatrix} 1 & \rho \\ \rho & 1 \end{bmatrix}$$

Similar to SDT, the psychometric functions are most practically thought about in terms of the (bivariate) standard normal cumulative distribution function, $\phi_2(u_x, u_y, \rho) = \int_{-\infty}^{u_x} \int_{-\infty}^{u_y} N([0, 0], \boldsymbol{\Sigma})$, in which $\boldsymbol{\Sigma}$ is defined as before. The upper integration limits can be found by calculating:

$$\begin{aligned}u_x &= -c_x + d'_x \\ u_y &= -c_y + d'_y\end{aligned}$$

Then, if responses are coded with $No = -1$ and $Yes = 1$, inputs for the ϕ_2 function are:

$$\begin{aligned}\psi_2(\mathbf{S}; \theta)_{(-1, -1)} &= \phi_2(-1u_x, -1u_y, \rho[-1 * -1]) \\ \psi_2(\mathbf{S}; \theta)_{(1, -1)} &= \phi_2(1u_x, -1u_y, \rho[1 * -1]) \\ \psi_2(\mathbf{S}; \theta)_{(-1, 1)} &= \phi_2(-1u_x, 1u_y, \rho[-1 * 1]) \\ \psi_2(\mathbf{S}; \theta)_{(1, 1)} &= \phi_2(1u_x, 1u_y, \rho[1 * 1])\end{aligned}$$

This can be defined generally with the equation:

$$\psi_2(\mathbf{S}; \theta)_{\mathbf{R}} = \phi_2([-c_x + d'_x]r_x, [-c_y + d'_y]r_y, \rho[r_x r_y]) \quad (6)$$

in which the $\mathbf{R} = [r_x r_y]$. All of the Stan (Stan Development Team, 2019a) programs in the appendices assume that responses are coded in this way.

2I-4AFC In the SDT section I discussed the 2I-2AFC task. Here I expand it to the 2I-4AFC task. There are still two temporal intervals, as was shown in the schematic in Figure 5, but now the participant has four response alternatives, since either signal could be in either interval: for example that pitch changed in the first interval and timbre in the second. The same dimension never changes in both intervals, that is, if there's a pitch change in the first interval, there never is a pitch change in the second interval.

Similar to Figure 8, Figure 11 shows the distributions of evidence from the two intervals. Note that each of the signals could be in either of the intervals, so there are four possible arrangements for the signals, but similar to the case in 2I-2AFC it is not important in which intervals the signals are concretely, what is important is the difference between the signal and noise distributions.

It is assumed in Figure 11 that κ_x^μ parameter is non-zero, which means that signal on y -dimension is able to increase the d' -value on the x -dimension—the interaction parameters, which κ_{muy} is one of, will be defined more comprehensively in the next section. This can be seen from the fact that in the left panel the mean of the noise distribution on the x -axis is slightly greater than zero. If such interaction was absent, the mean of the noise distribution would be zero on both dimensions, there would be on average no evidence for change in either dimension. The interaction shrinks the difference between the signal and noise distributions on that dimension.

As before, the participant is thought to base their decision on a decision variable that represents the difference of evidence in the intervals. As a consequence the d' values on both dimensions have to be divided by $\sqrt{2}$ (cf. section on 2I-2AFC). The distribution of the decisional variable can be seen in the right panel of Figure 11. The thick black lines mark boundaries between *hits* and *false alarms*. The participant, in this particular scenario, is very likely to score a *hit* on the y -axis, but on x -axis they are almost as likely to pick the wrong interval as the correct; the most likely response pairs being (FA, Hit) and (Hit, Hit).

As in the Yes/No task, responses can be coded with -1 and 1. Here, *false alarms* are coded with -1 and *hits* with 1. The general expression for the psychometric function is then

$$\psi_2(\mathbf{S}; \theta)_{\mathbf{R}} = \phi_2\left(\frac{d'_x}{\sqrt{2}}r_x, \frac{d'_y}{\sqrt{2}}r_y, \rho[r_x r_y]\right) \quad (7)$$

Comparison of the Yes/No and 2I-4AFC tasks The main difference between the Yes/No and 2I-4AFC tasks is what the participant's decision is directly related to. In the Yes/No task the decision is about the *appearance* of the stimulus. As was already discussed, this means that the performance of the observer cannot be decoupled from their decisional criterion. This adds one more parameter to be estimated from the data.

In the 2I-4AFC task the decision about the *interval*. This means that it is not necessary for the participant to adopt a criterion relating the strength of evidence to a response. Also, since the participants are choosing an interval, which is by its nature more neutral decision, it is usually thought that this kind of task is less susceptible to bias (Kingdom & Prins, 2010, Chapter 6).

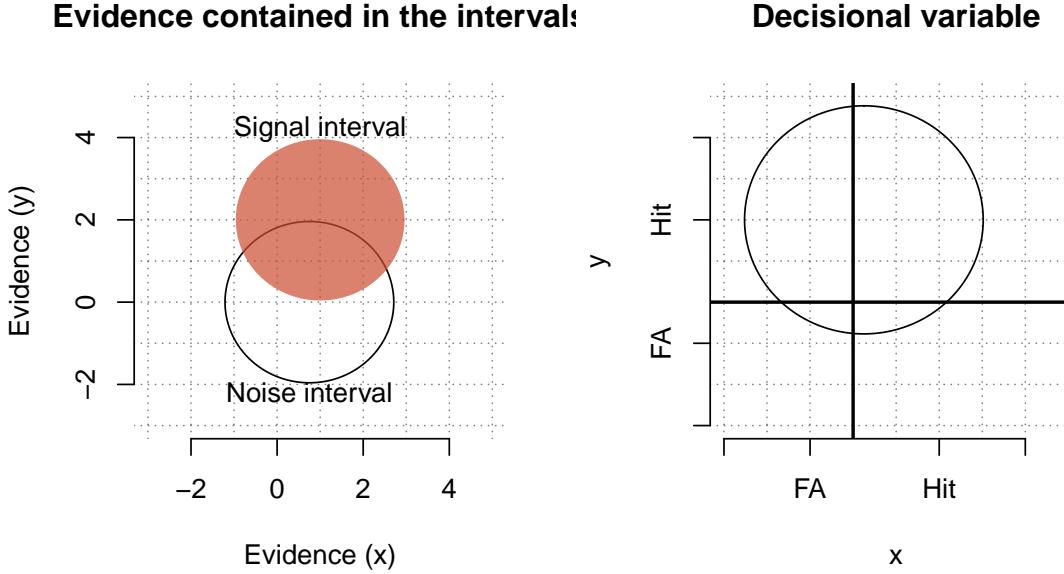


Figure 11: In the 2I-4AFC task, the participant is assumed to base their responses on the difference in the evidence from signal and noise distributions (left panel). Variance of the distribution of differences is the sum of the variances of the individual distributions (right panel).

Since on each trial two stimuli have to be presented in the 2I-4AFC task, it is somewhat slower to administer. However, if the observer knows, in the Yes/No task, that some non-zero signal is present on each trial, nothing prevents them from answering *yes* on each trial—they will always be correct, and seemingly able to detect the faintest of signals. To counter this, some amount of *catch* trials are usually included. These are trials during which the signal level is zero. During these trials no information about the sensory processing—the parameters σ and β —is accumulated. So even though individual trials are faster to complete, there has to be more of them.

The issue of catch trials is exacerbated in the multidimensional case. This is because in the one-dimensional case the issue is simply solved by including some noise stimuli as catch trials in the experimental run, resulting in two kinds of stimuli being presented: noise and signal stimuli. In the two-dimensional case there has to be three kinds of catch trials: noise only stimulus on either of the dimensions or both. Otherwise if the participant became aware of e.g. that the adaptive procedure only rarely selects a stimulus in which there's a noise stimulus on one dimension—i.e. $[0, S]$ or $[S, 0]$ —, they might become biased towards choosing the responses [No, No] and [Yes, Yes]. The linear decision bounds used in this work are not able to model this kind of response bias; an issue which will be discussed later.

Then, on theoretical grounds, the only difference to be expected is the omission of the criterion parameters from the 2I-4AFC model—reduction in sensitivity is taken into account by incorporating the term $\sqrt{2}$.

2.2.3 Interactions between the dimensions

As already discussed GRT aims to model *interactions* between the dimensions. When the one dimension has an effect on the other, this is called *interference*. There are three kinds of interactions considered in this thesis. Two of these interactions are defined as *couplings* between the psychometric functions, and the third interaction is the correlation of the evidence distribution, but since including correlation coefficient was already conclusively discussed in the preceding section, I will only discuss the couplings here.

Two kinds of couplings are modelled. These are demonstrated graphically in Figure 12. Panels in the left column demonstrate coupling between means while panel on the right side demonstrate coupling between standard deviations. Top row shows effects of the interactions on the evidence distributions while the bottom row shows the effect on the psychometric function on one dimension.² In all panels dashed lines show the situation in the absence of interaction while the solid lines are drawn in the presence of interaction.

Coupled means, κ_μ Coupling between the means of the distributions is modelled by letting the other dimension influence the d' values:

$$\begin{aligned} d'_x &= \left(\frac{S_x}{\sigma_x}\right)^{\beta_x} + \kappa_x^\mu S_y \\ d'_y &= \left(\frac{S_y}{\sigma_y}\right)^{\beta_y} + \kappa_y^\mu S_x \end{aligned} \quad (8)$$

As can be seen from the lower left panel of Figure 12 this kind of interactions shifts the psychometric on the affected dimension. Here, the presence of interference has increased the d' values and consequently probabilities for positive responses all along the range of the function.

Coupled standard deviations, κ_σ Another possibility is that the standard deviations are coupled. This kind of situation can arise if for example the other dimension is length: shorter signals contain less information, thus making frequency judgments noisier (see Houtsma (1995, p.274), also worse accuracy with shorter signal times has been reported by e.g. Townsend, Hu, and Kadlec (1988) implying more noise).

Similar to κ_μ , terms are included for letting the other dimension influence the σ parameter. However, since the σ parameters are constrained to positive real numbers (standard deviation can't be negative), this interaction is defined on the logarithmic scale (for a similar approach in probit modelling, which is closely related to SDT as will be discussed shortly, see Freeman, Keele, Park, and Salzman (2018)):

$$\begin{aligned} \sigma_x &= \exp(\log \sigma_x + \kappa_x^\sigma S_y) \\ \sigma_y &= \exp(\log \sigma_y + \kappa_y^\sigma S_x) \end{aligned} \quad (9)$$

When the model parameters relating to response boundaries or criteria—as is the case in here for the Yes/No model—they also have to be divided by the standard deviation. This is intuitively clear

²mathematically the union of the psychometric functions $\psi S_{(Yes, Yes)}$ and $\psi S_{(Yes, No)}$.

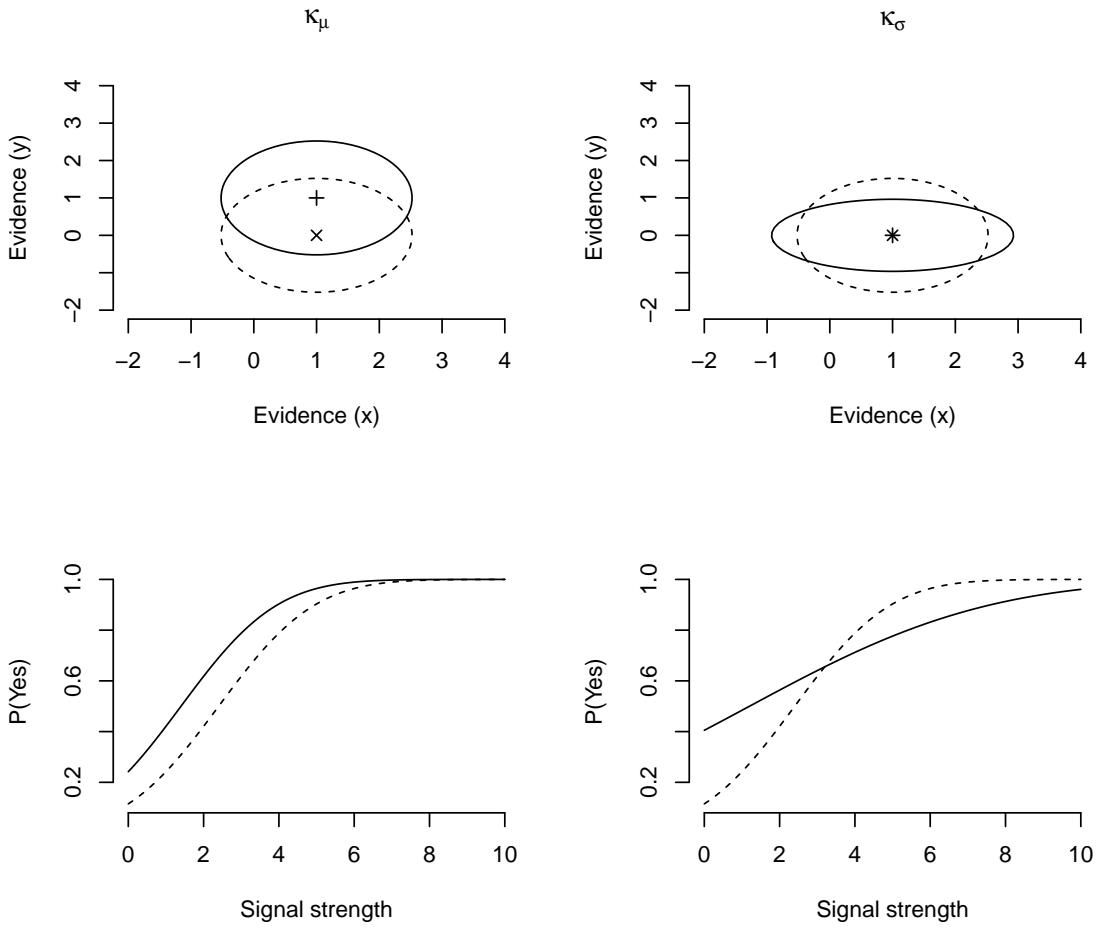


Figure 12: The two kinds of couplings in the model. The contours in the top row depict bivariate normal distributions. Dashed lines indicate the shape and location of the distribution in the absence of interaction; the solid lines depict them in the case that interaction is present. Curves in the lower row are psychometric functions for a single dimension.

from the observation that as the spread of the evidence distribution becomes larger, the parts of the evidence distribution falling into the response regions become increasingly equal. In other words response probabilities get closer to 0.25, and a larger denominator for the term c will achieve just this:

$$\psi_2(\mathbf{S}; \theta)_{\mathbf{R}} = \phi_2\left([- \frac{c_x}{\sigma_x} + d'_x]r_x, [- \frac{c_y}{\sigma_y} + d'_y]r_y, \rho[r_x r_y]\right) \quad (10)$$

This means that the definition of criterion is a bit different from that of when no coupling between standard deviations is present. In the model with constant variance, the criterion corresponds to the Z-score of the false alarm probability (see Kingdom and Prins (2010, Chapter 6)): criterion = $\phi^{-1}(P(FA))$, in which ϕ^{-1} is the standard normal inverse cumulative distribution function; in the model with coupled standard deviations the whole first term corresponds to the Z-score of false alarm probability: $\frac{\text{criterion}}{\sigma} = \phi^{-1}(P(FA))$. This means that in this model the criterion parameter is more accurately the geometric location of the decisional criterion in the decisional space, independent of the variance of the evidence distributions. However, since it is easier to interpret false alarm probabilities, prior in the models is placed on the term $\frac{\text{criterion}}{\sigma}$ —what this means in practice will be discussed in the section on Bayesian statistics

From the lower right panel of Figure 12 one can see that the effect is to draw the response probabilities towards 0.5 on a single dimension. Note that this is in agreement with what was said earlier about response probabilities being drawn towards 0.25: since the psychometric function in the figure is shown only for one dimension, it represents the probability of a positive response regardless of the other dimension.

Model with both kinds of couplings The most realistic case is one in which both kinds of couplings are included in the model. This is achieved easily by using Equation 8 and calculating the σ terms as in Equation 9.

Note on terminology The couplings, κ_μ and κ_σ , are sometimes called *between stimuli/distributions interference* (Silbert, Townsend, & Lentz, 2009) and the situation in which they are absent can be referred to as *sensory independence* (Ashby & Soto, 2015). Correlation can be called *within stimulus/distribution interference* (Silbert et al., 2009) and the absence of correlation as *perceptual independence* (Ashby & Soto, 2015). In addition, κ_μ can be called *mean-shift integrality* and κ_σ *variance-shift integrality* (Ashby & Maddox, 1994).

2.3 Lapsing behaviour

It is possible that during a psychophysical task, people sometimes *lapse*. That is, for some reason the response they give doesn't reflect the cognitive process of interest. This might be due to e.g. lapse in attention, a coding error in the program, or a slip of a finger. In the classic GRT models, lapsing behaviour has not been taken into account, even though in the unidimensional case lapses have been shown to be able to exert considerable bias on the parameter estimates of the psychometric function (Wichmann & Hill, 2001).

Estimating lapsing rate from data can be problematic (Treuein and Strasburger (1999); Wichmann and Hill (2001)), and this problem can be worse for adaptive procedures, at least for trying

to adaptively estimate the lapsing rate (Prins, 2012). Since the values of the psychometric function that are closer to zero are affected both by the lapses and the decision criterion (in the Yes/No task), most of the information gained about lapsing behaviour comes from lapses happening at high intensities. What makes matters worse is that the lapsing response doesn't necessarily differ from what the participant would've answered otherwise, so the distinction between lapses and genuine response is ambiguous (for further discussion on these problems, see Prins (2012)). For these reasons I fixed the values of λ and γ during the adaptive estimation phase. However, during the analysis of the psychophysical data I estimated the parameter λ from the data, a decision, which will be discussed in more detail in the analysis section.

I model lapsing behaviour by using a (hierarchical) mixture model (see Zeigenfuse and Lee (2010)), which will be discussed in Section 3.2 *Hierarchical models*.

2.4 Critical discussion

2.4.1 Relationship to other models

Classic GRT I will begin by discussing what I will call a *classic* GRT model, and argue for the two main modifications implemented in this work: first abandoning the ubiquitous categorization task and second incorporating some functional form between the physical signal levels and the internal quantities.

In the classic model the signals of interest are two-dimensional, the participant is required to hold one decisional criterion per dimension, and the method of constant stimuli (MOCS) is employed. Method of constant stimuli means that the levels of stimuli are fixed prior to the experiment. When two levels are used per dimensions, the experiment is called a 2x2 identification experiment. A possible stimulus set is demonstrated in Table 1. I call this the classic model since it is the most widely—and in some cases the only—discussed type of model in GRT related literature (see e.g. Ashby and Soto (2015); Ashby and Townsend (1986); Cohen (2003); Kadlec and Townsend (1992); Silbert (2010); Silbert and Thomas (2013)).

Table 1: An exemplar set of stimuli that could be used in a 2X2 identification experiment. The exact levels of e.g. high and low pitch are chosen a prior by the experimenter.

		Pitch	
		Low	High
Timbre	Bright	Low pitch + Bright timbre	High pitch + Bright timbre
	Dark	Low pitch + Dark timbre	High pitch + Dark timbre

The logic in this kind of task differs from that of discrimination task. Instead of responding to differences in the stimuli, the participant's task is to categorize³ the stimuli as if categorizing apples and oranges into their respective piles. The image of categorizing something into piles is surprisingly apt, since in the classic speeded classification experiments—which have been influential to the study

³Note that the term *categorization* has its own specific meaning distinct from identification in GRT literature (see Ashby and Soto (2015)), but here I'm using the term rather colloquially—and in any case, I believe the problems discussed here to apply both categorization and identification tasks.

of multidimensional perception, including GRT–(Garner, 1974) the participants would categorize visual stimuli printed on cards.

Categorization is sensible in the case of clear differences, e.g. categorizing extremely different pitches such as 100 Hz and 5000 Hz. However, in the case of minute differences that are close to the detection threshold, it is more likely that the subject bases their decisions on perceived *differences* between consecutive stimuli. One of the assumptions of the model is that single responses are statistically independent (see e.g. Wickens (1992, p. 218)), but if indeed the responses are based on perceived differences between consecutive stimuli, this assumption is gravely violated.

Another difficulty is that the participant would have to hold accurate representations of the stimuli in their memory. Again this is simple when categorizing apples and oranges, but much more difficult in a psychophysical experiment in which the levels are to begin with selected in such a way that they are easily confusable. As already discussed, Signal Detection Theory is used specifically in a situation in which there's a high probability for confusing the stimuli.

Of course all kinds of assumptions are always violated, since any statistical model is a simplification of the real and complex cognitive processes, but this is not an excuse to *not* use procedures that are more in line with what is known *a priori*. It is for these reasons that I've opted to base the model presented here explicitly on discrimination instead of categorization or identification.

As per the second modification, the addition of functional relationships, in the classic GRT model each stimulus is represented by its own bivariate normal distribution (see e.g. Ashby and Soto (2015)). For example, in the aforementioned 2X2 categorization task parameters for four bivariate normal distributions would be estimated. If more stimuli are used also more decision boundaries are needed in the categorization task: a separate set of decision boundaries between each category is required.

This can lead to difficulties when trying to interpret the model. Consider for example the hypothetical case in Figure 13: it is very difficult to come up with a simple summary of how the dimensions x and y are related. This is not only a problem for concise communication of results, but it is a theoretical problem: in my view modelling should strive towards more complete models, and staying agnostic about the functional forms of interactions does not advance this goal.

Related difficulty is that it is hard to know, e.g. in the case depicted in Figure 13 which of the features of the data are purely coincidental and which of them accurately represent real dimensional interactions. Statistically speaking one could be *overfitting*; a problem that is also recognized in the context of classic GRT models by F. A. Soto, Zheng, Fonseca, and Ashby (2017). Also, often the fit of the model is assessed by comparing the predicted response probabilities with the observed proportions (see for example Silbert et al. (2009, Figure 4)), but if a separate distribution is fitted to each stimulus, the model is overly flexible, and it is likely that *any* data set would be matched very closely by the model. As such, this close match is not evidence for the veracity of the model; it is evidence for its flexibility.

The current model rectifies these problems. Interactions are summarized by the κ and ρ coefficients, so it's easy to make inferences. Since the model is constrained by the explicit functional relationship between signal levels and d' , it's less prone to the type of overfitting described.

Process model The interactions in the model can be conceptualized by hypothesizing a process structure to the model, as in Ashby (1989), Ashby (2000) or Cohen (2003). This is demonstrated

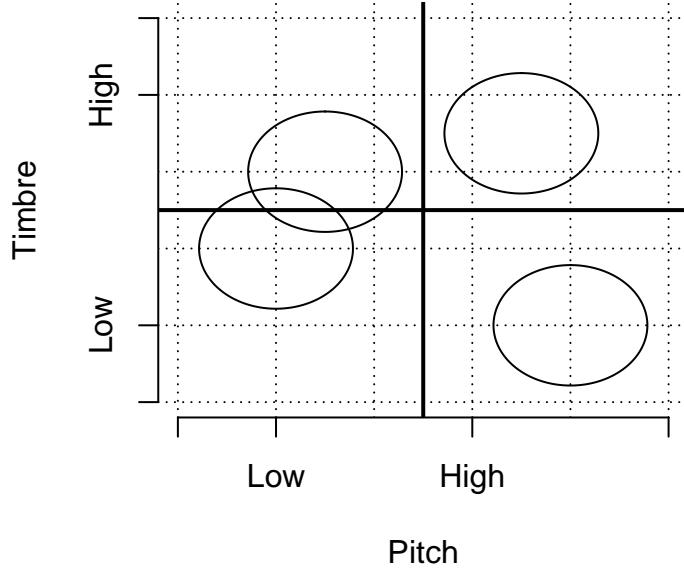


Figure 13: An example of a classic GRT model fitted to a 2X2 categorization data. Each of the stimuli are fitted with their own bivariate normal distribution.

in Figure 14. On each channel the physical signal, S , is transduced to the internal representation of the signal, s , after which it is judged against the decisional criterion, c , and the participant gives their response, R . The straight lines from a stage to the next—for example from S_x to s_y —denote an influence to the average level of the variable while curved lines inside a stage—for example from s_y and s_x —denote correlated noise.

The effect of the signal level on the average level of the internal signal on the other channel is usually referred to as stage 1 interaction, the correlation of noise between the internal representations is referred to as stage 2 interaction, and the dependence of the criteria on the internal level of the signal or correlation between their noise is referred to as stage 3 interaction. Stage 1 and 2 interactions are referred to as *perceptual* interactions, and stage 3 interaction is referred to as *cognitive* interaction. (Ashby, 1989, 2000; Cohen, 2003).

In the current model, stage 1 interactions are modelled by the κ terms while the stage 2 interaction is modelled by the correlation parameter ρ . Stage 3 interactions are not modelled, since, as the model currently stands, they are not identifiable. This issue will be discussed in greater detail later.

Relationship to Generalized Linear Models Generalized Linear Models (GLM's) are widely implemented in statistical analysis (see e.g. Kruschke (2015); Skrondahl and Rabe-Hesketh (2004)). The theoretical links to well-known models in statistics mean that the model should be easily generalizable to different situations, and that there exists a wide range of theoretical and empirical literature.

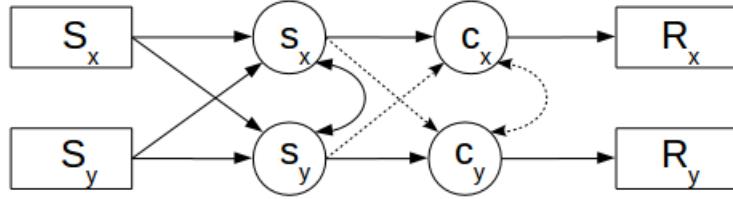


Figure 14: A process interpretation of the model. Variables inside rectangles are observed, whereas variables inside squares are latent.

Decarlo (1998) describes the unidimensional SDT model as a generalized linear model. In the case of normally distributed noise, a *probit model*, which is how already Gustav Fechner—who is sometimes regarded as the founder of psychophysics—modelled categorical judgments in the 19th century (discussed in Stigler 2003, Chapter 7).

The models—with the caveat mentioned at the end of this section—used in this thesis can be thought of as a *non-linear multinomial* probit model. Non-linear (Box, Hunter, & Hunter, 2005, p. 379) since the relationship between S and d' included the term β ; multinomial, since in a binary choice probit model the responses are assumed to follow the binomial distribution, but in the multidimensional model, since there are more than two choices, the responses are assumed to follow the multinomial distribution Skrondahl and Rabe-Hesketh (2004).

As Generalized Linear Models in general, also the model here can in theory use any distribution as the model for the errors (the use of other than normal distributions in GRT has already been discussed in Ashby and Townsend (1986), but to my knowledge, other distributions besides normal have *never* been used). The main limitation is mathematical and computational: how to effectively calculate the two-dimensional integrals needed for the response probabilities. There can also be other issues, e.g. for the typical multinomial logistic model, no correlation is usually modelled (Skrondahl & Rabe-Hesketh, 2004), which can prove to be a too limiting assumption.

In the context of GLM's, the modelling of lapses is can be called *robust regression* (Kruschke, 2015, p. 635) since it relaxes the distributional assumptions somewhat, and isn't so prone to estimation errors due to outliers. However it should be noted that this requires the lapsing parameters to be fixed Skrondahl and Rabe-Hesketh (2004). This is not true for all of the models used in this thesis.

2.4.2 Nonidentifiabilities and other limitations of the model

What do the latent distributions really represent? The main assumption of SDT, and consequently GRT, is that the participant bases their decisions on noisy latent quantities. Therefore, it is important to discuss what those latent quantities actually represent.

In SDT literature writers are generally careful to call the latent quantities evidence (Verde, MacMillan, & Rotello, 2006; Wickens, 2002) or judgments (Stigler, 2003, p. 247). In GRT related

literature, however, the latent quantities are thought to be more closely related perceptions. Ashby and Soto (2015) call them *perceived values*; in Ashby and Townsend (1986), Kadlec and Townsend (1992) and Silbert et al. (2009) they are *perceptual effects* and the space (in which the distributions are defined in) is defined as *perceptual space*; F. A. Soto et al. (2017) uses the term *perceptual representation*.

As is clear from the preceding discussion on SDT and GRT, I've taken the more traditional way of calling the latent variables evidence, but even in this case one has to be careful in considering *what* the evidence is about. Concretely the latent quantities are a sum of *everything* that creates variability in responses: noise in encoding the signals, internal noise due to e.g. blood pressure or breathing, transient noises from sneezing or swallowing etc., lapses of attention, criterial changes, non-stationarity of threshold, effects of learning and so on.

It is clear that for example in a pitch discrimination task many of the aforementioned factors don't necessarily directly effect the perceived amount of pitch difference, but can and will affect the judgments or amount of evidence in other ways, for example by masking the signal and limiting information that way.

Often there seems to be two strong implicit assumptions. First, that either the first factor would be significantly greater than any of the others, and second that any effect is sufficiently free from other influences. While these assumptions might not be that dangerous in SDT models, the problem is exacerbated in GRT: there is virtually no information about how the aforementioned factors influence e.g. inferences about interactions between the dimensions, a problem recognized already by Silbert et al. (2009). It is my impression, based on my own unpublished simulation studies, that the factors just mentioned can in some situations lead to significantly incorrect inferences.

Validity of the process model described earlier relies heavily on these strong assumptions and our ability to differentiate between different sources of variation. It is my belief that in the current formulation, using straightforward perceptual experiments, one can not properly identify these, and indeed has to rely on strong prior assumptions. For this reason I would, at this point, see the process model more as a helpful conceptualization than anything else.

Conceptualization of the latent variables affects the way interactions are understood theoretically. Based on the preceding discussion I hope it is clear that it is different to say that e.g. timbre influences the distribution of *evidence* about pitch than it is to say that timbre influences the distribution of *perceptual effects* of pitch.

Criterial noise Criterial noise can be defined as random variation in the criteria around some central value, in the present context the central value would most naturally be thought of as the parameter c in the model. If this variation is assumed to be Gaussian, it is not identifiable from other sources. This issue has been discussed in SDT literature already by Wickelgren (1968)⁴; in the context of GRT Ashby (2000) has discussed the role of criterial noise. In the context of SDT there have been proposals for identifying the magnitude of criterial noise using experimental manipulations (see e.g. Benjamin, Diaz, and Wee 2009; Cabrera, Lu, and Dosher 2015; Kellen, Klauer, and Singmann 2012), but the application of these methodologies to GRT will not be considered further here.

⁴Wickelgren (1968) talks about *unidimensional strength theory*, but the discussion is directly applicable to the present situation.

In practice, any criterial noise will be included in the σ estimates, since, as already discussed, these represent to sum total of all variation. With regard to GRT the problem is if the criterial noise between the dimensions is correlated: this will be included in the ρ estimates. Another possibility is that criterial changes—e.g. between sessions or during a long session—would also affect the κ estimates. To my knowledge, no systematic exploration of this possibility exists; in my own non-systematic simulations I have noticed that indeed in some cases such criterial shifts can affect the parameter estimates and lead to erroneous inferences, however, this will not be explored any further in this thesis.

Non-orthogonality of decisional boundaries Another topic tying into the modelling of decisional processing is the (possible) non-orthogonality of the decisional boundaries. In the GRT literature non-orthogonal decisional boundaries are seen as failure of *decisional separability* (Ashby & Soto, 2015), since non-orthogonality implies that the decision is contingent on the level of the signal on the other dimension.

In the present paper the focus is on orthogonal decisional boundaries. Non-orthogonal boundaries could be incorporated in the model, Ennis and Ashby (2003) have shown a simple method for calculating the non-orthogonal integrals to evaluate the response probabilities. However, what could prove to be problematic is that the method, in my experience, is rather slow.

Another problem is the identifiability of non-orthogonal decisional boundaries: F. Soto, Vucoovich, Musgrave, and Ashby (2015) claim that their model makes decisional separability identifiable—contrary to the classic model (Silbert & Thomas, 2013)—but their result has been refuted on mathematical grounds by Silbert and Thomas (2016): any non-orthogonal boundaries can be thought of as a transformation of the space to one in which the boundaries are orthogonal and correlation is non-zero.

It follows that, in the models with orthogonal decisional boundaries, any non-orthogonality would be seen as non-zero values of the ρ parameter, which should be taken into account when interpreting this parameter.

3 Bayesian statistics

Bayesian statistics is at the core of this thesis: the adaptive algorithm to be used is based on the idea of representing uncertainty about parameters as probability distributions over them. Bayesian statistics will also be used for making inferences about data, and Bayesian philosophy will guide the model development process. In this section I will provide a short introduction to the central concepts that are needed for understanding the adaptive algorithm and, later, inferences made about the data.

3.1 Basics of Bayesian statistics

The core idea of Bayesian statistics is the updating of prior information with observed data to arrive at a synthesis of them. This synthesis is called the posterior probability distribution, which brings me to another idiosyncrasy of Bayesian statistics: whereas in "classical" statistics one would calculate point estimates, that is for example the average difference between conditions, in Bayesian statistics the goal is evaluate how well *any* plausible value fits (I use the word *fit* colloquially) the observed difference resulting, indeed, in a probability *distribution* which is taken as the estimate. This can be taken as an intuitive quantification of uncertainty: if wide range of values fit the data almost as well, this indicates that there's more uncertainty about what can be said for example about the difference between the conditions, in contrast to situation in which only a narrow range of values fit the data.

Both the prior information and the posterior distribution are represented as probability distributions over the parameters, θ . Usually some common probability densities with simple functional forms are used, for example normal distributions. Equation 11 represents updating prior information mathematically (Kruschke, 2015):

$$P(\theta|Data) \propto P(\theta)P(Data|\theta) \quad (11)$$

In this equation $P(\theta|Data)$ is the posterior distribution, $P(\theta)$ is the prior distribution and $P(Data|\theta)$ is the likelihood function, i.e. the probability of observing the data given the parameter values.⁵ *Data* refers to observations collected during a psychophysical experiment, that is, pairs of stimuli and responses, and θ are the parameters of any model in this thesis.

The likelihood function determines *how* the observed data updates the prior information. In practice, this means that there needs to be a model which determines how observed data translates into probabilities about parameter values, or how the model learns from observed data. Here, the likelihood is defined by any of the models (M in Equation 12) introduced earlier in this thesis. Consequently the likelihood for a particular θ is simply the product of the probabilities of the observed responses given the θ (and the model)⁶:

⁵In the full Bayes' theorem the posterior distribution is normalized with the integral over the possible observed values, $P(Data)$, assuring that it always integrates to one and is a proper probability distribution: $P(\theta|Data) = (P(\theta)P(Data|\theta))/P(Data)$ (Kruschke, 2015), but I think the equation without the normalizing constant captures the main idea more clearly. This is why the symbol \propto ("proportional to") is used in Equation 11.

⁶Remembering that multiplying probabilities corresponds to the logical operand *and*, this means the we are interested in finding out the probability of the response on the first trial AND the response on the second trial and on the third trial and so on, depending on how many responses there are.

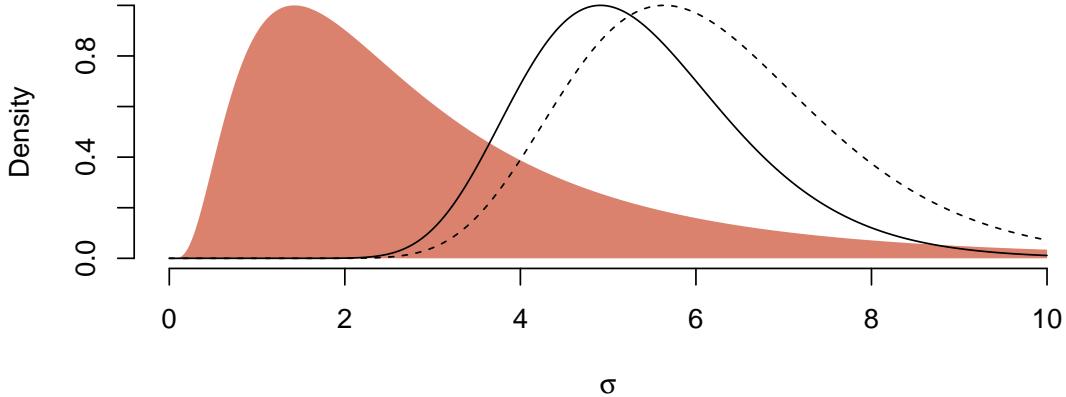


Figure 15: A simplified example of the basic concepts of Bayesian inference. The shaded area represents prior knowledge about the parameter σ . Dashed line is $P(\text{Data}|\theta)$ and the solid line is the posterior probability density, $P(\theta|\text{Data})$

$$\prod_{t=1}^n \psi_2(\mathbf{R}^t, \mathbf{S}^t; \theta, M) \quad (12)$$

Note that usually the logarithmic likelihood function is used of numerical stability, $\sum_{t=1}^n \ln \psi_2(\mathbf{R}^t, \mathbf{S}^t; \theta, M)$, which is how the likelihood is defined in all of the R and Stan programs. Stan programs can be found from the appendices.

Updating prior information is demonstrated for a single parameter in Figure 15. The shaded region represents the prior distribution $P(\theta)$, which in this case, is a log-normal distribution. The prior distribution is centred around $\sigma = 2$, indicating that values close to 2 are seen as most probable prior to observing any data.

After some data has been observed the prior distribution is updated with evidence, $P(\text{Data}|\theta)$, represented by the dashed line in Figure 15, which then results in the posterior distribution $P(\theta|\text{Data})$, represented by the solid line. Observing data, in this particular case, has reduced our uncertainty about likely values for the parameter σ , which is evident in how the probability density has become more concentrated. In addition to this, the mode of the distribution has shifted, indicating shift in what values are found to be likely on average.

How the prior distribution is determined is a complex question. First of all, the prior distribution should take into account mathematical constraints, for example here the fact that σ must always be positive, which is why I used the log-normal distribution in the example. Another aspect is that one can encode current knowledge in a meaningful way by choosing the prior appropriately. A *weak* or *non-informative* prior mostly acts as a way of regularizing the inferences and increasing computational stability, while a *strong* or *informative* prior can encode specific subject-matter knowledge (for accessible discussion on priors and discussion on weak and strong priors see *Prior*

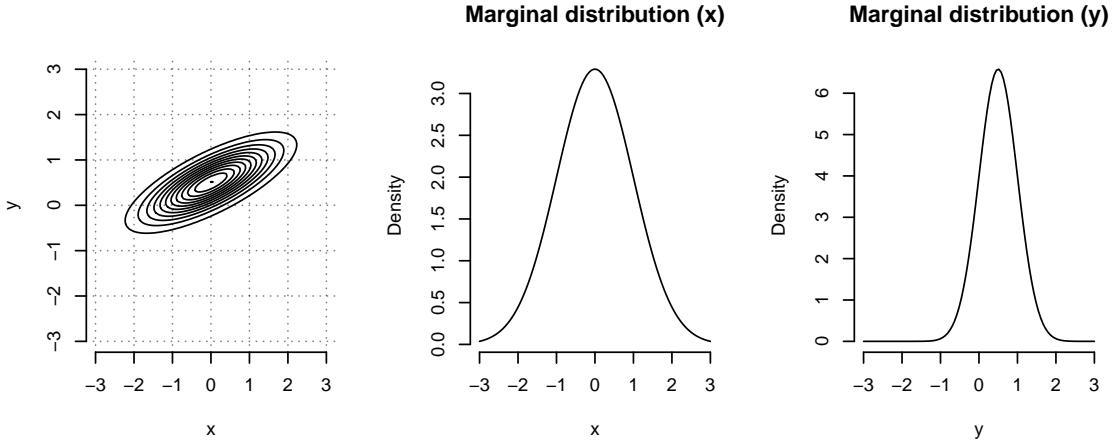


Figure 16: A two-dimensional posterior probability density and its marginals.

choice recommendations (n.d.)).

For determining priors for psychometric functions Lee (2018) recommends informative priors, and using prior predictive distributions—that is, distributions of psychometric functions—as a way of checking that the priors make sense and lead to psychometric functions that are realistic. One problem with this approach is that usually priors that assume independence among the parameters are used (an issue that is also discussed in *Prior choice recommendations* (n.d.)), but it is likely that the parameters of the psychometric functions have strong correlations between them.

For example psychometric functions with high values for σ but low values for β produce psychometric functions that can be unrealistically slow in reaching high performance levels, especially if values for κ_μ or κ_σ parameters are significantly large. The issue of encoding such dependencies in the prior is left for future work, I will be using priors which are independent out of convenience, and choosing values that make sense for the parameters in isolation, which is admittedly not optimal.

Lastly, one particularly useful feature of priors is that it's easy to expand models into having a multilevel structure through them, which will be discussed in more detail in the next section.

Of course, in models of any complexity prior and posterior probability densities are defined over a multidimensional space, e.g. in this thesis the dimensionalities range from 7 to 27. This means that we are not interested in, for example, probabilities for values of σ parameter alone but probabilities for combinations of parameter values, for example, that $\sigma = 2.0$ and $\beta = 0.7$ and $\kappa_\mu = 0.1$ and so on for all of the parameters of the specific model.

Often it is useful to summarize multidimensional posterior density as its marginals. Most analyses done in this thesis will use the marginals of the posterior probability densities somehow. Figure 16 demonstrates the relationship between a two-dimensional posterior probability density and its margins. The leftmost panel shows the two-dimensional distribution from above, and the two other figures show its marginals, which in this case correspond to viewing the two-dimensional "bump" from either axis.

3.2 Hierarchical models

The term *hierarchical model* can refer to many different kinds of models⁷. I will be using two kinds of hierarchical models: 1) Model in which the parameters are assigned distributions 2) Models embedded in a higher level mixture model. Both kinds of models will be discussed separately.

Models in which parameters have distributions

Often these kinds of models are described as being models in which priors have (hyper)priors (e.g. Kruschke (2015, p. 225) talks about ". . . *hierarchical chain of dependencies among parameters.*".). However, I think it's clearer to think this as an extension of fitting e.g. separate linear models to different data sets: instead it is possible to fit the individual linear models together on a higher level, and assign the groups of parameters (e.g. all intercepts for all of the models) distributions. This is sometimes known as *random effects modelling* in the frequentist setting.

The general idea is demonstrated in Figure 17. The index i is an index for the *group* to which the observation y_{ij} belongs to, j indexes observations inside that group. Each observation is assumed to be drawn from a normal distribution, but each group gets a unique set of parameters of the linear model, i.e. there are as many α and β as there are groups. Without the hierarchical structure this would simply correspond to as many separate linear models as there are groups.

Here, however, these parameters are each assigned a distribution (prior), which in this case is the normal distribution, whose parameters are unknown and are inferred from the data. Note that for the σ parameters the normal distributions are truncated at zero. At the highest level each parameter of these distributions gets a hyperprior, which in this case for each parameter is the standard normal distribution—taking into account again that σ is truncated at zero.

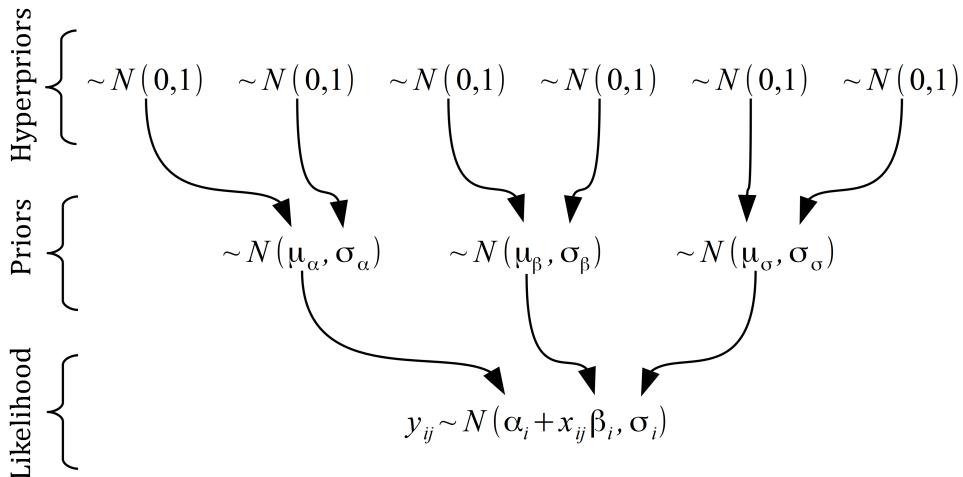


Figure 17: A schematic illustration of a hierarchical linear model in which each of the parameters is given a distribution for which the parameters are inferred from the data. This creates a distinct multilevel structure to the model.

⁷Andrew Gelman has collected a bunch of different names for hierarchical models in to a blog post: <https://statmodeling.stat.columbia.edu/2019/09/18/all-the-names-for-hierarchical-and-multilevel-modeling/> – a fact which highlights the multifacetedness of hierarchical models

In other words, there are—for example—separate α parameters for each group, but *at the same time* as the values⁸ for the individual α parameters are found, the parameters for the prior distribution from which they are assumed to be drawn from are also found; in this way the individual parameters depend on the priors, and the parameters of the priors depend on the individual values.

Since the parameters are related on a higher level, some information is shared between them. One feature is that all of the means are adjusted towards the common mean, which functions analogously to multiple comparisons adjustments in frequentist statistics (Gelman, Hill, & Yajima, 2012), which is the main reason for including the hierarchical structure here.

Models embedded in a mixture model

Often the same data set can be fit by multiple models, or there are many plausible candidates for models of the cognitive processes of interest. These are often also called *mixture* models, since the idea is to model observed data as a mixture of the different models (for the general statistical idea see Keller and Kamary (2018)). Since the individual models are embedded in a larger scale structure that models the mixing proportions, this kind of model is naturally also thought of as a hierarchical model (see Kruschke (2015, Chapter 10)).

The simple two-level model I describe corresponds with how lapsing behaviour is often modelled in the psychophysical literature (see e.g. Zeigenfuse and Lee (2010) for a theoretical summary or e.g. Lesmes et al. (2015) for a practical implementation). The model is presented in figure 18. At the lowest level, the figure shows two models, one labelled *lapses* and the other *cognitive*, embedded in the higher level model. Both of the lower level models aim to explain the same set of observations, $y_1, y_2 \dots y_i$, the y terms, in this case, consisting of pairs of stimuli ($S = [S_x, S_y]$) and responses ($R = [R_x, R_y]$).

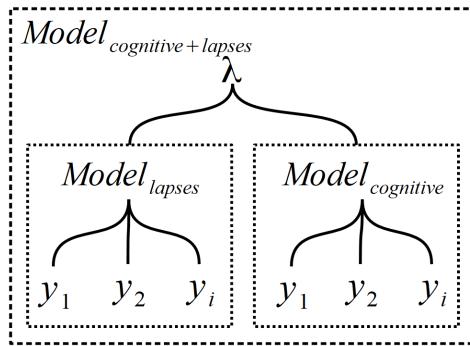


Figure 18: A simple hierarchical model, which shows how the two models, one modelling lapses and the other the cognitive processes of interest, are related on a higher level through the coefficient λ .

The cognitive model can be any of the models defined in Section 2.2 *General Recognition Theory*. The lapsing model (which was discussed in general terms in Section 2.3 *Lapsing behaviour*) is defined by the single parameter γ , which is a multinomial distribution containing the response probabilities for the individual response categories. These probabilities do not depend on the signal levels, or in

⁸To be precise, not *values* but posterior distributions, as was discussed earlier.

fact, on anything else outside them (see Section 2.2.2 *Modelling responses in two dimensions* for how the responses are coded):

$$P(y_i) = \begin{cases} \gamma_1, & \text{if } \mathbf{R} = [-1, -1] \\ \gamma_2, & \text{if } \mathbf{R} = [1, -1] \\ \gamma_3, & \text{if } \mathbf{R} = [-1, 1] \\ \gamma_4, & \text{if } \mathbf{R} = [1, 1] \end{cases} \quad (13)$$

This simply says that if, for example, on a certain trial the response $[1, -1]$ is observed, the likelihood is incremented by γ_2 . I will be assuming that all responses are as likely, i.e. $\gamma = [0.25, 0.25, 0.25, 0.25]$. This implies that the parameters of the lapsing model are not inferred from the data, partially due to the aforementioned (see Section 2.3 *Lapsing behaviour*) difficulty of inferring details of lapsing behaviour from psychophysical data.

The contributions of these two models are defined by the coefficient λ , which defines the weight of each model to the total likelihood:

$$\begin{aligned} P(y_i|\Theta_{\text{cognitive + lapses}}, \text{Model}_{\text{cognitive + lapses}}) = & \lambda P(y_i|\theta_{\text{lapsing}}, \text{Model}_{\text{lapsing}}) + \\ & (1 - \lambda)P(y_i|\theta_{\text{cognitive}}, \text{Model}_{\text{cognitive}}) \end{aligned} \quad (14)$$

The vector $\Theta_{\text{cognitive + lapses}}$ contains all of the parameters for the lower level models and λ : $\Theta_{\text{cognitive + lapses}} = [\lambda, \theta_{\text{lapsing}}, \theta_{\text{cognitive}}]$.

3.3 Approximating the Posterior Probability Densities

Unfortunately posterior probability distributions usually lack closed-form solutions, and they can only be approximated. A problem specific to Bayesian adaptive testing is that selecting the optimal stimulus on each trial requires one to have an approximation not only of the current posterior distribution but also for the possible posteriors given potential stimulus values (this will be elaborated in Section). This means that the approximations have to be quick.

This problem has been solved in a few different ways. Kontsevich and Tyler (1999) used *grid approximation* which is based on dividing the posterior distribution into equally spaced intervals and then calculating posterior probability density at each point⁹. In essence, this is similar to approximating a function by drawing its graph by calculating its value at some finite number of points. Thus, grid approximation is simple, but its downside is that it becomes computationally costly when dealing with high dimensional posterior distributions. E.g. if one has a 9-dimensional posterior distribution and wishes to represent it with 25 discrete points per dimension, one would need to calculate 25^9 values, which is infeasible even given the computing capabilities of modern computers. The dimensionality of the cognitive models presented in this thesis range from seven to twelve.

For faster computation with higher dimensional posteriors in the context of Bayesian adaptive testing, e.g. Kujala and Lukka (2006) and DiMatta (2015) have proposed the use of random sampling to build fast approximations of the posteriors. This is the main methodology I also used.

⁹For a general description of this method see Kruschke 2015, p.144)

But how can a distribution be approximated by drawing random samples from it?

As an example, let us say that the distribution we are interested in approximating is a $\text{Gamma}(3, 1)$ distribution. We might be interested for example it's expected value

$$E[\text{Gamma}(3, 1)] = \int_0^\infty x \text{Gamma}(x|3, 1) \quad (15)$$

and if we did not have access to a closed-form solution, we could draw random samples from it and use those to evaluate that integral. Figure 19 shows the idea behind this. The rightmost panel shows the target distribution and the center panel shows the samples drawn from it. You can see that there are more samples where the density of the target distribution is higher (around where x is between 2 and 4). In the leftmost panel a histogram of the samples from the center panel is superimposed on the target density.

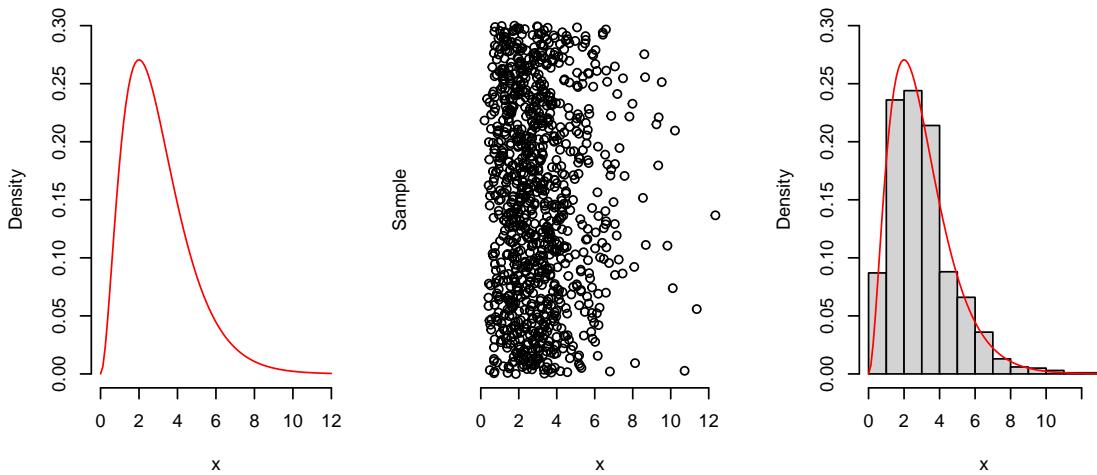


Figure 19: Using random samples to approximate a $\text{Gamma}(3, 1)$ distribution. Right panel: the target distribution. Center panel: random samples drawn from that distribution.

In practice, however, we are usually not able to draw samples from the posterior distribution directly. For this reason different kinds of algorithms have been developed for drawing random samples from arbitrary distributions. In this work I used two different algorithms for drawing random samples from the posteriors. For the final posterior analyses (e.g. when analyzing data from the psychophysical experiments) the models were programmed in Stan language, which uses a No-U-Turn-Sampler (Stan Development Team (2019b, Chapters 15 & 16)). Stan is easy to use, it generally produces accurate approximations of the posteriors—consequently approximations build using Stan are used as a ground truth. The downside is that it's too slow for implementing it during an experiment.

For this reason, during both the simulated and real experiments, a resample-move algorithm was used for quick approximations of the posterior distributions.

Stan's sampling method is taken here as a black-box method, and won't be discussed in further

detail, however I will describe the Resample-Move algorithm since they it has direct consequences on the implementation of the adaptive algorithm. I will also provide short discussion on Laplace approximation, which is based on approximating posterior distributions as multivariate normal distributions, since it was used as a practical back-up method for the adaptive algorithm.

3.3.1 Resample-Move algorithm

The Resample-Move algorithm, as described by Chopin (2002) involves using *sequential importance sampling* and the occasional *rejuvenation* step (see Algorithm 1).

Central idea is to represent the posterior distribution as weighted *random samples*. This is demonstrated in the left panel of Figure 20, in which a one-dimensional normal distribution, as represented by the solid line, is approximated with a set of uniformly sampled values along the x -axis, called particles. The heights of the particles (the dashed lines) correspond to the weights of the particles (proportionally; the weights should always sum to unity). One can use the locations of the particles and their weights to estimate e.g. the mean and variance of the underlying distribution, using standard formulae for discrete random variables.

Algorithm 1 Sequential importance sampling with rejuvenation

- 1: At the beginning of the experiment, draw N particles from the prior distribution and set uniform weights.
 - 2: Observe a data point.
 - 3: Update the particle weights according to the observation.
 - 4: **if** $N_{\text{eff}} < N/4$ **then**
 - 5: Resample particles.
 - 6: Move resampled particles.
 - 7: Return to step 2.
-

The complete particle sets thus is most conveniently thought of as a matrix with parameter values and their weights. Table 2 demonstrates this idea for the 2I-4AFC model. The particle set consists of randomly sampled values of the parameters, θ . Each row corresponds with one set of parameter values and a weight, w . In Figure 20 weights were shown for a single parameter, here weights correspond with their "heights" in a 7-dimensional space.

Table 2: An example of a particle set, which consists of sets of parameter values, θ , and associated weights.

θ								
	σ_x	σ_y	β_x	β_y	κ_x	κ_y	ρ	w
1	1.2	2.5	0.2	2.6	-0.2	0.5	-0.1	0.10
2	0.5	0.3	1.1	0.6	0.4	0.1	0.2	0.05
3	0.2	1.7	1.8	1.5	0.3	-0.1	0.3	0.15
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
N	0.4	1.5	0.7	0.8	-0.1	-0.5	0.75	0.20

As more data is observed, fewer and fewer particles have non-zero weights. This is natural, since usually at the beginning of an experiment there's more uncertainty about the specific values of the

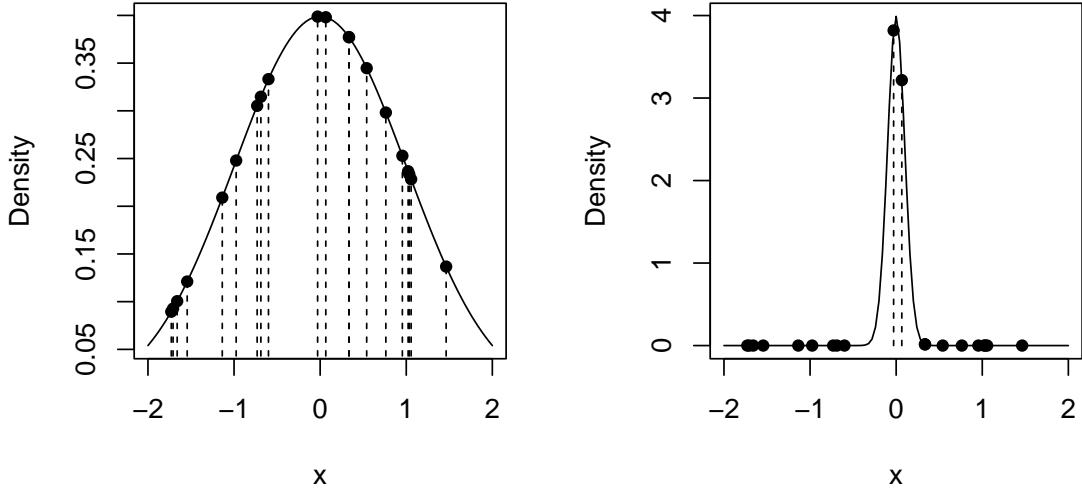


Figure 20: Demonstration of sequential importance sampling and particle degeneracy. When the posterior probability becomes more concentrated, after some data is observed, less particles have weights that differ significantly from zero.

parameters, but as the experiment progresses, more and more particles can be singled out, or given smaller weights. When the underlying distribution becomes more concentrated, the weights of most particles approach zero, as is shown in Figure 20. This is called *particle degeneracy*. A simple way of quantifying it is the effective sample size (Speekenbrink, 2016):

$$N_{\text{eff}} = 1 / \sum_i^N w_i^2 \quad (16)$$

When the particle set becomes too degenerated—here defined as the point when $N_{\text{eff}} < N/4$ —it has to be rejuvenated.

Rejuvenation is done by first resampling the particles. Resampling means sampling new particles with replacement proportional to their weights. I used *multinomial resampling*. This is similar to how a roulette wheel works, the difference being that in a roulette wheel each number has a uniform probability of being chosen; here the "sectors" for some particles are wider, and thus they have greater chance of being chosen. This results in particles with greater weights being replicated, and particles with lower weights being removed from the particle set. This does not effectively combat particle degeneracy which is why the algorithm has a *move* step. (Chopin, 2002)

During the move step proposed particles are drawn from what is called a *proposal distribution*. As the proposal distribution I used a multivariate normal distribution with means and standard deviations corresponding to those of the current posterior distribution, as suggested by Chopin. These proposals are accepted with the probability: $p(\text{proposal}_i|y)/p(\text{old particle}_i|y)$, that is the

ratio of the posterior probabilities of an old particle and a proposed particle. It is easy to see that for proposals, which have higher posterior probability than the old particles, the probability is greater than one, implying that they are always accepted. (Chopin, 2002).

Particle methods have been used in adaptive estimation for example by DiMatta (2015), although they did not implement the rejuvenation step, making their implementation unrealistic for the longer experimental runs used here, and by Kujala and Lukka (2006) who use a more sophisticated rejuvenation step, which, due to its increased computational cost, was also not realistic.

3.3.2 Laplace Approximation

Laplace approximation can be conceptualized as a two-step process: 1) find the mode of the posterior, and 2) find the second derivatives of the posterior at the mode. The matrix of the second derivatives at the mode is known as the *Hessian* matrix. Since derivatives are related to the rate of change, the Hessian matrix can be used to construct the covariance matrix for the (multivariate) normal distribution at the mode. Figure 21 shows two univariate implementations, in both panels the solid lines represent the target distributions (which are to be approximated) and the dashed lines show the approximation. In the left panel the target distribution is a t-distribution ($df = 1$), and in the right panel a Gamma distribution ($\alpha = 3$, $beta = 1$).

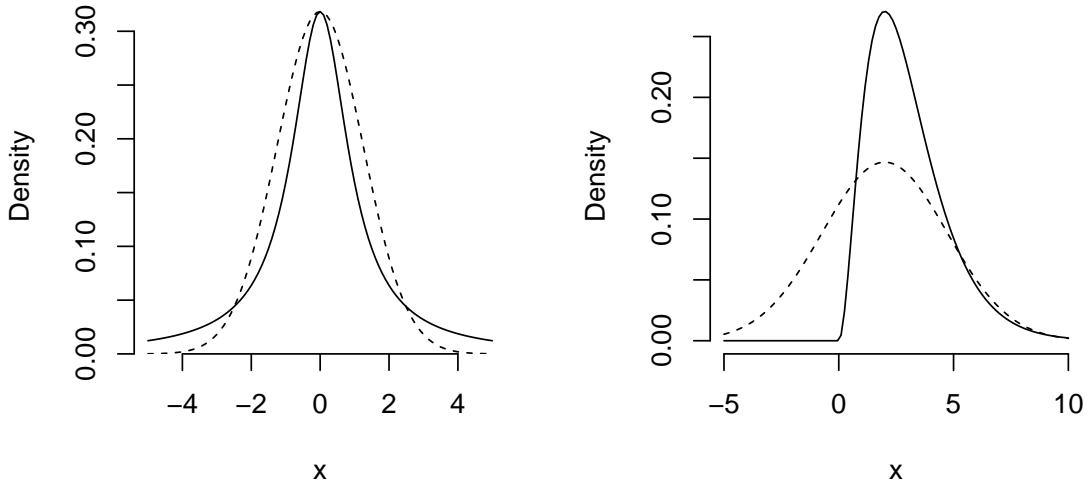


Figure 21: Laplace approximations of two univariate distributions: t-distribution (left panel) and a shifted Gamma-distribution (right panel). Solid lines show the target distributions while the dashed lines show the approximations.

There are two things to note: 1) the farther the target distribution is from normal, the worse the approximation will be; 2) the approximation might extend beyond the support of the target distribution. The second point is evident in how the approximation for the Gamma distribution

extends to negative numbers, even though the target distribution's support is constrained to positive real numbers.

Even though Laplace approximation has been used in the context of Bayesian adaptive estimation (Y. Shen and Richards (2013) apply it to the adaptive estimation of auditory filters) I use it here as just a backup: for details see Section [see 4 Some practical considerations](#). Even though the models used here can be reparametrised in a way that all dimensions have support on the real number line (this will be discussed shortly), in my experience correlations between the dimensions can be non-linear, which is something that can't be approximated well with a multivariate normal distribution.

3.4 Bayesian inference in the General Recognition Theory

In this work Bayesian approach is also used during the analysis of the data. Classic GRT studies have been dominated by frequentist methods, and e.g. definitions of interactions have relied on testing for the statistical significance of the parameter values associated with them (e.g. Ashby and Soto (2015); Wickens (1992)), to my knowledge, Silbert (2010) is the only GRT related work using Bayesian analysis framework. Contrary to the majority of studies, I won't be doing any explicit significance tests, nor am I using the usual *Type I/II* error framework (Christensen, 1997, pp. 470 - 471), since there are many problems associated with this.

For example if interactions are selected based on their statistical significance, effect sizes are likely to be exaggerated (Gelman, 2018); but the more damning criticism is that the *p*-value doesn't differentiate between *effect* or *no effect* (Greenland et al. (2016)), i.e. between *interaction* or *no interaction*, and in general, there has been a substantial amounts of criticism targeted towards focusing on binary decisions based on *p*-values and instead a push to instead focus on the effect sizes and uncertainties associated with them (see e.g. Greenland et al. (2016); Kline (2004); Steiger and Fouladi (1997)).

This shift should not be seen only as a trivial data analytical decision: it should be seen as reflecting a wider push in the behavioural sciences to move away from binary decisions (see e.g. Amrhein, Korner-Nievergelt, and Roth (2017)), and—in the light of the topic—relating to the interpretation of interactions as being graded properties of the stimuli (as discussed in Kemler Nelson (1993)), instead of something that either is or isn't. On the other hand, in Null Hypothesis Testing (NHST, see Kruschke (2015, Chapter 11)), as it is usually implemented, the main goal is to reject some uninteresting default model, while in Bayesian analysis—as it's defined in Gelman et al. (2014)—the explicit goal is to develop models that describe the data generating processes realistically. I find this goal to be much intuitively meaningful—rejecting an uninteresting default model only tells us that the specific model is bad, it does not tell us what models are good. In GRT this usually means rejecting some model that assumes no interactions between the dimensions—the important question of *what kind* of interactions there then actually are can't be reliably answered by just using NHST.

Posterior predictive checks An important part of Bayesian work flow is *model critique* (Gelman et al., 2014, Part II). The starting point is the Boxian philosophy that "*All models are wrong . . .*" (Box et al., 2005): since we can be a prior certain that all models are wrong to a degree, it becomes important to check the model against the data to see *how* it is wrong, and if the mismatch is significant.

In the classical GRT models, model checking has usually been limited to comparing the predicted and observed response probabilities (see for example Silbert et al. (2009, Figure 4)). However this suffers from the fact that, as discussed earlier in Section 2.4 *Critical discussion*, the classic GRT models are prone to over-fitting, which means that the predicted and observed response probabilities will *always* be fairly close to each other. This is not evidence for the validity theoretical assumptions behind the model, but rather just an artifact of the flexibility of the model.

In Bayesian setting, *posterior predictive checks* can be used to assess the fit of the model to the data. The idea is to use the posterior distribution to predict new data. The distribution of predicted data can be thought of as what the model "thinks" about the data. If there's large discrepancies between what the model thinks and data, this is a clear indication that the model does not sufficiently capture all relevant features of the data. (Gelman et al., 2014, Chapter 7).

Developing good posterior predictive checks can be difficult, and there is not a set way to conduct them. In this thesis I will be using simple checks that are based on dividing the data into discrete categories and counting the number of positive responses in each category. This is not exhaustive, but, as will be seen in the analysis section, is sufficient for some making some preliminary observations.

Unconstrained parameterization By their very nature, σ parameters are constrained to positive real numbers, since they represent standard deviations. Such hard constraints don't exist for the β terms, but since the psychometric functions are assumed to be monotonically increasing, I also constrained the β parameters to positive real numbers too. The correlation parameter, ρ , is bound between -1 and 1.

As already mentioned, these bounds have to be taken into account when deciding what distributions to use as priors for the parameters. One solution is to choose distributions with matching supports, i.e. for example log-normal distributions for the parameters bound to positive real numbers. Another possibility is to *reparametrise* the model in such a way that all of the parameters have support on the real number line. The latter approach is taken here.

Priors for σ and β —and as a consequence the parameters themselves—are defined on the logarithmic scale. Prior for ρ is defined on the inverse hyperbolic tangent scale. In other words

$$\begin{aligned}\sigma' &= \log(\sigma) \\ \beta' &= \log(\beta) \\ \rho' &= \text{atanh}(\rho)\end{aligned}\tag{17}$$

Both of these are widely used transformations in statistics for parameters with this kinds of constraints (see Stan Development Team (2017, Chapter 22)). In practice this means that the *inverse* transformations have to be implemented in the likelihood calculations: since e.g. since β' is on the natural log scale, we recover the original parameter by using the exponential function: $\beta = \exp(\beta')$.

For the d' calculations this implies that Equation 4 becomes

$$\begin{aligned} d'_x &= \left(\frac{S_x}{\exp(\sigma'_x)} \right)^{\exp(\beta'_x)} + \kappa_x S_y \\ d'_y &= \left(\frac{S_y}{\exp(\sigma'_y)} \right)^{\exp(\beta'_y)} + \kappa_y S_x \end{aligned} \quad (18)$$

and for the psychometric function (compare this with Equation 6):

$$\psi_2(\mathbf{S}; \theta)_{(R_x, R_y)} = \phi_2([-c_x + d'_x]r_x, [-c_y + d'_y]r_y, \tanh(\rho')[r_x r_y]) \quad (19)$$

A similar transformation would be applied also to Equation 7.

An important aspect of this process to note is that since the priors are normal on the *transformed scale*—i.e. $\log \sigma \sim N(\mu, \sigma)$ —, they are not normal on the *original scale* (see Kruschke 2015, pp. 729 - 732). Indeed, the current parameterisation implies that for example the prior for σ is a log-normal distribution.

Since the parameters in this reparameterised version have support on the real number line, this means that symmetric proposal distributions can be used more readily during the move-step of the Resample-Move algorithm. Remember that new values for the parameters are proposed during the move-step. If the parameters have constraints, and a symmetric proposal distribution is used, the proposals may lay outside the constraints, which means that they are automatically rejected¹⁰.

¹⁰Another possibility would be to use a non-symmetric distribution, but this would make the move-step somewhat more complicated, and this possibility is not explored in this thesis.

4 Adaptive estimation

The motivation behind adaptive psychophysical testing can be traced to intuitive enough starting point. Suppose that the researcher is interested in finding out what is the faintest possible sound that the subject is able to detect. It wouldn't make sense to waste time presenting them with stimuli that they can always detect, rather, if the subject has very low threshold for detection, it would make sense to reduce the level of the stimuli until the subject starts making errors, which is how adaptive *staircase* methods work (c.f. (Kingdom & Prins, 2010, Chapter 5)).

A popular subclass of adaptive methods are based on Bayesian statistics. Watson (1983) introduced the *QUEST* procedure which was influential in its use of Bayesian statistics in determining which stimuli to choose. The *QUEST* procedure was used only for estimating one parameter, threshold (analogous to the parameter σ in this thesis), and it was later expanded by King-Smith and Rose (1997) in their *ZEST* algorithm for estimating the threshold and slope of the psychometric function (analogous to the parameters σ and β in this thesis).

However, it wasn't until the work by Kontsevich and Tyler (1999) that Bayesian Adaptive Estimation as how it would be known today would emerge. The aforementioned *QUEST* and *ZEST* procedures used *ad hoc* heuristics for determining the placement of the stimuli: in the *QUEST* procedure stimuli are placed near the threshold, in the *ZEST* procedure, in order to estimate the slope of the psychometric function, stimuli are placed above and below the threshold. A symptom of such *ad hoc* methods is that Treuwein (1995) lists 22 adaptive procedures in his Table 1—and this does not include procedures for two-dimensional stimuli.

In Kontsevich and Tyler (1999) the stimulus selection heuristic is entirely based on principles of Bayesian statistics, and by that virtue entirely abstract and thus generalizable—at least in theory, no *ad hoc* rules for choosing stimuli are needed. Indeed, this principle has been applied with success to different tasks and models in psychophysics, and also to estimation of multidimensional psychometric functions.

The approach I use differs from previous work on adaptive methods with multidimensional signals. In the approaches by DiMattina (2015), Lesmes, Jeon, Lu, and Dosher (2006), J. Shen, Sivakumar, and Richards (2014); Y. Shen and Richards (2013) and Kujala and Lukka (2006) Bernoulli distributed responses are modelled, that is, only two response categories are used. In the *QUEST+* procedure introduced by Watson (2017) it is possible to use multinomial distribution as a model of the responses, but the approach is quite different, e.g. the correlation is not modelled and the approach is not based on GRT.

The algorithm in Kontsevich and Tyler (1999) chooses stimuli that minimize the *expected entropy* of the posterior distribution. In simple terms, entropy refers to the evenness of a probability distribution (Kruschke, 2015, p. 365). This is demonstrated in Figure 22. The figure depicts two multinomial distributions, in which the categories refer to the possible responses for two-dimensional stimuli (see Section 2.2 *General Recognition Theory*). The distribution on the left is maximally even, while the distribution on the right side has more probability assigned to the responses (0,0) and (1,1). Consequently, the distribution on the right has less entropy.

In terms of posterior distribution this means that stimuli, that make the posterior probability densities as concentrated around a single point as possible are chosen. Recall from the previous section on Bayesian statistics that the spread of the posterior distribution for the parameters quantifies uncertainty related to the values of the parameters. Thus, minimizing the spread of the posterior

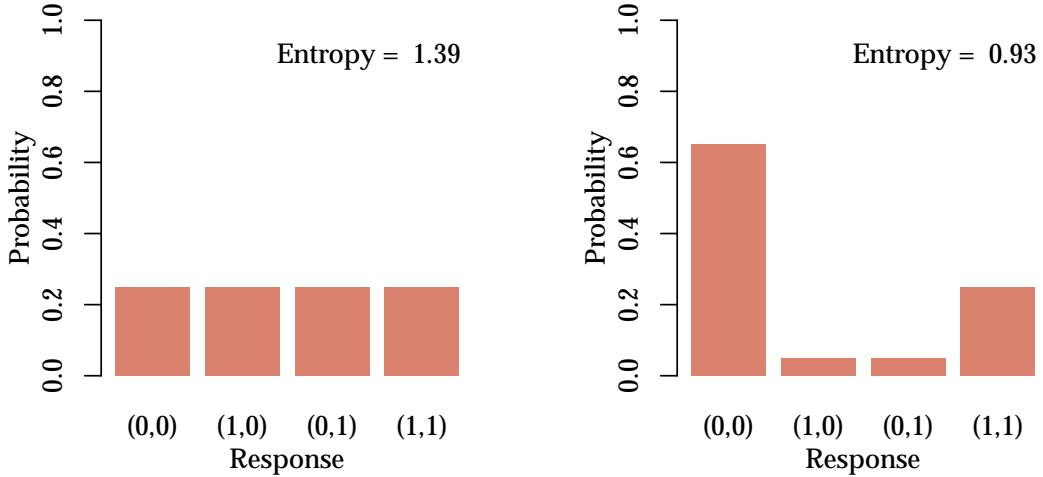


Figure 22: Two multinomial distributions that represent the probabilities for the possible response categories in the model. The distribution on the right is less even, and as a consequence has less entropy.

entropy minimizes uncertainty about their values.

The reason for using the distribution of responses in the above example is that Kujala and Lukka (2006) and Kujala (2011) show that minimizing the entropy of the response distribution is equivalent to minimizing the entropy of the posterior distribution of the parameters. This is, in my opinion, simpler than minimizing the entropy of the posterior distribution directly as done by Kontsevich and Tyler (1999).

The intuition behind this is that the response probabilities for a given stimulus are calculated from distributions of parameter values. Instead of continuous distributions, it is easier to think randomly sampled vectors of parameters as was shown in Table 2 when posterior approximation methods were discussed in Section 3 *Bayesian statistics*. If there's only one parameter vector with non-zero weight (indicating low entropy in the posterior distribution), there's only one possible response; if there's a wide range of possible parameter vectors (indicating high entropy in the posterior distribution), these will also result in wide range of possible responses.

Kujala and Lukka (2006) and Kujala (2011) give equations for calculating information gain using a set of IID particles—that is, particles with uniform weights. Since a weighted set is used by the algorithm I use (see Section 3 *Bayesian statistics*), the equations are modified to accommodate this:

$$h\left(\sum_{i=1}^N w_i \psi_2(\mathbf{S}; \theta_i)\right) - \sum_{i=1}^N h(w_i \psi_2(\mathbf{S}; \theta_i)) \quad (20)$$

in which h is the entropy of a discrete distribution, in this case, a distribution the response probabilities. Each θ_i is a "slice" of the matrix of particles (a single row of Table 2), defining a

single set of parameter values. h is (Kontsevich & Tyler, 1999):

$$h(p) = - \sum_{i=1}^N p_i \ln p_i \quad (21)$$

On each trial the stimulus that maximizes Equation 20 is chosen.

Some practical considerations In the preceding discussion I described the adaptive algorithm theoretically. In implementing the algorithm in R (R Core Team, 2019), some practical issues came up.

First, a look-up table of the bivariate integrals (the function ϕ_2) was pre-computed to speed up the calculations.

Second, maximizing the information gain function using optimization algorithms proved prohibitively slow and unreliable. For this reason the possible stimuli were chosen from a grid of 15 stimuli. This allowed me to pre-compute response probabilities for all of the stimuli, which sped up the algorithm considerably.

Third, during the adaptive run, whenever the particle set was rejuvenated, the ranges for the grid, from which the stimuli were chosen were re-computed. The lower and upper limits for stimuli were found from inverting the d' function, and then finding stimulus levels that would correspond to $d' = 0.1$ for the lower limit and $d' = 2$ for the upper limit (Equation 22). Posterior means for the σ and β parameters were used, in addition, to take into account posterior uncertainty, these were shifted by 1.96 standard deviations to produce psychometric functions with either steep slope (high β) and low threshold (low σ) or shallow slope and high threshold. These corresponded to the most extreme scenarios: if the observer's psychometric function has a shallow slope and high threshold, response probabilities change relatively slowly and the range of stimuli has to be larger; the inverse is true if the threshold is slow and the slope steep.

$$\begin{aligned} S_{\text{lower}} &= 0.1^{(1.0/\exp(E[\beta]+1.96*SD[\beta]))\exp(E[\sigma]-1.96*SD[\sigma])} \\ S_{\text{higher}} &= 2.0^{(1.0/\exp(E[\beta]-1.96*SD[\beta]))\exp(E[\sigma]+1.96*SD[\sigma])} \end{aligned} \quad (22)$$

Fourth, to ensure that the particle approximation is accurate, three parallel particle algorithms were run. If the algorithms diverged—here defined as the marginal means differing more than .2 on the prior scale—Laplace approximation was used to start the particle sets again.

5 Simulations

In the simulated experiments, the performance of the model was evaluated using simulated observers. A set of parameter values— σ , β , κ_μ and so on—were chosen for the simulated observer, and then used to generate simulated responses. This results in simulated data sets that otherwise function similarly to datasets generated by human observers, but since the generating parameters are chosen by the experimenter the ground truth is known. This is useful for 1) diagnosing problems with the model and 2) observing how sample size affects parameter estimates and their standard errors.

Evaluating models by simulating data is used in Bayesian statistics in general (Gelman et al., 2014, Chapter 10) and in psychophysics for studying the performance of the models (e.g. in the presence of lapsing trials in Prins (2012)) or adaptive algorithms (e.g. by Kontsevich and Tyler (1999)). Oron (2007) has criticized simulation studies as a substitute for thorough analytical solutions: the results of simulations—as I mentioned in the Introduction—are contingent on several *ad hoc* assumptions, and as such it is not always clear if the results are generalizable beyond the initial study. However this can also be seen as the strength of simulation studies: it's possible to easily test models and assumptions that escape analytical solutions. In this case, no analytical solutions exist, so simulation is the only tool at use.

I am using simulations to answer two main questions:

1. How much more efficient the adaptive algorithm is in relation to sampling stimuli from a fixed grid—if at all?
2. How well can generating parameters be recovered?

These questions are closely related, since relative efficiency of the algorithms (Question 1) is defined here by the quantities that are also used to evaluate Question 2.

There are two quantities of interest. First is defined by taking the means of the marginal posterior distributions as point estimates and calculating the squared differences to the generating parameters. This quantifies squared bias and variance of these estimators. The second quantity is the standard deviations of the marginal posterior distributions. This is used as a measure of how much uncertainty about the parameters is left after the data collection process. The goal, as already stated, is to minimize uncertainty about the parameters.

Question 1 is evaluated by inspecting if there are differences between the algorithms in how quickly the aforementioned quantities approach zero. Question 2 is answered by looking at the same quantities, but the focus is on the overall performance, not differences. The first question is more closely related to the topic of this thesis, but the second question has more general value regarding the estimation of GRT models.

Also, the simulations were intended as providing preliminary information for the psychophysical experiment (Section 6 *Psychophysical experiments*) in which participants' task was to detect pitch and timbre differences. Thus although not mentioned beyond the choice of priors, the dimensions of stimuli in the simulation are implicitly pitch and timbre.

During the simulations the model with κ_μ (which will be labelled *Model 1* in Section 6 *Psychophysical experiments*) was used to determine the stimulus choices for two reasons. First, it's the most widely used model in the GRT related literature (or, the classic model that is analogous to it, see Section 2.4 *Critical discussion*: The models used e.g. in the 2x2 categorization experiments

commonly model interactions as interaction between means of the distributions. Second, there are some results suggesting timbre affects the pitch of signal (Allen & Oxenham, 2014; Platt & Racine, 1985), a phenomenon that would most naturally be modelled as shift in the mean of the evidence distribution; on the other hand, according to Silbert et al. (2009) the main source of interaction between pitch and timbre arises from correlated noise, which again is included in the model as the parameter ρ .

5.1 Methods

The general method was the following: first a set of generating parameters for the simulated observer was drawn randomly, and then either the adaptive or the non-adaptive algorithm was run (the simulated responses were generated during this part, since in the case of the adaptive algorithm, the chosen stimuli depend on the responses given, so the stimuli can't be generated in advance). This was done for both the Yes/No and 21-4AFC procedures, resulting in four different conditions. I have chosen the number 800 for the total number of trials fairly arbitrarily; it represents a number of trials that, I think, could still be administered relatively continuously to a participant, without taking into account non-stationarities induced by a multi-session design¹¹: as stated in Section 2.4 *Critical discussion*, these models assume that there are no shifts in, say, the response strategy or detection thresholds during the experiment. Longer experimental runs, or those that may span multiple days, almost certainly violate this assumption, and thus, I believe shorter experiments to—ironically—provide more accurate information.

Prior distributions Prior distributions and distributions from which generating parameters for the simulations were drawn from are shown in Figure 23 and in tabular form in Table 3. The same prior and generating distribution is used for both dimensions, that is the prior labelled σ is used for both of the σ parameters in the model. Note that the scale for criterion is given in false alarm probabilities for easier interpretation.

Priors for the parameters were chosen based on prior information from Dai and Micheyl (2011); Silbert et al. (2009) and from pilot testing. Priors for the σ parameters were chosen to on average reflect the expected sensitivity of the observers in this kind of task, with two caveats: 1) timbre changes were assumed to be divided by ten, that is a timbre change of 2 would correspond to a change of 20Hz on the original scale (see Section 6 *Psychophysical experiments* for more information about how the stimuli were defined) and 2) standard deviation for the prior was large to reflect the possibility of widely differing thresholds.

Priors for the interactions were centred on zero, acting as a way to keep inferences conservative.

Generating parameters for the simulations were drawn from bimodal distributions. The idea was to draw values that are covered by the prior, but which do not necessarily correspond with the mode of the prior distribution. Another motivation was to have qualitatively different simulated observers: some that have high values for some of the parameters and others that have low values.

¹¹800 was also the number of trials completed by both participants during one session of the psychophysical experiment conducted for this thesis, see Section 6 *Psychophysical experiments*

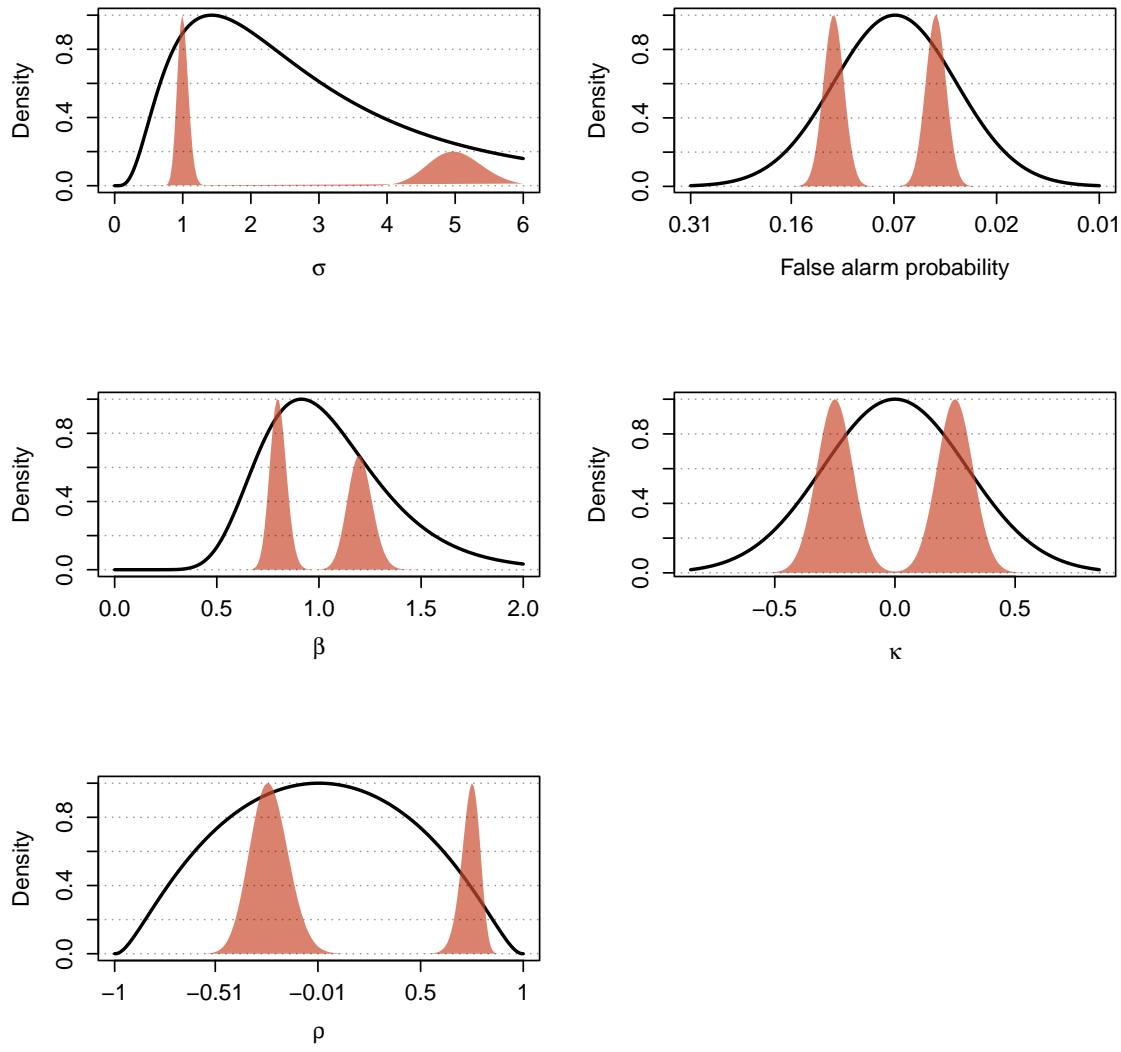


Figure 23: Prior distributions for the parameters of the model (solid black lines) and distributions for generating parameters (regions shaded with red). Note that all of the densities are normalized to have maximum value of 1.0.

Table 3: Parameters used for prior distributions and distributions of generating parameters. M^l and M^u , respectively, are for the lower and upper peaks of bimodal distributions.

	Prior		Generating		
	M	SD	M^l	M^u	SD
σ	$\log(2.5)$	0.75	$\log(1)$	$\log(5)$	0.08
C	1.5	0.3	1.2	1.7	0.05
β	$\log(1)$	0.3	$\log(0.8)$	$\log(1.2)$	0.05
κ	0	0.3	-0.25	0.25	0.075
ρ	0	0.7	$\text{atanh}(0.75)$	$\text{atanh}(-0.25)$	0.1

5.2 Results

Total number of simulations per condition are shown in Table 4. Squared errors and marginal standard deviations are presented in two ways: 1) on trial-by-trial basis and 2) by estimating the average differences on the last trial.

Table 4: Conditions and number of simulations in each.

Procedure	Algorithm	N
Yes/No	Adaptive	184
Yes/No	Random	284
2I-4AFC	Adaptive	174
2I-4AFC	Random	130

Trial-by-trial estimates Trial-by-trial results from the simulations are summarized in Figures 24 to 31. The figures show squared errors in relation to generating parameters and marginal standard deviations after N trials, all the way to trial number 800. In all plots black color is used for the randomly sampled stimuli while red is used for adaptively sampled stimuli. The shaded regions indicate 50%-quantiles (from 25% to 75%); solid lines indicate medians. Sensory (σ , β , crit) and interaction (κ_μ , ρ) are shown in their own figures.

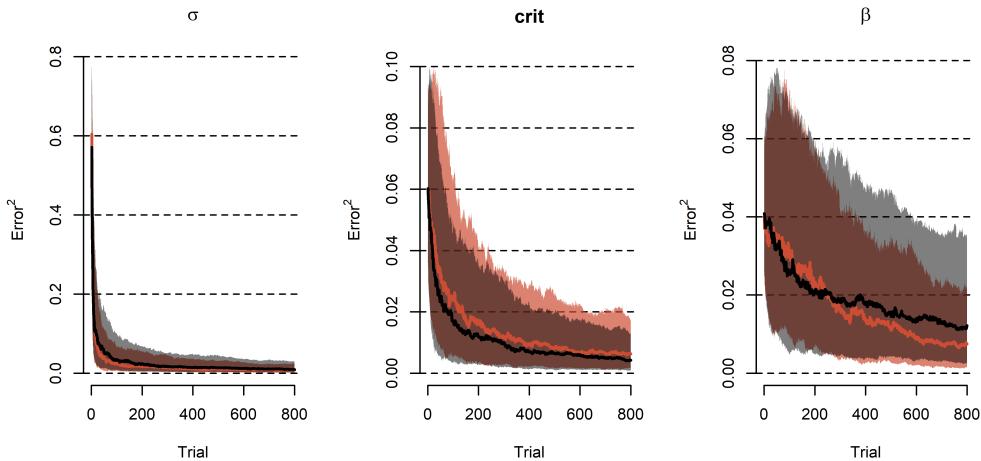


Figure 24: Procedure: Yes/No; sensory parameters. Trial-by-trial squared error between marginal means of the posterior distribution and generating parameters. Red: adaptive algorithm; black: random stimuli.

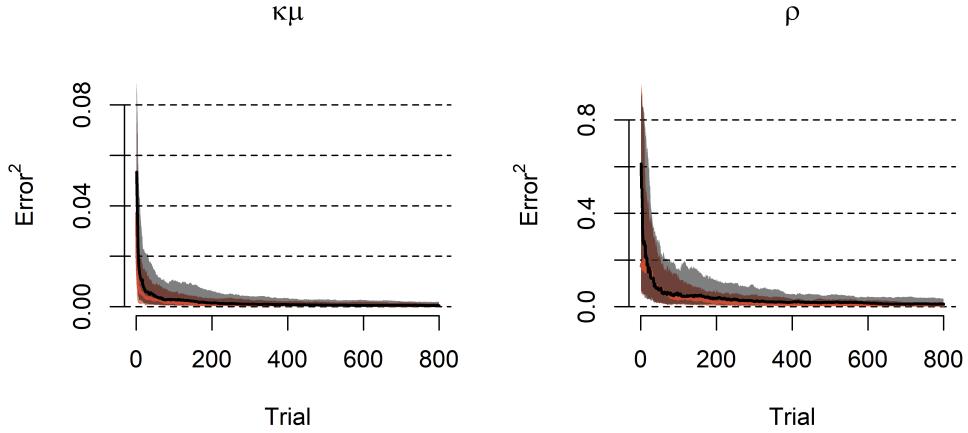


Figure 25: Procedure: Yes/No; interaction parameters. Trial-by-trial squared error between marginal means of the posterior distribution and generating parameters. Red: adaptive algorithm; black: random stimuli.

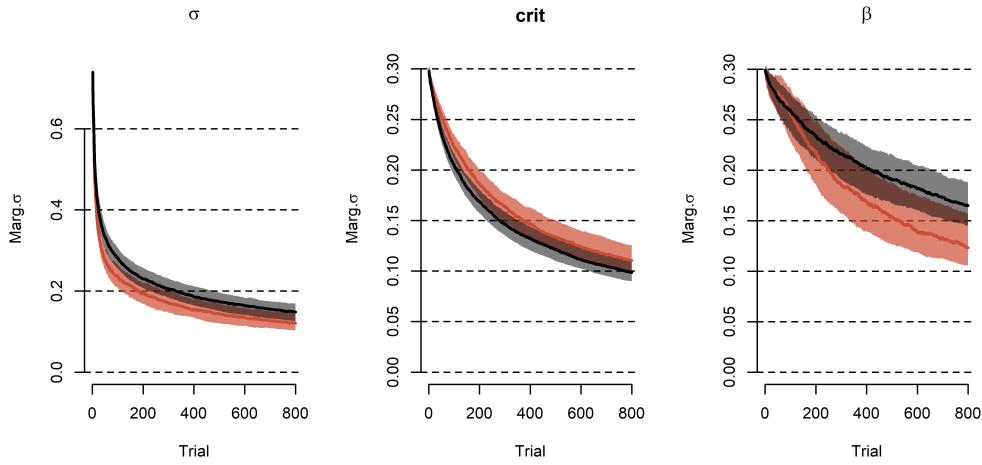


Figure 26: Procedure: Yes/No; sensory parameters. Trial-by-trial marginal standard deviations. Red: adaptive algorithm; black: random stimuli.

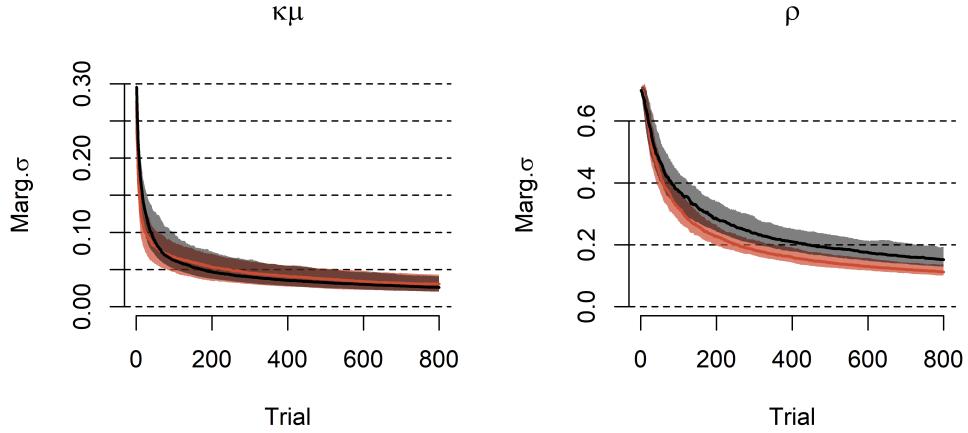


Figure 27: Procedure: Yes/No; interaction parameters. Trial-by-trial marginal standard deviations. Red: adaptive algorithm; black: random stimuli.

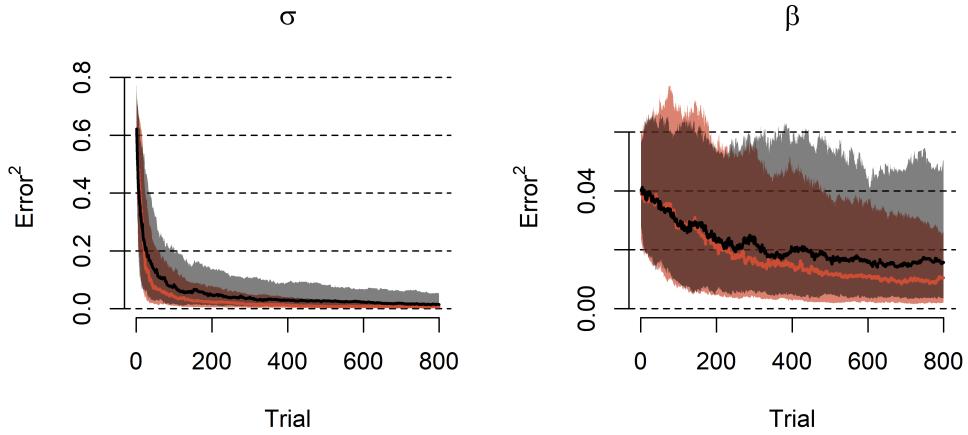


Figure 28: Procedure: 2I-4AFC; sensory parameters. Trial-by-trial squared error between marginal means of the posterior distribution and generating parameters. Red: adaptive algorithm; black: random stimuli.

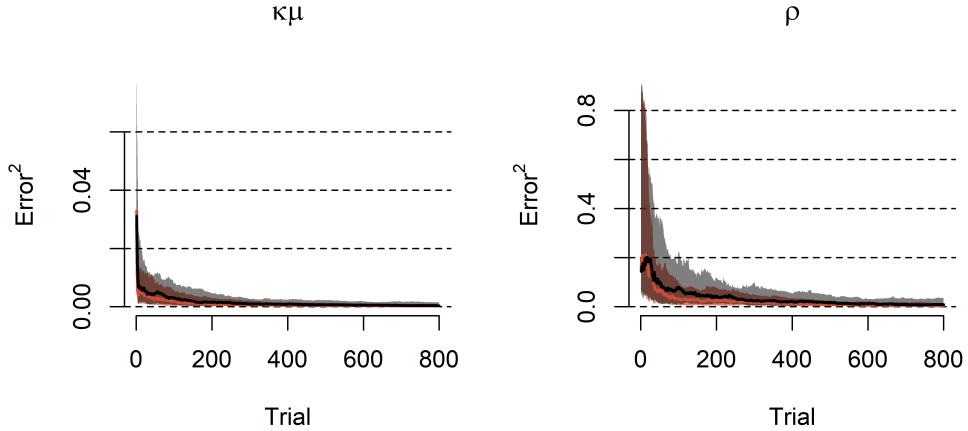


Figure 29: Procedure: 2I-4AFC; interaction parameters. Trial-by-trial squared error between marginal means of the posterior distribution and generating parameters. Red: adaptive algorithm; black: random stimuli.

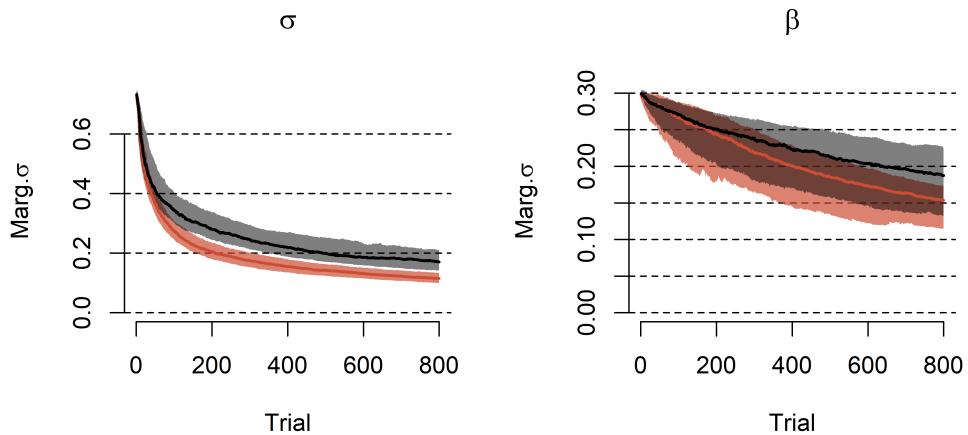


Figure 30: Procedure: 2I-4AFC; sensory parameters. Trial-by-trial marginal standard deviations. Red: adaptive algorithm; black: random stimuli.

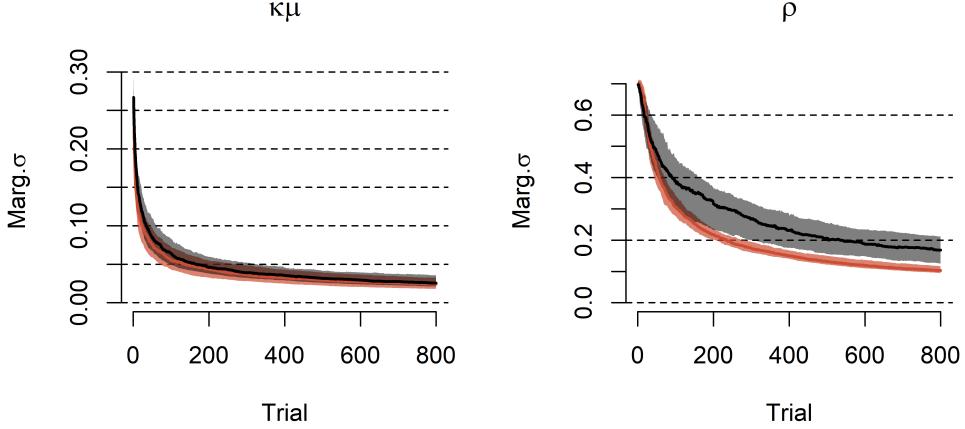


Figure 31: Procedure: 2I-4AFC; interaction parameters. Trial-by-trial marginal standard deviations. Red: adaptive algorithm; black: random stimuli.

Hierarchical model A hierarchical model was fit to both the marginal standard deviations and squared errors at the last trial to get a more quantitative understanding of the differences at that point. This should be considered as a first-order approximation of a more complete model of performance. Why only the estimates at the last trial are presented is because for any simple model which breaks the performance down in two factors, namely the final values and the speed at which they were arrived at, only one of the factors needs to be known. This is because all of the algorithms start from the same values (the prior distribution) and are run for the same amount of time (800 trials). Of course some more sophisticated model, that would perform more complex decomposition of performance could be used, but the development of such algorithm is beyond the scope of this thesis.

A common dummy coded linear model with Gaussian errors was used. The slopes, intercepts and standard deviations were pooled inside each condition, which means that e.g. all of the marginal standard deviations in the Yes/No condition with adaptively selected stimuli were given a common hyperprior. The models were fit using Stan.

Data was standardized before fitting ($y' = (y - \bar{y})/\text{SD}(y)$). Parameter-specific β' coefficients (slopes fit to standardized data) of the linear model are shown in Figures from 32 to 35. In all plots the thicker parts indicate 50% and narrower lines 95% equal-tailed intervals¹². Positive values indicate better performance by the adaptive algorithm.

¹²Equal-tailed interval (ETI) means the interval between some percentiles of the posterior distribution (Kruschke, 2015, p. 342), here for example the interval between 2.5% and 97.5% for the 95% ETI.

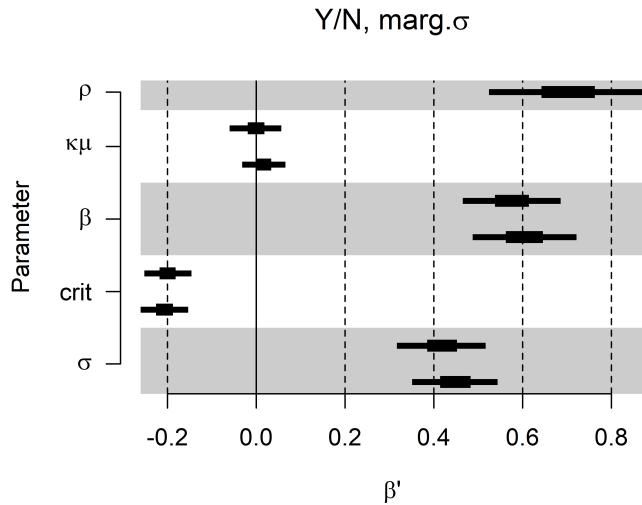


Figure 32: Procedure: Yes/No. β' coefficients for the difference between adaptive and random algorithms in marginal standard deviations.

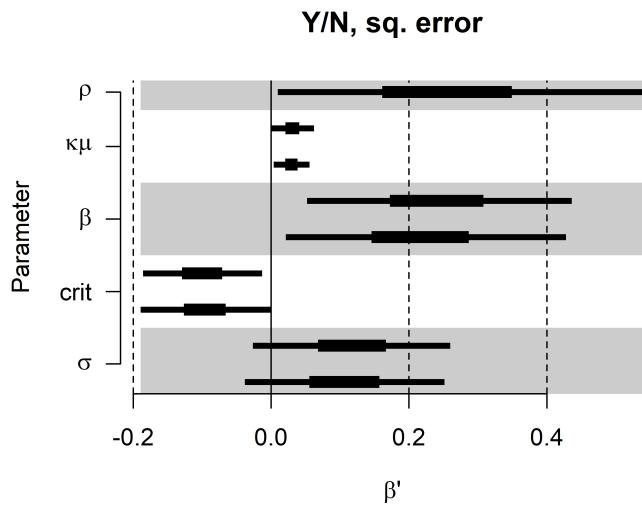


Figure 33: Procedure: Yes/No. β' coefficients for the difference between adaptive and random algorithms in squared errors between marginal means and generating parameters.

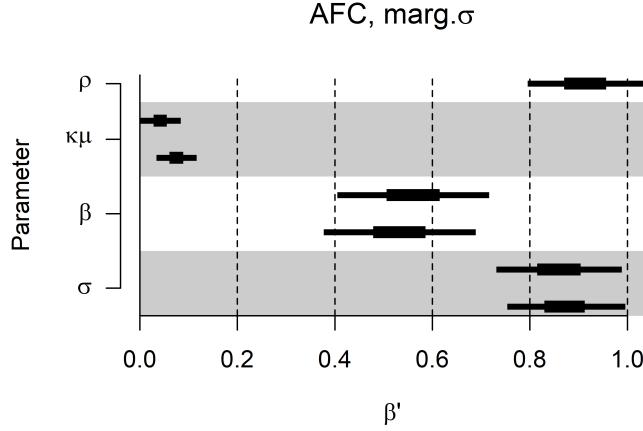


Figure 34: Procedure: 2I-4AFC. β' coefficients for the difference between adaptive and random algorithms in marginal standard deviations.

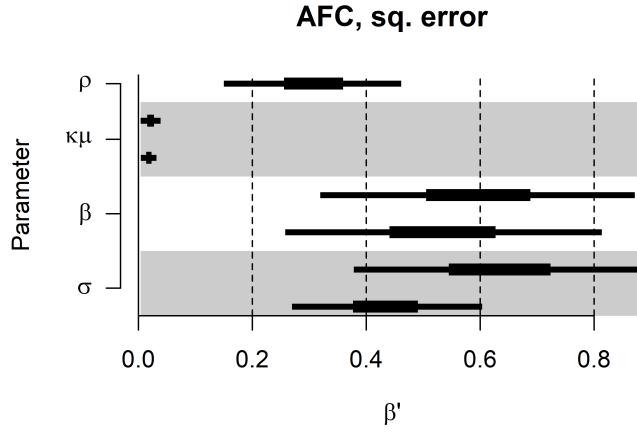


Figure 35: Procedure: 2I-4AFC. β' coefficients for the difference between adaptive and random algorithms in squared errors between marginal means and generating parameters.

5.3 Discussion

Question 1: is the adaptive algorithm more efficient? Judging from Figures 32 and 34, by the time of 800 completed trials the adaptive algorithm has managed to reduce marginal standard deviations more for parameters σ , β and ρ in both conditions. Effect sizes range from around 0.4 to almost 1.0. However, as can be seen from Figures 26, 27, 30 and 31, differences in raw scores are fairly modest, around 0.05 to 0.20.

There doesn't seem to be that much difference in κ_μ parameters, which is probably explained by the fact that, as can be seen from Figures 31 and 27 the marginal standard deviations for these parameters reduce fairly quickly.

The most surprising result is that random sampling seems to be more effective in reducing uncertainty about the criteria (Figure 32).

Similar pattern can be observed for the squared errors (Figures 33 and 35) but the differences are swamped by a lot more variability.

Question 2: how well are generating parameters recovered? It seems that the majority of improvement for most parameters happens before 400 trials. After that, information gain seems to slow down considerably.

The most problematic parameters would seem to be the β parameters. This contrasts results in for example Kontsevich and Tyler (1999), and could indicate that when in this more complex model inferences about the non-linearity of the d' function become more uncertain. Note also that the variance of squared error in estimating the β parameter increases somewhat during the first 100 trials (Figures 24 and 28, but the effect is more pronounced in the Yes/No condition), implying that for some simulations the posterior means drift further from the generating parameters.

With regard to the interaction terms, squared errors and marginal standard deviations for the κ_μ parameters seem to approach zero fairly quickly, most of the improvements seems to have happened well before 100 trials, but the ρ parameters seem more problematic (Figures 25, 27, 29 and 31). Marginal standard deviations for the ρ s don't get, on average, much smaller than .2 which still implies a lot of uncertainty—keeping in mind that correlation coefficients are bound between -1 and 1 .

6 Psychophysical experiments

In order to test the model in practice, I conducted a small scale psychophysical experiment ($N = 2$). Obviously, given the small sample size, the goal wasn't to generalize the results to a population, but to see if, even in this small preliminary sample, obvious problems would arise. The same model as in the simulations was used: the model with coupled means (κ_μ , which will be later called *Model 1*), reasons for this were detailed in Section 5 *Simulations*.

The stated goal can further be divided into two parts:

1. How well does the model behave?
2. Are there differences in parameter estimates from the two tasks (Yes/No and 2I-4AFC)?

What I mean by the model being *well-behaved* is that all of the parameters should be identifiable, and capture relevant features of the data. Furthermore, by the parameters being *identifiable* I mean that there shouldn't be strong posterior correlations between parameters, or the joint-posterior probability density should not be multimodal.

If there are strong posterior correlations between the parameters, this means that inferences about one parameter, to a degree, depend on the other. Multimodality can mean, for example in the case of the interaction parameters, that both positive and negative values would be supported by the data, which is obviously problematic.

Capturing relevant features is related to posterior predictive capabilities of the model: as a simple but artificial example consider a data set in which the conditional expectation follows second degree polynomial. If one would fit a linear model to that data set, and compare predicted data against observed data, it would be clear that the model doesn't capture the parabolic nature of the expected value. This question is slightly more complicated when it comes to the cognitive models in this thesis, but it will be discussed in greater detail in conjunction with the analyses from the psychophysical experiment.

The second question follows from the prediction that the only difference between the Yes/No and 2I-4AFC tasks should be that the sensitivity (parameter σ) in the 2I-4AFC task is reduced by $\sqrt{2}$. This prediction has been often used to test Signal Detection Theory models, see for example Wickens (2002). As such this question is also related to the validity of the model: deviations from this prediction would indicate that some of the assumptions are incorrect.

6.1 Methods

Procedure The experiments were programmed and run in R (R Core Team, 2019). Both participants completed 400 trials for both tasks during the same session. Before each trial participants were prompted to press enter to hear the stimulus, and feedback was provided after each trial.

The participants responded separately to both pitch and timbre: two inputs were required after each trial. In both tasks the first input always corresponded with pitch and the second with timbre. In the Yes/No task the participant would type 0 for *No* and 1 for *Yes*. For example, if they thought that pitch didn't change but timbre changed, they typed the string "01". In the 2I-4AFC task the participant would type the intervals in which they thought the dimensions changed. For example, if they thought pitch changed in the second interval and timbre in the first interval, they typed the string "21".

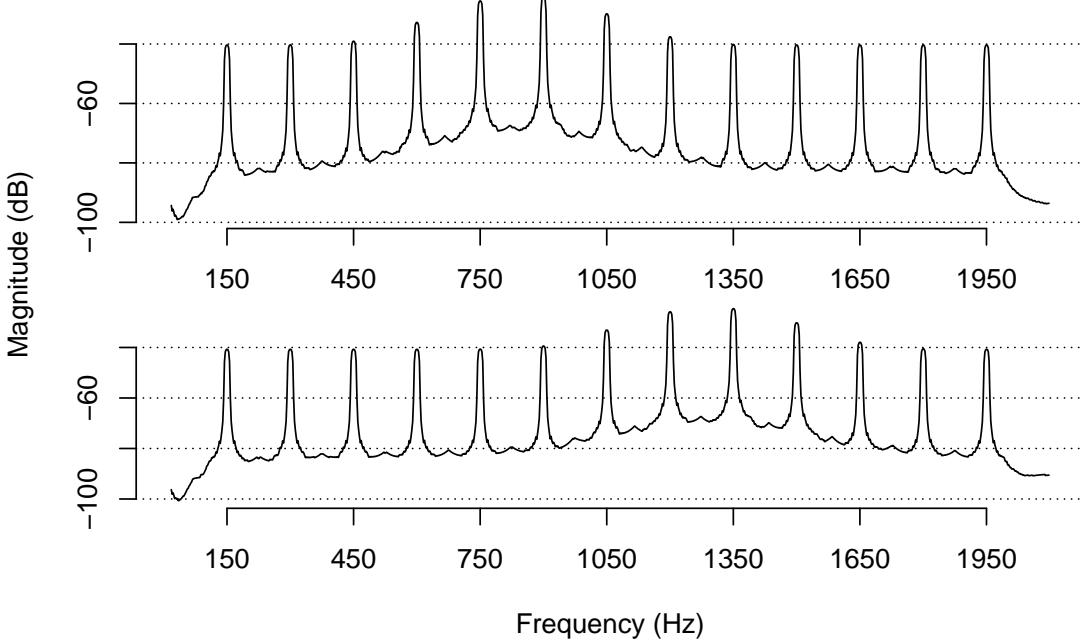


Figure 36: Examples of stimuli used in the psychophysical experiment. In both panels $F_1 = 150\text{Hz}$. Spectral prominence in upper panel is 850Hz and in lower panel 1300Hz, making the stimulus in the lower panel brighter in timbre.

Stimuli A single stimulus consisted of a reference tone and a test tone, as shown schematically in Figure 2.

Similar to Silbert et al. (2009) the individual tones consisted of 13 components that were integer multiples of the first component: pitch was manipulated by changing the frequency of the first component and timbre by filtering the harmonics with a Gaussian filter ($\sigma = 150\text{Hz}$).

Reference level (pitch of the first component) for pitch was 150Hz; reference level for spectral prominence (μ of the Gaussian filter) was 850Hz .

Both reference and test tones were 300ms in length, with a 100ms gap in between. In the 2I-4AFC task the intervals were divided by a 300ms gap.

6.2 Results

The initial model (model with only κ_μ) did not fit the experimental data very well. To alleviate this problem, I developed the alternative models, which did provide better fit to the observed data. All of the models that I will discuss in this section are presented in Table 5.

Model fits for both participants are presented separately. Furthermore, the analyses are divided into two parts: in the first part I describe fits for models from 1 to 3; in the second the model in which both kinds of interactions are present is discussed (this model is labelled *Both* in Table 5).

Table 5: All of the models fitted to the psychophysical data. If other information is not present, an X indicates that the parameter(s) are free in the model.

Task	Model	Parameters						
		σ	C	β	κ_μ	κ_σ	ρ	λ
Yes/No	Model 1	X	X	X	X		X	Fixed ($\lambda = 0.02$)
	Model 2	X	Wider prior	X	X		X	X
	Model 3	X	Wider prior	X		X	X	X
	Both	X	Wider prior	X	X	X	$-0.8 < \rho < 0.8$	X
	Both (Truncated)	X	Wider prior	X	X	$\kappa_\sigma > 0$	$-0.8 < \rho < 0.8$	X
2I-4AFC	Model 1	X		X	X		X	Fixed ($\lambda = 0.02$)
	Model 2	X		X	X		X	X
	Model 3	X		X		X	X	X
	Both	X		X	X	X	$-0.8 < \rho < 0.8$	X
	Both (Truncated)	X		X	X	$\kappa_\sigma > 0$	$-0.8 < \rho < 0.8$	X

After I've discussed the model development and model fits, I will discuss introspective remarks collected from the participants. This is intended as qualitative model criticism, in contrast to the quantitative model development that I will go through first.

Models one to three present the initial stage of model development: Model 1 is the same as the one used in simulations. For models 2 and 3 the lapsing rate parameter λ was estimated from the data, to accommodate the (presumably) higher proportion of lapsing trials than assumed a prior. Also the prior for the criterion was widened: the original prior turned out to be weight larger false alarm probabilities than what was intended. In Model 3 the variances of the evidence distributions were coupled, it uses the κ_{sigma} style model of interference.

The model labeled *Both* has, as the name implies, both kinds of couplings in it. The last model, *Both (Truncated)* is otherwise identical to the previous model, but the κ_{sigma} parameter is constrained to positive values. For these models also the parameter ρ was constrained between -0.8 and 0.8 to improve computational stability: it seemed that for the model with normal prior on tanh transformed scale the chains (of the MCMC sampler) sometimes get stuck to extreme values of correlation.

Stan code for all of the models are included in the appendix.

6.2.1 Posterior predictive plots

I did model checking by dividing the two dimensional space of stimuli into three bins per dimension, resulting into 9 total bins, and calculating the proportion of positive responses in each bin for both dimensions. In the data from the Yes/No task the first bin always included stimuli with strength 0.

This method resulted in unequal bin sizes and I would recommend that future works use more sophisticated methods for finding boundaries for the bins.

In all of the posterior predictive plots (see for example Figure 39) the black dots joined by the dashed lines indicate observed data. For example in the top left of Figure 39 the first three black dots indicate responses to increasingly large pitch changes pitch when timbre category is 1 (no change in timbre), the next three dots indicate response to pitch when timbre category is 2 (some change in timbre) and so on.

Table 6: Number of observations per bin in the posterior predictive distribution.

Bin	Pitch	1	1	1	2	2	2	3	3	3	
	Timbre	1	2	3	1	2	3	1	2	3	
N	JP	YN	108	58	49	53	38	6	50	6	32
		AFC	48	81	21	68	38	24	26	34	60
OK	YN	111	39	50	49	38	14	50	18	31	
	AFC	105	27	20	45	80	16	12	16	79	

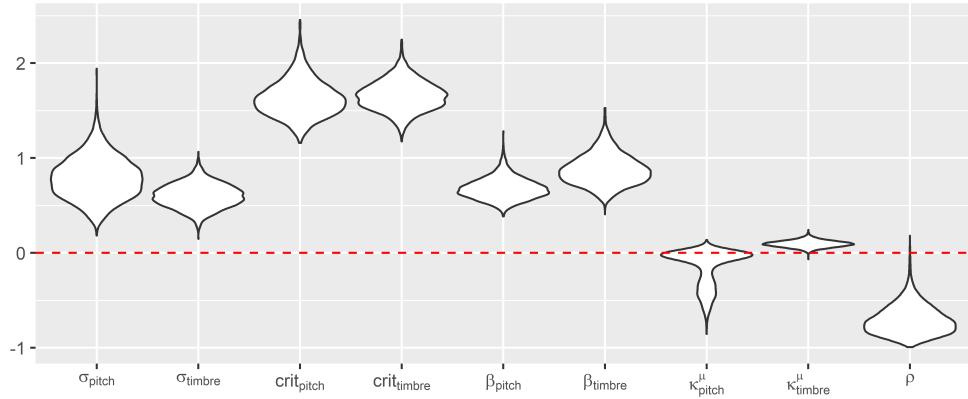
In this exemplar figure one can see that in each timbre category the probability of a positive response to a pitch change increases when the pitch differences get larger. However, the magnitude changes with the timbre category, indicating some kind of interference between pitch and timbre.

The data, in each case, is also divided by the category of the irrelevant dimension. For example on the upper row probability of a positive response to pitch dimension (relevant dimension) is plotted. These are broken down with regards to timbre (irrelevant dimension). Inside each timbre category, pitch category goes up from left to right.

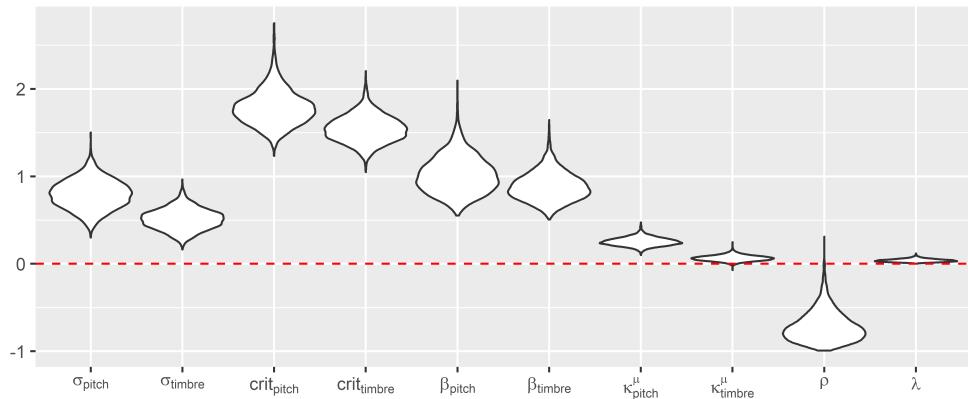
Since completely straight and orthogonal boundaries were used for binning the data, the number of observations per bin varied somewhat; for the participant JP, there didn't seem to be bounds that would result in better distribution. Another factor leading to uneven bin sizes was my decision to bin stimuli with strength 0 in the YN task separately. Number of observations per bin are reported in Table 6.

6.2.2 Participant OK, Models 1 to 3

OK/YN/Model 1



OK/YN/Model 2



OK/YN/Model 3

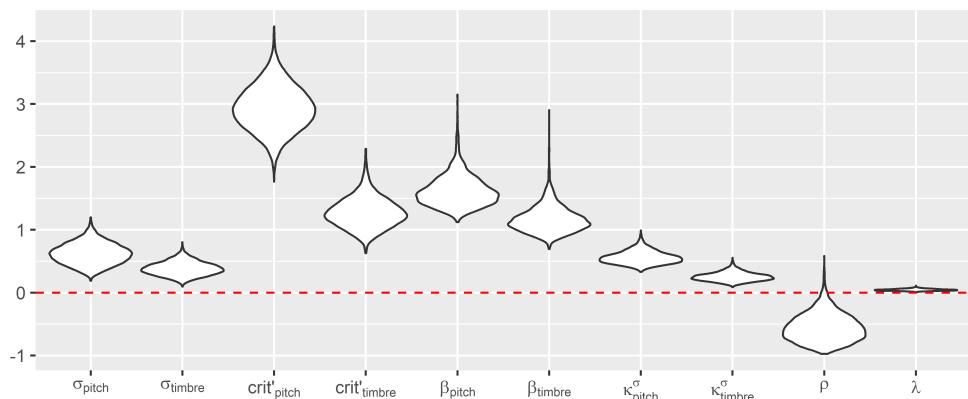
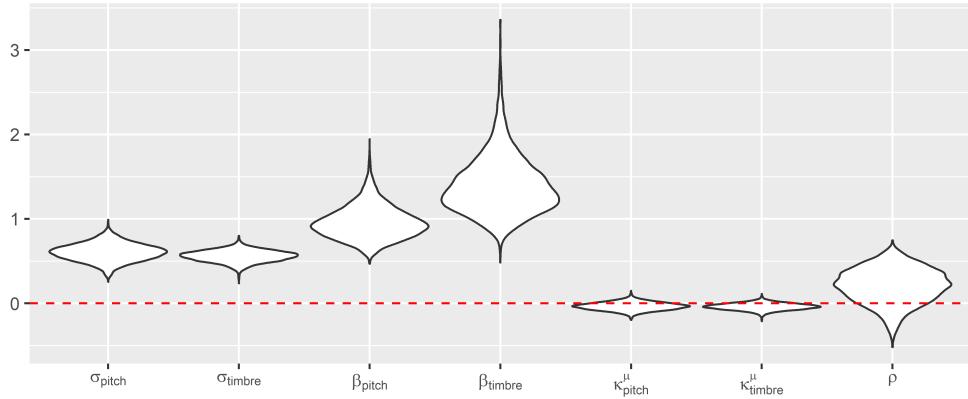
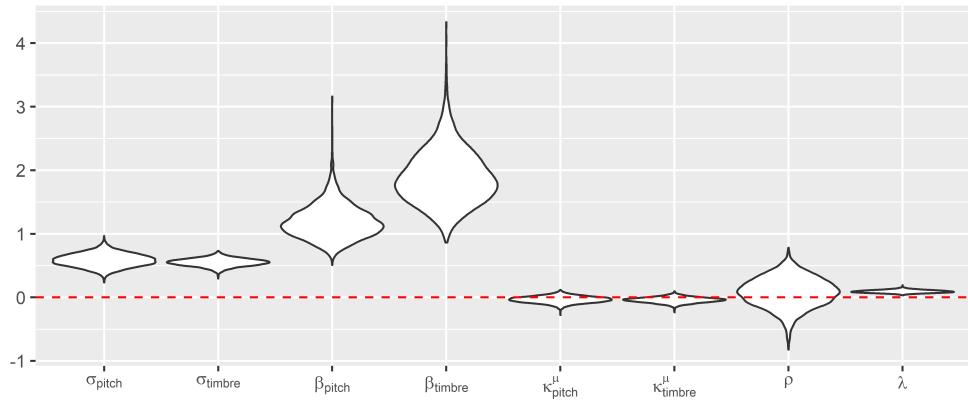


Figure 37: Task: Yes/No; Participant: OK. Marginal posterior distributions for parameters of Models 1 to 3.

OK/AFC/Model 1



OK/AFC/Model 2



OK/AFC/Model 3

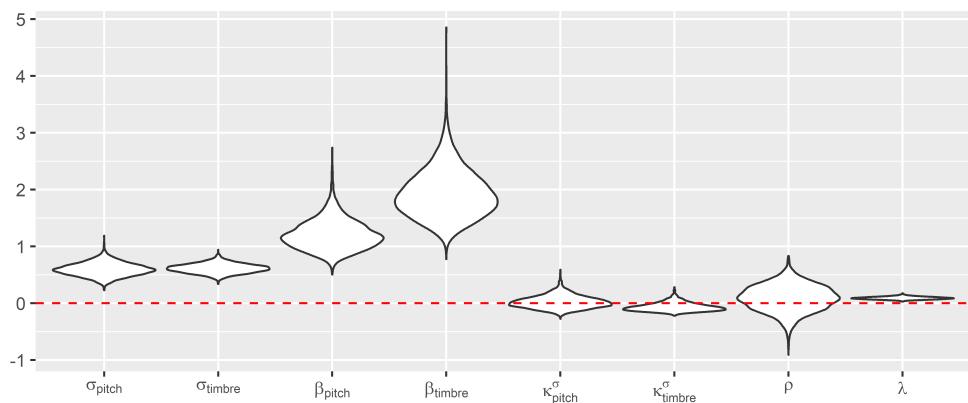


Figure 38: Task: 2I-4AFC; Participant: OK. Marginal posterior distributions for parameters of Models 1 to 3.

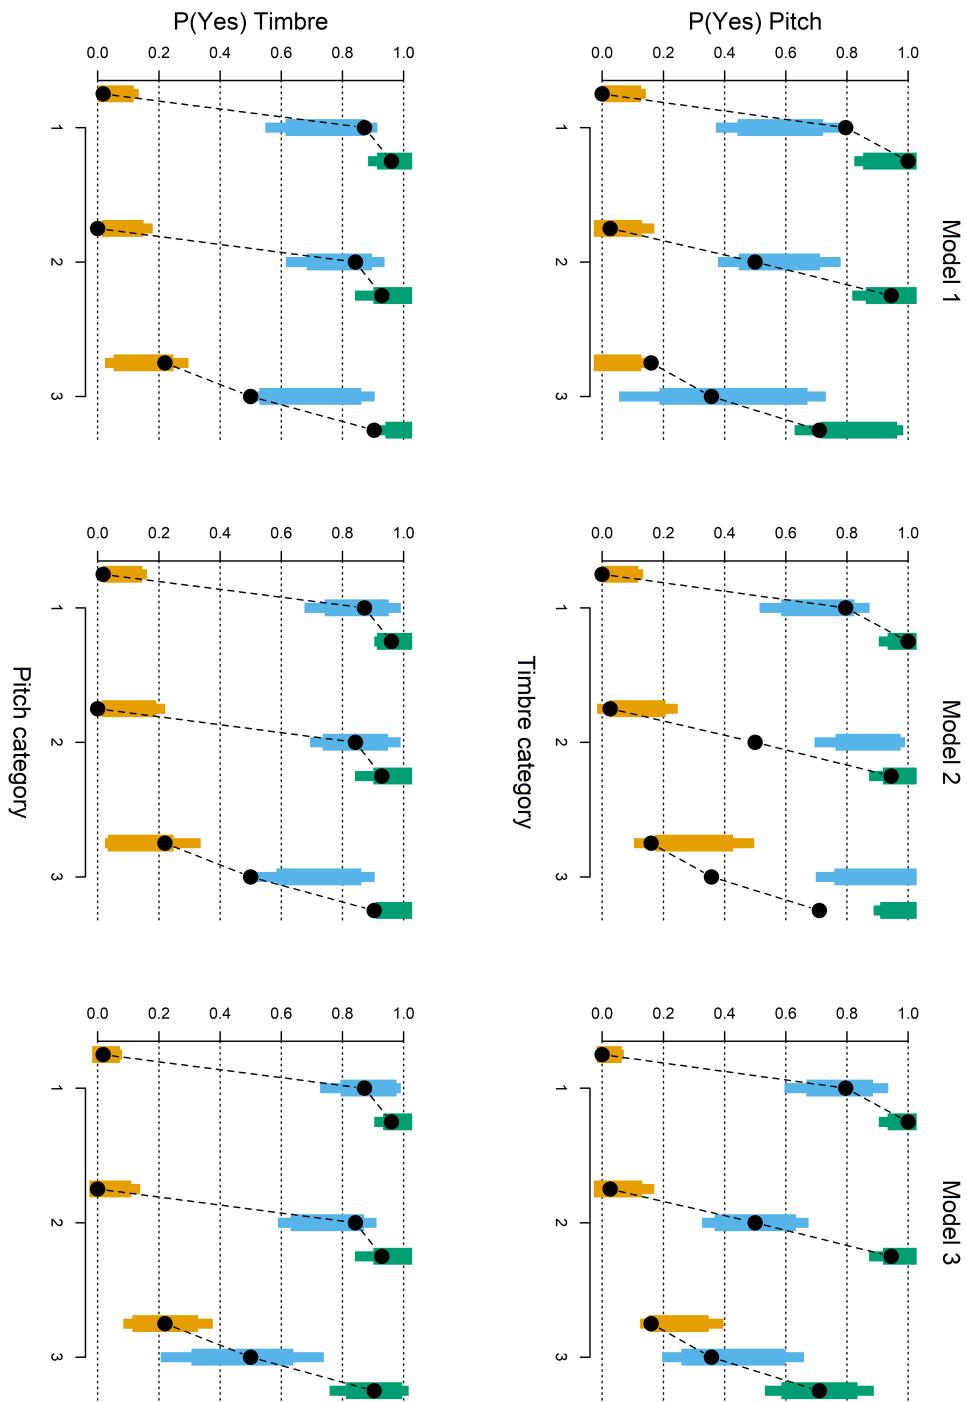


Figure 39: Task: Yes/No; Participant: OK. Posterior predictive distributions for parameters of Models 1 to 3.

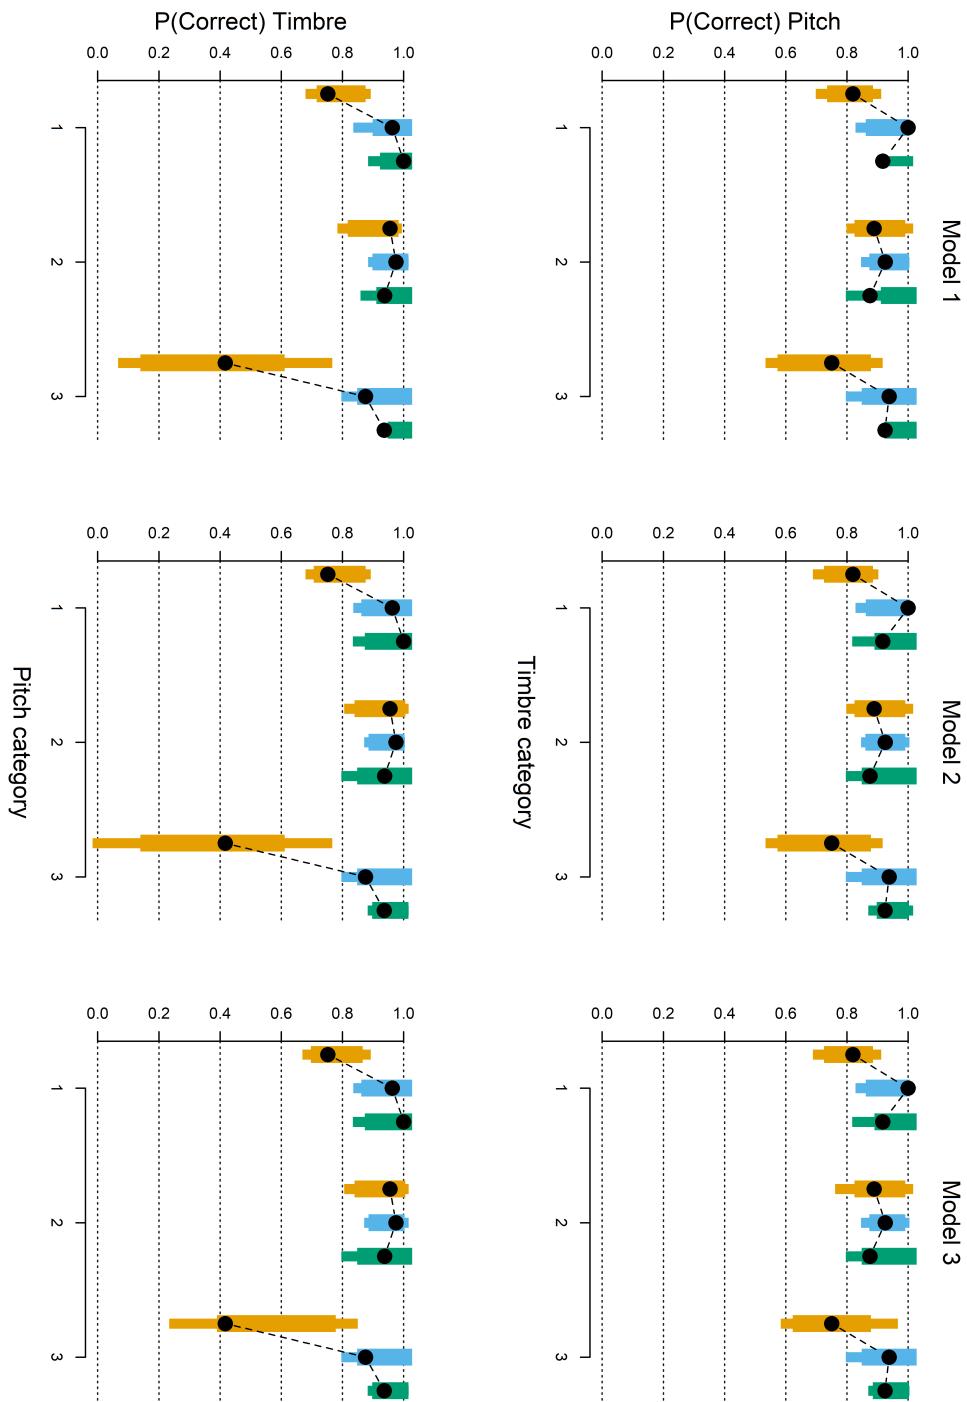


Figure 40: Task: 2I-4AFC; Participant: OK. Posterior predictive distributions for parameters of Models 1 to 3.

Participant OK, Models 1 to 3: Discussion In the Yes/No task, it is apparent from the bimodality of the posterior distribution for the κ_{mu} parameter (Figure 37, top panel) that there are problems with Model 1. Freeing the λ parameter (Model 2) manages to fix the bimodality for that parameter (centre panel in the same figure), but this affects the posterior predictive performance

negatively (Figure 39, centre panel).

From looking at the observed responses (Figure 39) it is clear that as the irrelevant signal level increases, the probability of a *Yes* response goes towards 0.50 which is more consistent with a shift in *standard deviation*.

Due to this, model with coupled standard deviations (Model 3) was fit to the data. It is clear from the posterior predictive plots (Figure 39, right panel) that this has the closest match to the observed data.

The prior used for the criterion proved to be too tight: the data set provided less information about the false alarm probabilities than I expected *a priori*. This lead to poor posterior predictive performance. To alleviate this problem I widened the prior for the criterion for all of the models besides Model 1, which is the original model.

The same set of models was fit to the data from the 2I-4AFC task. However there is much less to be said; none of the models seem to do significantly worse than the others (Figures 38 and 40).

6.2.3 Participant OK, Model with both interactions

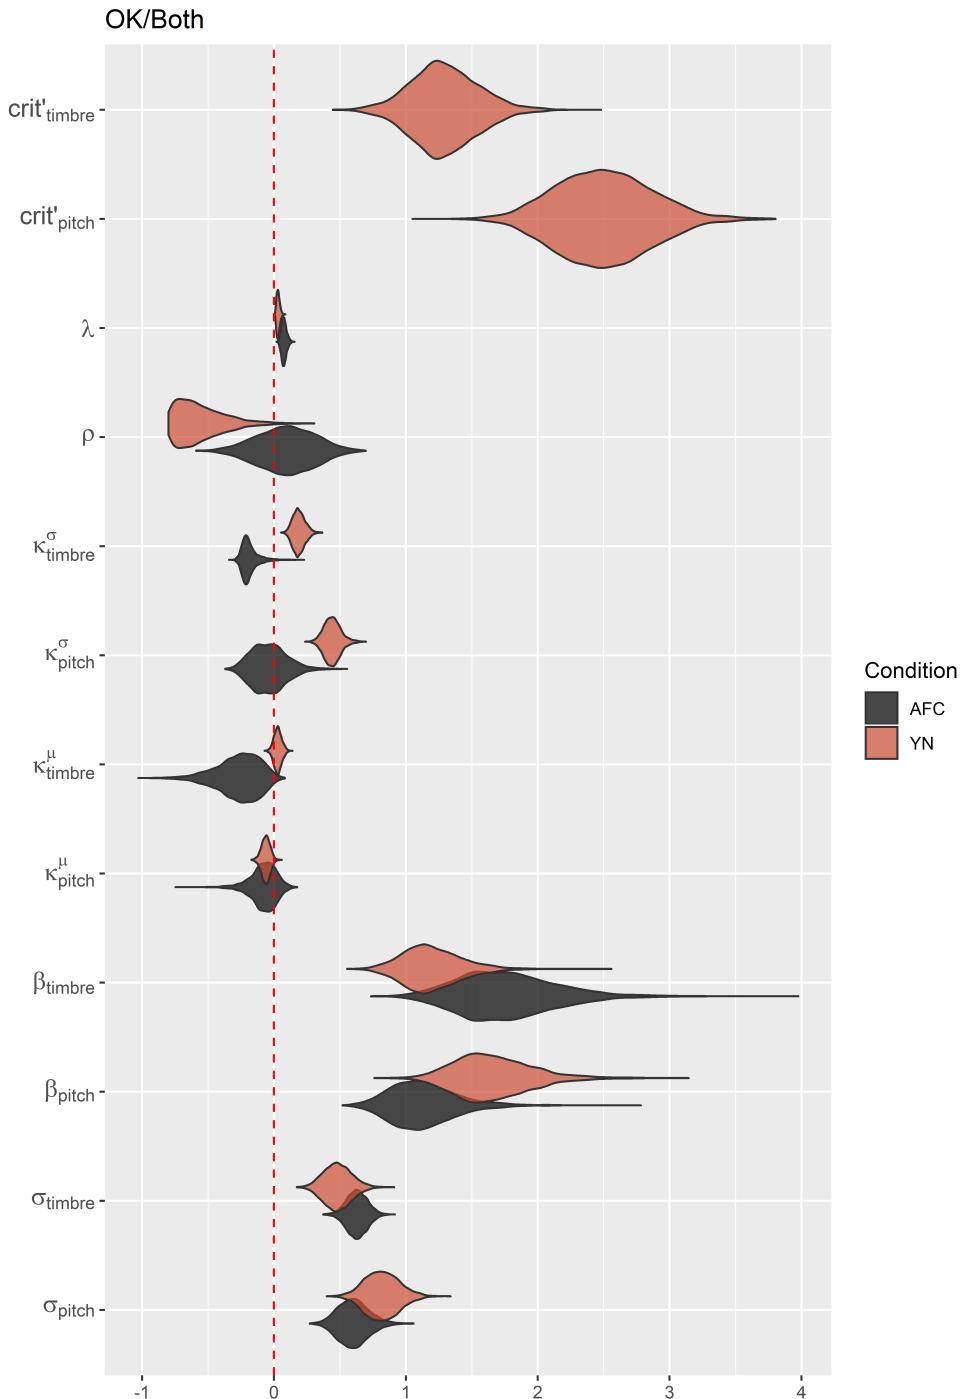


Figure 41: Both tasks; Participant: OK. Marginal posterior distributions for parameters of model with both interactions

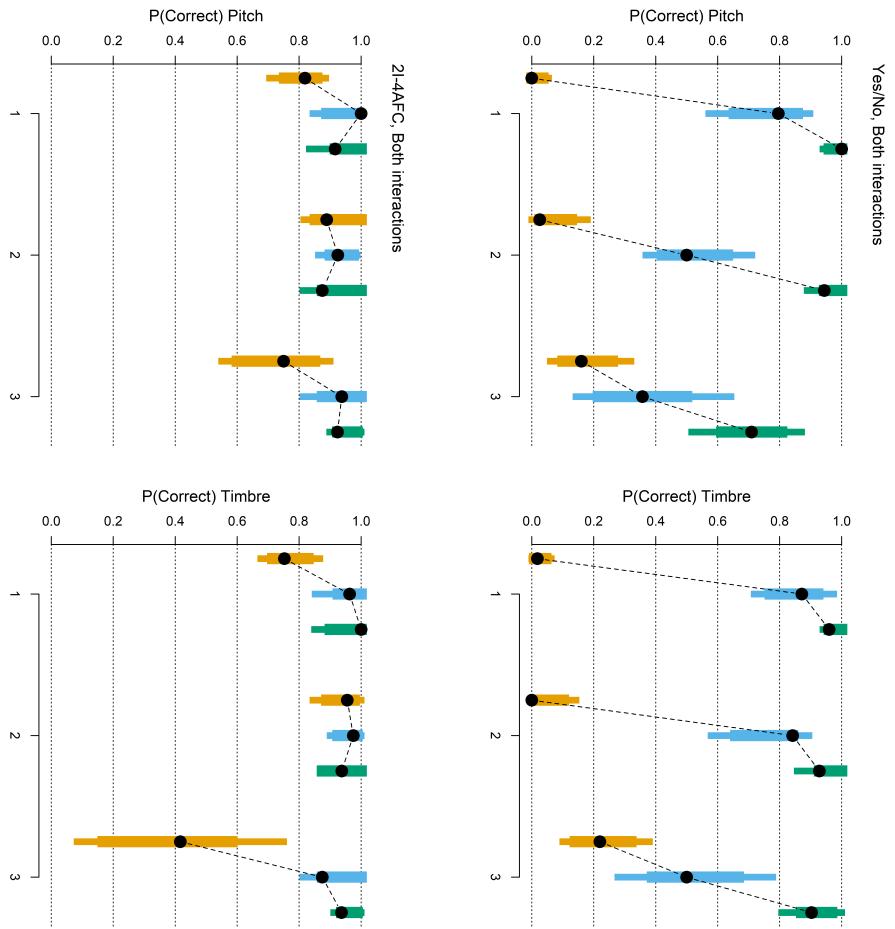
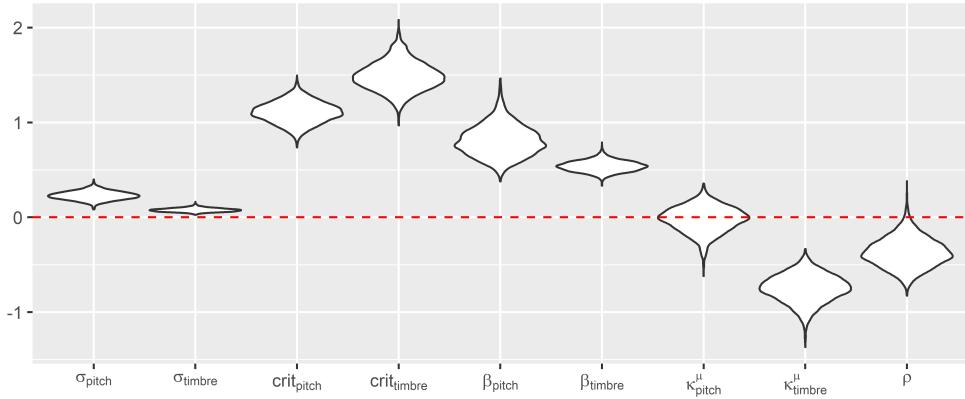


Figure 42: Both tasks; Participant: OK. Marginal posterior distributions for parameters of model with both interactions

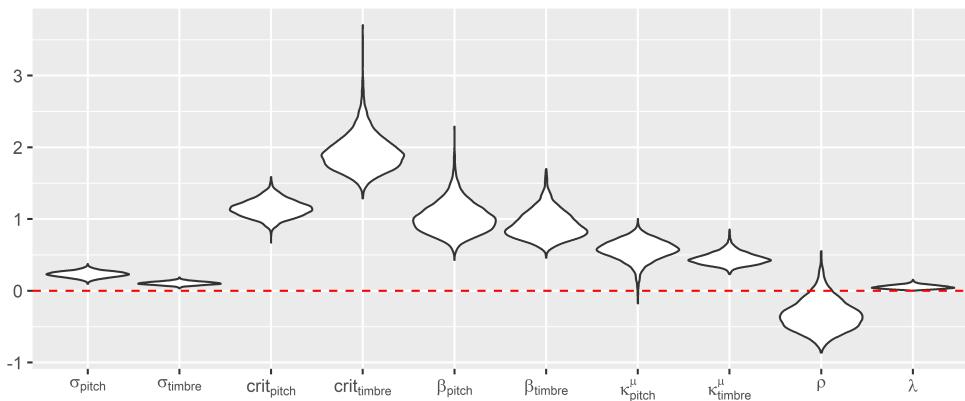
Participant OK, Model with both interactions: Discussion Both kinds of interactions were combined to a single model, in order to compare the relative contributions of both interactions to the observed patterns of inference. Also, one other modification was made: the prior for ρ was changed to a uniform distribution ($\rho \sim \text{Uniform}(-0.8, 0.8)$) since when running the Stan programs, a chain, if it wandered too close to extreme values, would get stuck to those extreme values; the motivation in this modification, then, was to improve computational stability of the model.

6.2.4 Participant JP, Models 1 to 3

JP/YN/Model 1



JP/YN/Model 2



JP/YN/Model 3

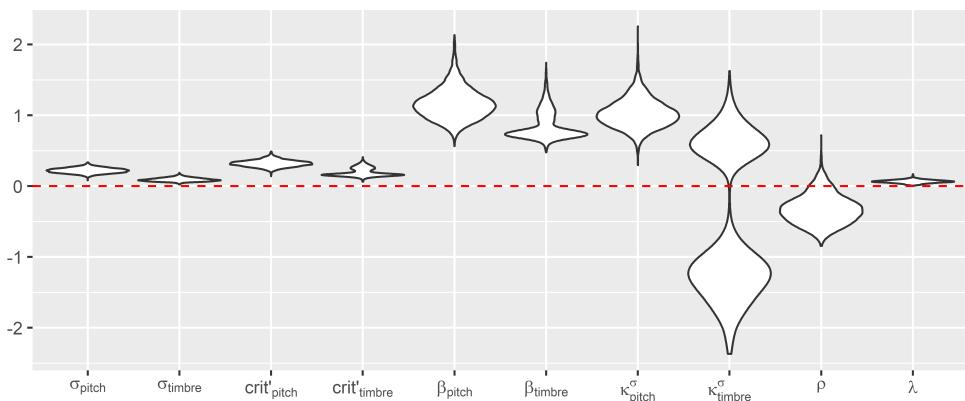
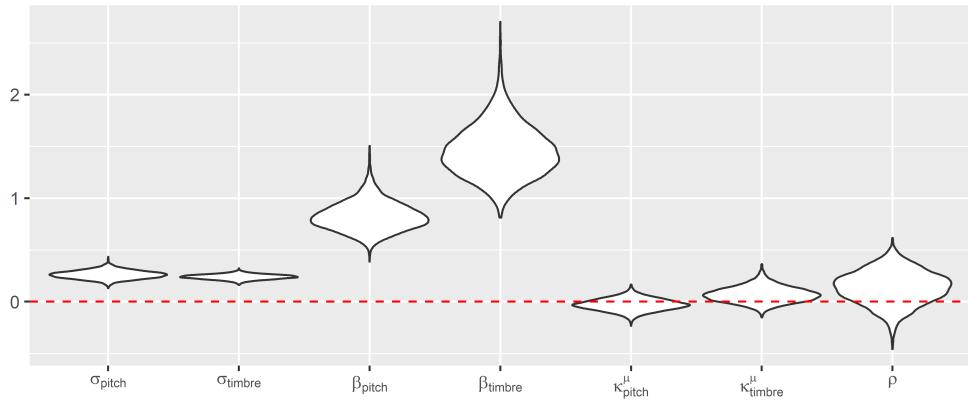
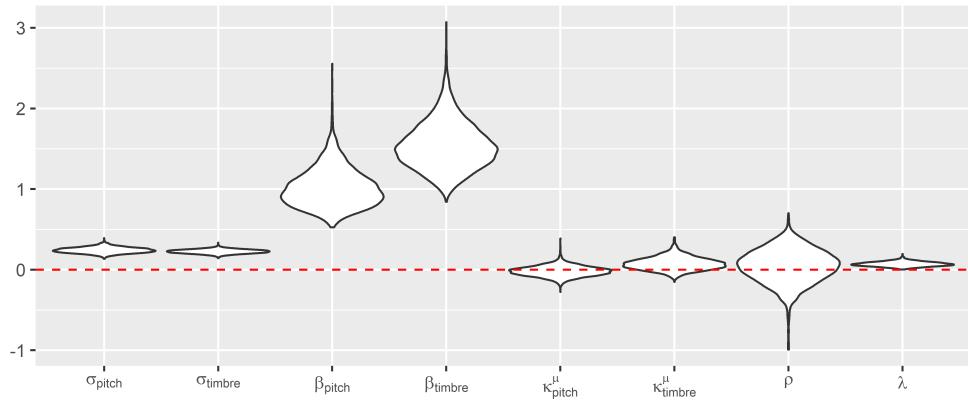


Figure 43: Task: Yes/No; Participant: JP. Marginal posterior distributions for parameters of Models 1 to 3.

JP/AFC/Model 1



JP/AFC/Model 2



JP/AFC/Model 3

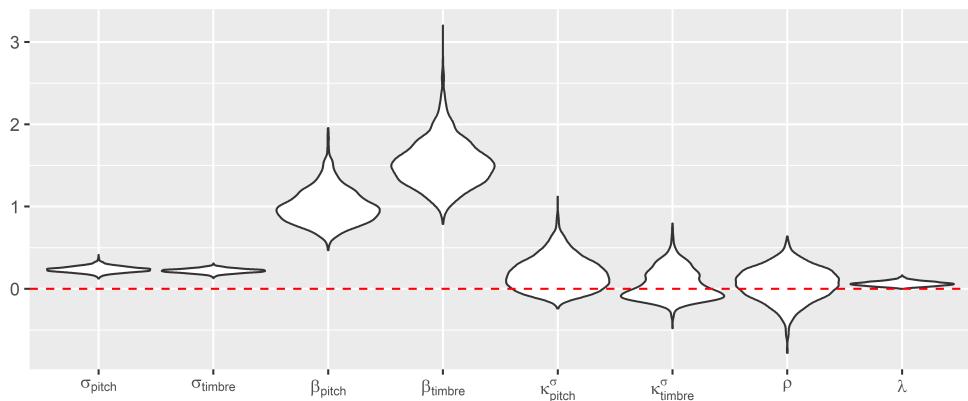


Figure 44: Task: 2I-4AFC; Participant: JP. Marginal posterior distributions for parameters of Models 1 to 3.

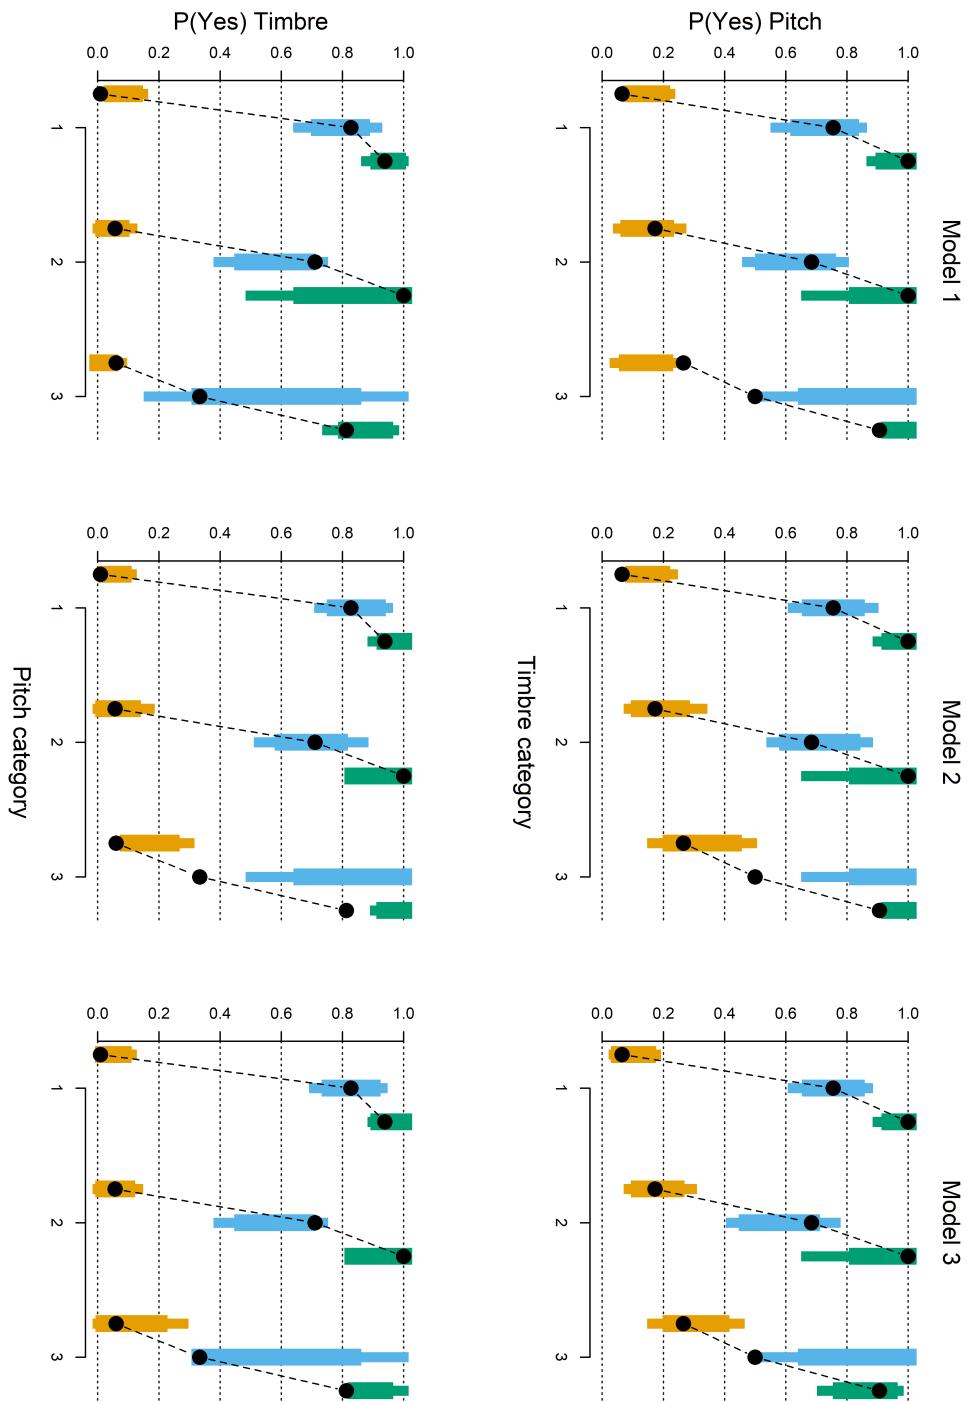


Figure 45: Task: Yes/No; Participant: JP. Posterior predictive distributions for parameters of Models 1 to 3.

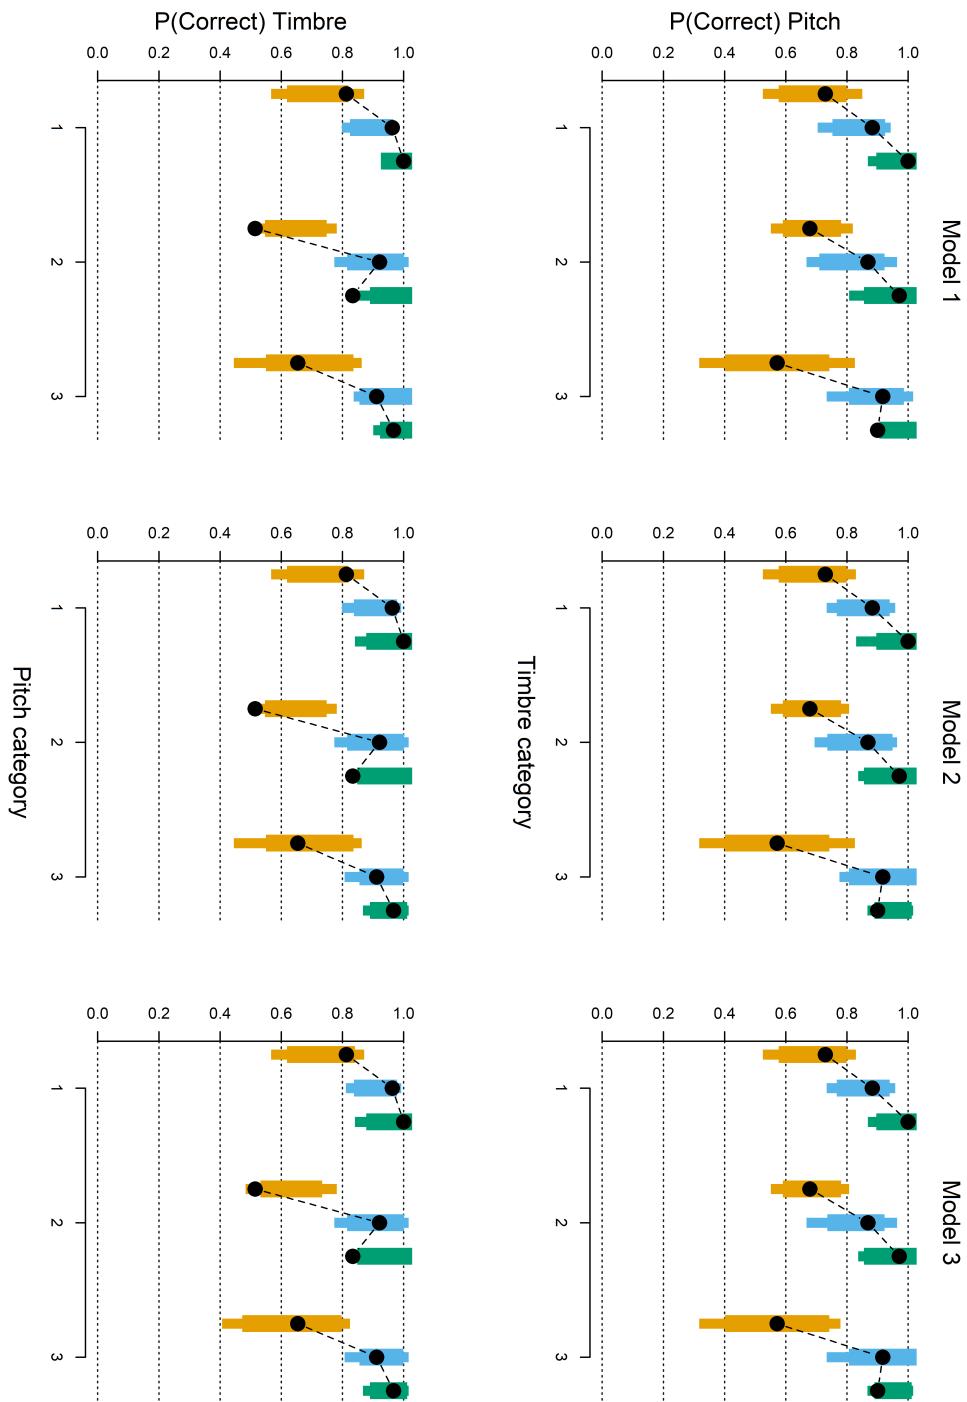


Figure 46: Task: 2I-4AFC; Participant: JP. Posterior predictive distributions for parameters of Models 1 to 3.

Participant JP, Models 1 to 3: Discussion Yes/No: In contrast to participant OK, here Model 3 (Figure 43) seems to suffer from non-identifiabilities: sign for the κ_σ parameter for *timbre* dimension seems to not be identifiable from the data. The bimodality from the κ_σ parameters seems to have bled to the criterion parameter on the *timbre* dimension.

6.2.5 Participant JP, Model with both interactions

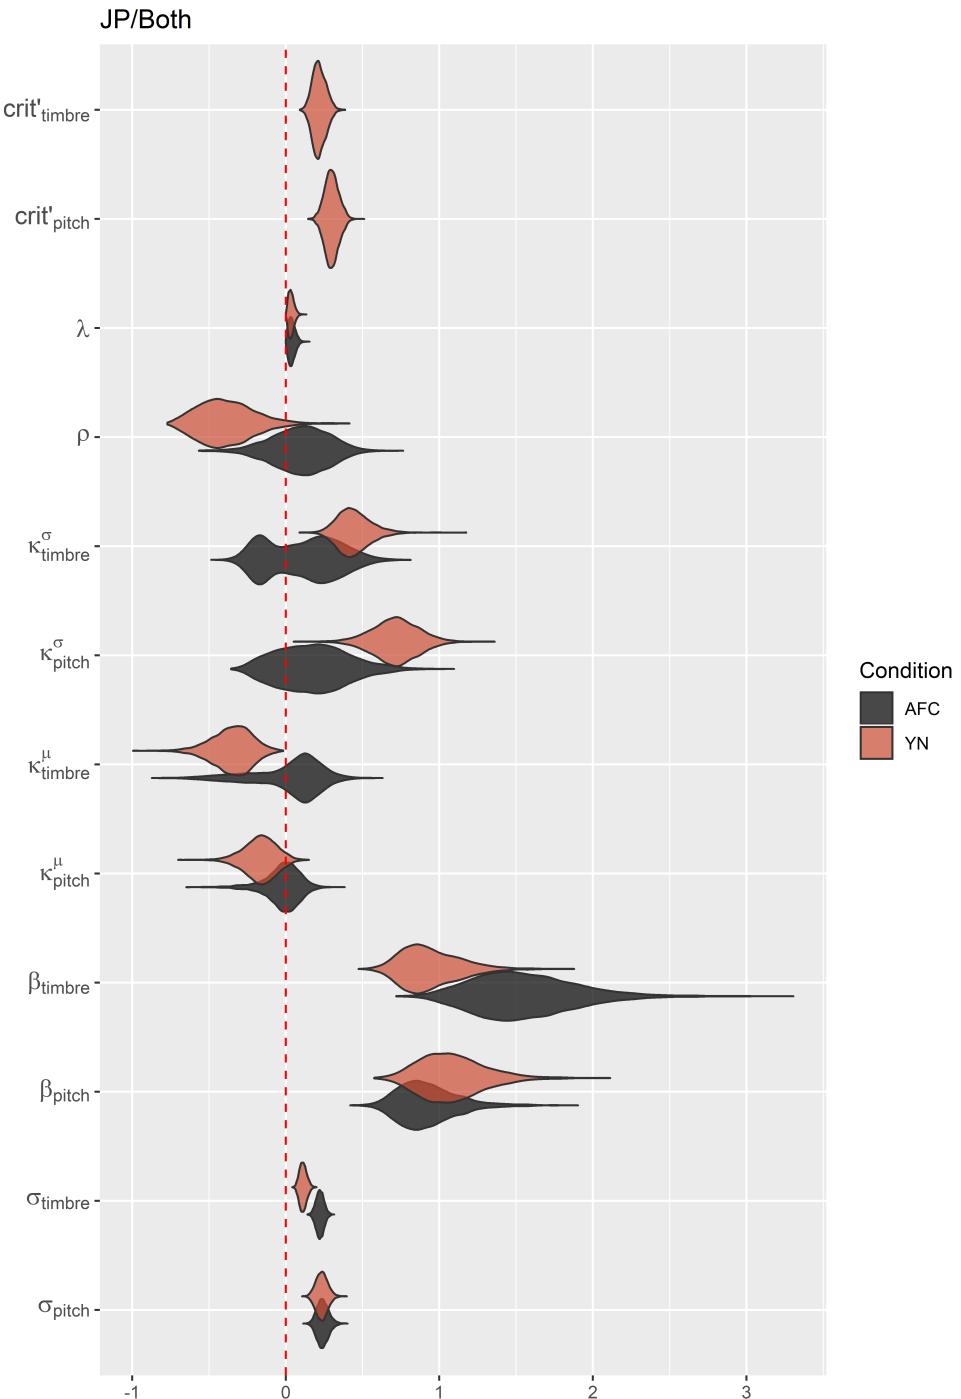


Figure 47: Both tasks; Participant: JP. Marginal posterior distributions for parameters of model with both interactions

Participant JP, Model with both interactions: Discussion As can be seen from the parameter estimates (Figure 47) the sign of κ_σ in the Yes/No task is not identifiable here either. To alleviate this problem, I constrained κ_{sigma} to positive values. This seemed reasonable, since the parameter was identified as positive in the 2I-4AFC model, and on the grounds that pitch and timbre are usually thought to "draw" the signal level in the same direction, that is, higher timbre results in slightly higher pitch estimates (Platt & Racine, 1985).

6.2.6 Participant JP, Model with both interactions, $\kappa_\sigma > 0$

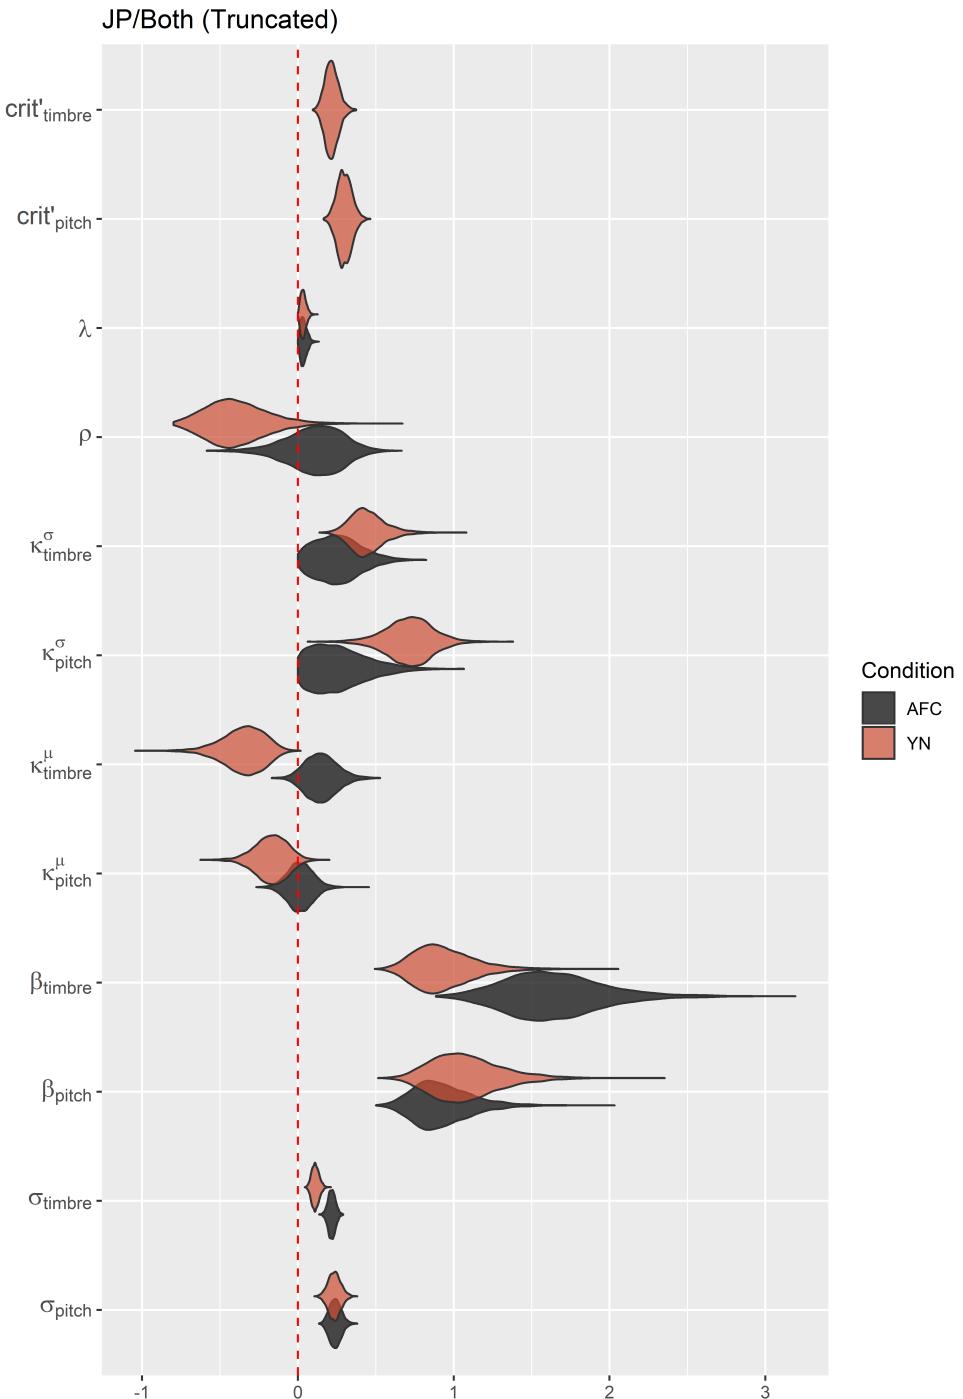


Figure 48: Both tasks; Participant: JP. Marginal posterior distributions for parameters of model with both interactions, and $\kappa_\sigma > 0$

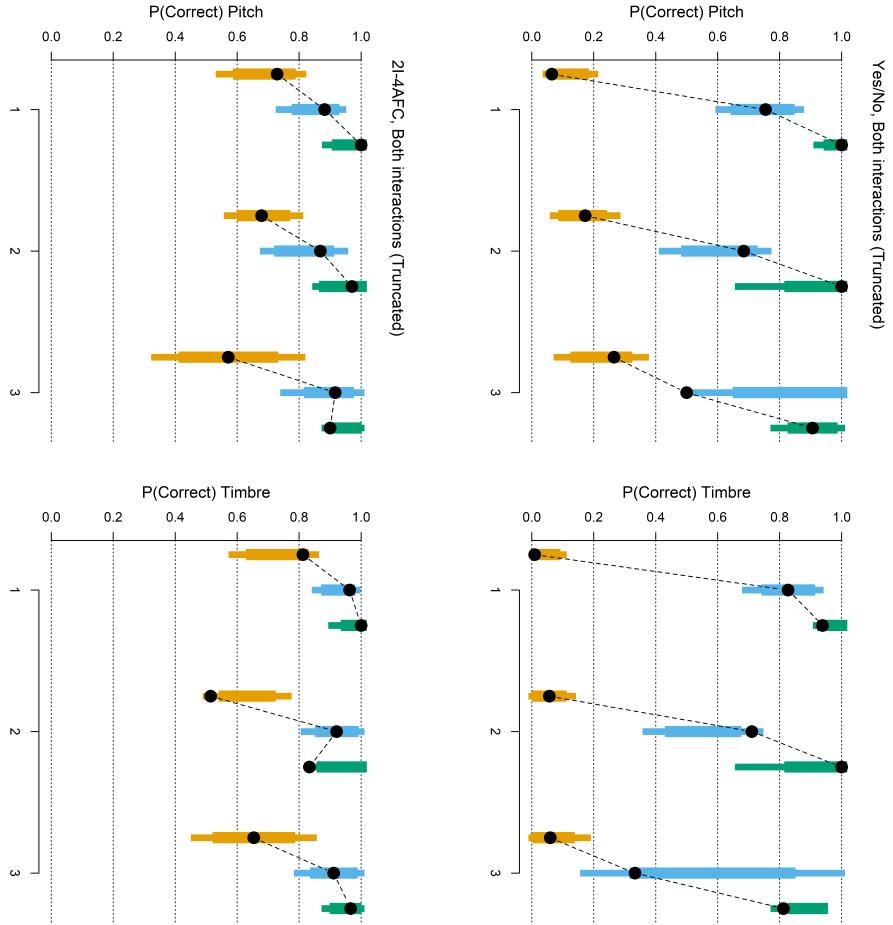


Figure 49: Both tasks; Participant: JP. Marginal posterior distributions for parameters of model with both interactions, $\kappa_\sigma > 0$.

Participant JP, Model with both interactions, $\kappa_\sigma > 0$: Discussion The posterior distributions don't exhibit obvious problems with the model (Figure 48) and the posterior predictive distribution (Figure 49) indicates acceptable fit with the observed data.

6.3 Some introspective remarks

While posterior predictive checking acts as a quantitative way to check (some) aspects of the models, I will be also incorporating more qualitative model critique in the form of introspective remarks. This is done in the spirit of Newell's adage that "*Know the method your subject is using to perform the experimental task*" (Newell, 1973). Of course, this kind of introspection was already present in the earlier discussion about the pitfalls of an identification task (vs. discrimination).

I've compiled and edited some of them that I feel are most important regarding violations of the assumptions of the model in this section.

I've chosen not to include the information about from which participant each remark is for three reasons: first, these are highly subjective and contingent on how each participant verbalizes their internal processes; second, some of the remarks can be conflicting, reflecting dynamic changes

strategies or attention; and third, since there were only two participants, no inferences about the generalizability can be made anyway. However, these provide important information about the limitations of the models and tasks as implemented here; which ones are important and how they are to be remedied are questions left for future work.

Remark 1. *The response input style used for the 2I-AFC task here was difficult, and potentially lead to increased amount of lapses; at least it meant that the cognitive burden was greater in that task.*

Remark 2. *In the 2I-4AFC task the overall sounds of the intervals were compared against each other more so than stimuli inside the intervals.*

Remark 3. *Attention and decisional criteria were calibrated by immediately preceding stimuli. For example the internal idea of how stimulus in which there's change on both dimension sounds like was affected by preceding stimuli and responses to them.*

Remark 4. *Stimuli were attended sequentially: one of the dimensions was primary, and the decision whether it changed/in which interval it changed was made first.*

Remark 5. *In the 2I-4AFC task easily discriminable stimuli terminated attention to that dimension. E.g. if there was a clear change in timbre during first interval, there was no need to listen to it on the second interval; instead, all attention could be diverted to listening for changes in pitch.*

Remark 6. *In the 2I-4AFC task if one of the dimensions was not discriminable and the other was, often the default choice was to pick the interval in which the more discriminable stimulus was, thinking that maybe the large stimulus masked the fainter.*

Remark 7. *There was a lot of variability in how the response categories were used during the experiment.*

6.4 General discussion

Question 1: Was the original model (Model 1) sufficient? For participant OK Model 1 resulted in posterior distribution with two modes (Figure 37), which is a clear sign of problems with the model. The marginal posterior distributions for participant JP are less problematic, but it is clear from the posterior predictive plots for both participants (Figures 39 and ??) that Model 1 fails to capture the fact that for both participants changes in the other dimension increase false alarm probability while decreasing the probability of a hit. A pattern that is only explained by including the term κ_σ to the model.

For the 2I-4AFC models there aren't as great differences.

The pattern of results in which the Yes/No model has more problems probably is due to that model being more diagnostic.

Question 2: Did coefficients in the Yes/No and 2I-4AFC tasks correspond with each other? Due to problems with Models 1 to 3, I used models with both kinds of interactions (and with $\kappa_\sigma > 0$ for participant JP) to test the prediction that the coefficients from both tasks should be identical (when taking into account the reduced sensitivity in the 2I-4AFC task).

Results are summarized in Figures 41 and 48. For both participants it seems that the sensory parameters (σ and β) are fairly close to each other in both tasks. However, the interactions are more varied.

Parameter κ_μ seems to be the most well-behaved in this sense. For participant OK the κ_μ parameters suggest—strongly for the *timbre* dimension—different signs for the coefficient; for participant JP these coefficients for the AFC model seem to lean towards zero, while for the YN model positive values are suggested.

For both participants the ρ parameters are drawn towards zero in the AFC model while strongly suggesting negative values in the YN model.

Question 3: What inferences can we make about the processing of pitch and timbre?

As already discussed, the interaction coefficients varied significantly across the two models, which suggest that the coefficients are strongly influenced by e.g. attention and decisional processes. This makes it hard to draw strong inferences.

Discussion will again be based on Figures 41 and 48. For both participants the κ_σ coefficients are positive for both *pitch* and *timbre*. This suggest that changes on the other dimension increase the variability in the evidence distributions.

For participant JP both κ_μ coefficients are negative, suggesting that changes on the other dimension decrease the mean of the evidence on the other dimension, that is for example large timbral changes make pitch changes seem smaller (assuming that the evidence distribution reflects mainly perceptual processing).

In the AFC task $\kappa_{\text{timbre}}^\mu$ is mostly on the positive side, which would suggest that changes in pitch would make timbral changes seem larger; or that the participant becomes more prone to answering positively to *timbre* dimension.

Conclusions from introspection It is clear from the introspective remarks that many of the (im-plicit) assumptions of the models were violated. The remarks indicate for example non-stationarity, serial processing of the dimensions and correlation between subsequent stimuli. Especially the 2I-4AFC task proved to be problematic due to the more complicated response scheme and complex stimuli. To better diagnose these—and other—problems, more methods for posterior predictive checking should be developed, to increase the coverage of model criticism.

7 Conclusions

The main theoretical contribution of this thesis is to add psychometric functions to GRT and define interactions through the coupling of these functions. Primary reason for this was to model response probabilities as a function of continuous stimulus values; the classic GRT models are used for categorical stimuli. This adds some structural assumptions to the model, such as how to parameterize the relationship between physical signals and d' values and how to couple the psychometric functions. The additional structural constraints help make more focused predictions, that—when they violate observed data—can be used to diagnose problems with the models, which in turn is important for model development. The classic categorical models suffer from over-fitting, which can mask such problems.

I used simple posterior predictive checks based on binning the data using orthogonal boundaries and binning stimuli with zero signal strength as their own categories. This resulted in uneven bin sizes, however, even this simple method was effective in doing some preliminary comparisons between the models and observed data. Developing effective posterior predictive checks should be taken as a goal in itself, since, as already mentioned, this is crucial for model criticism and, consequently, for developing more realistic models.

Another important difference to classic GRT is that instead of identification I modelled discrimination. Since the perception of faint stimuli is more often based on the differences between stimuli, I believe models of (auditory) perception should take this into account, as it is a rather simple thing to do—although this does limit the selection of experimental tasks.

As shown by the simulations, the adaptive method, unsurprisingly, was able to achieve better performance on average, however, the difference was rather small and was engulfed in a lot of variability. This suggests that adaptive sampling of stimuli offers little in terms of precision in parameter estimates. What I would believe to be sufficient would be to estimate psychometric functions for each dimension in the absence of interactions and adapt stimuli to these functions, separately for each dimension; similar to how adaptation was implemented for one-dimensional stimuli in Kontsevich and Tyler (1999).

There's a lot of variability in how well the generating parameters of the simulated observers are recovered. The κ_μ terms were recovered fairly fast, but the posterior distributions for ρ parameters were left relatively wide. The central implication for future research is that it would be extremely difficult to for example try to infer the dependence of ρ parameter on the stimulus values (which is sometimes done in classic GRT models, see Ashby and Soto (2015); F. A. Soto et al. (2017)).

Data from the psychophysical experiments indicated rather poor fit of the initial model (Model 1), especially for observer OK. For this observer the data indicated that the main source of interaction between the dimensions was in how they influenced the standard deviations of the evidence distributions. For both observers I ended up using a model that included coefficients for both coupled means and standard deviations, however differences were less clear in the 2I-4AFC task.

The challenge that adaptive methods are facing when it comes to GRT related models is not necessarily how to minimize the entropy of the posterior distribution—as was done here—but rather how to make reliable inferences about the structural aspects of the models, and how to deal with e.g. the kind of non-identifiabilities that were discussed during the analysis of psychophysical data: for observer JP the sign of the parameter κ_σ was not identifiable which required me to constrain the parameter to positive values; on the other hand the parameter κ_μ was problematic for the observer

OK. This suggests that strong prior information might be needed to identify the models completely—at least for some data sets. At the minimum this implies that one should use posterior approximation methods that can deal with multimodal distributions.

The response input mechanism for the 2I-4AFC task proved to be problematic. I do believe that it is useful to compare parameter estimates from different tasks, since these provide simple tests of some predictions of the models, but the 2I-4AFC task, as it was implemented here, was lacking and thus severely limits its applicability. More intuitive methods for inputting responses should be developed.

In addition to this I made the assumption that the observer isn't significantly biased in the 2I-4AFC task. Due to the coding scheme (responses were coded as *hits* and *false alarms*), and not recording the actual intervals testing this assumption was not possible. The effects of bias on the parameter estimates should be explored more thoroughly in future work.

Despite these problems, the sensory (σ and β) parameters from the both tasks were fairly similar, suggesting some agreement with the theoretical predictions, however there were significant differences in the interaction parameters. This indicates strong influence from non-sensory sources to interactions, given the assumption that the interactions, if they arise from sensory sources, should not be affected by the task structure. Identifying sensory and non-sensory sources of variability has been proven problematic even in SDT, as was discussed in Section 2 *Theories of Detection*, and it would seem, based on these data, that the problem might be even worse in GRT.

The assumption that task structure should not affect sensory interactions needs some further clarification. The stimuli in the Yes/No and 2I-4AFC tasks are similar only in an abstract sense: concretely the stimuli are different. During the introspection it was brought up that this difference in the concrete appearance in the stimuli lead to somewhat differing ways of attending to the stimuli. Whether this means that the stimuli are different to the observers is of course a more complicated issue. The assumption that the level at which perception operates are the atomic stimuli themselves, that is the single sound bytes (reference and test tones), has to be made. Indeed, under that assumption it shouldn't matter what the configuration in which the stimuli are presented in is. However, if perception is more holistic and instead operates on larger scale units such as the intervals as wholes, then we should not be surprised if the results are different.

One theme that I will bring up once more is the theoretical meaning given to the latent distributions. I have been talking about sensory and non-sensory sources of interaction. However, if we take the view that the latent distributions do not model sensory effects but instead evidence, the question is not if the estimated parameters model sensory interactions: they absolutely do NOT model sensory sources, they model evidence, which is a higher level concept and more obviously modulated by for example attention.

Garner (1974) defined dimensional interactions as a set of *converging operations*, that is from an array of different tasks, some of which required the participant to attend to both dimensions and some to just one (for a review see Burns (2014)). Indeed, this research by Garner greatly influenced the seminal paper on GRT (that is Ashby and Townsend (1986)) in which one of the motivations was to unify these concepts under one theory. Since then, the task in which the participant is required to attend to both dimensions simultaneously has dominated the field—indeed, such task was also used in the psychophysical tasks in this thesis. Insights from the introspection point towards this assumption of simultaneous attention being problematic, and perhaps being one factor in why the

interaction estimates from the two tasks diverged. For greater unification a model of what *mode of attention* (that is, are they attending to the dimensions sequentially, is one of the primary, or are they attending to the wholes as a unity etc.) the participant is using in the task might be needed.

References

- Allen, E., & Oxenham, A. (2014). Symmetric interactions and interference between pitch and timbre. *Journal of Acoustical Society of America*, 135(3), 1371 - 1379.
- Amrhein, V., Korner-Nievergelt, F., & Roth, T. (2017). The earth is flat($p > 0.05$): significance thresholds and the crisis of unreplicable research. *PeerJ*, 5:e3544.
- Ashby, F. (1989). Stochastic general recognition theory. *Perception & Psychophysics*, 44, 195 - 204.
- Ashby, F. (2000). A stochastic version of general recognition theory. *Journal of Mathematical Psychology*, 44, 310 - 329.
- Ashby, F., & Maddox, W. (1994). A response time theory of separability and integrality in speeded classification. *Journal of Mathematical Psychology*, 38, 423 - 466.
- Ashby, F., & Soto, F. (2015). Multidimensional signal detection theory. In J. Busemeyer, Z. Wang, J. Townsend, & A. Eilders (Eds.), *The Oxford handbook of computational and mathematical psychology* (chap. 2).
- Ashby, F., & Townsend, J. (1986). Varieties of perceptual independence. *Psychological Review*, 93(2), 154 - 179.
- Benjamin, A. S., Diaz, M., & Wee, S. (2009). Signal detection with criterion noise: Applications to recognition memory. *Psychological Review*, 116(1), 84-115.
- Box, G., Hunter, J., & Hunter, W. (2005). *Statistics for experimenters: Design, innovation, and discovery* (2nd ed.). New Jersey: John Wiley & Sons.
- Boys, R. (1989). Algorithm AS R80: A remark on algorithm AS 76: An integral useful in calculating noncentral t and bivariate normal probabilities. *Journal of the Royal Statistical Society. Series C (Applied Statistics)*, 38(3), 580 - 582.
- Burns, D. (2014). *The many faces of Garner interference* (Unpublished doctoral dissertation). Indiana University.
- Cabrera, C., Lu, Z.-L., & Dosher, B. (2015). Separating decision and encoding noise in signal detection tasks. *Psychological Review*, 122(3), 429 - 440.
- Chopin, N. (2002). A sequential particle filter method for static models. *Biometrika*, 89(3), 539 - 551.
- Christensen, L. (1997). *Experimental methodology* (7th ed.). Allyn and Bacon: Boston.
- Cohen, D. (2003). Direct estimation of multidimensional perceptual distributions: Assessing hue and form. *Perception & Psychophysics*, 65(7), 1145 - 1160.
- Dai, H., & Micheyl, C. (2011). Psychometric functions for pure-tone frequency discrimination. *Journal of the Acoustical Society of America*, 130(1), 263 - 272.
- Decarlo, L. (1998). Signal detection theory and generalized linear models. *Psychological methods*, 3(2), 186 - 205.
- DiMattina, C. (2015). Fast adaptive estimation of multidimensional psychometric functions. *Journal of Vision*, 15(9), 1 - 20.
- Ennis, D., & Ashby, F. (2003). *Fitting decision bound models to identification or categorization data*. <http://www.psych.ucsb.edu/~ashby/cholesky.pdf>.
- Fastl, H., & Zwicker, E. (2007). *Psychoacoustics: Facts and models*. Springer.
- Freeman, E., Keele, L., Park, D., & Salzman, J. (2018). The plateau problem in the heteroscedastic probit model. arxiv:1508.03262v1.

- Garner, W. (1974). *The processing of information and structure*. Hillsdale, NJ: Erlbaum.
- Gelman, A. (2018). The failure of null hypothesis significance testing when studying incremental changes, and what to do about it. *Personality and Social Psychology Bulletin*, 44(1), 16-23.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., & Rubin, D. B. (2014). *Bayesian data analysis* (3rd ed.). Chapman and Hall, CRC.
- Gelman, A., Hill, J., & Yajima, M. (2012). Why we (usually) don't have to worry about multiple comparisons. *Journal of Research on Educational Effectiveness*, 5, 189 - 211.
- Greenland, S., Senn, S. J., Rothman, K. J., Carlin, J. B., Poole, C., Goodman, S. N., & Altman, D. G. (2016). Statistical tests, P values, confidence intervals, and power: a guide to misinterpretations. *European journal of epidemiology*, 31(4), 337-350.
- Houtsma, A. J. (1995). Chapter 8 - pitch perception. In B. C. Moore (Ed.), *Hearing* (p. 267 - 295). San Diego: Academic Press. Retrieved from <http://www.sciencedirect.com/science/article/pii/B9780125056267500108>
- Kadlec, H., & Townsend, J. (1992). Signal detection analyses of dimensional interactions. In F. Ashby (Ed.), *Multidimensional models of perception and cognition* (chap. 8).
- Kellen, D., Klauer, K., & Singmann, H. (2012, 04). On the measurement of criterion noise in signal detection theory: The case of recognition memory. *Psychological review*, 119, 457-79.
- Keller, M., & Kamary, K. (2018). Bayesian model selection and averaging via mixture model estimation. arxiv:1711.1006v2.
- Kemler Nelson, D. (1993). Processing integral dimensions: the whole view. *Journal of Experimental Psychology: Human Perception and Performance*, 19(5), 1105 - 1113.
- Kingdom, F., & Prins, N. (2010). *Psychophysics: A practical introduction*.
- King-Smith, P., & Rose, D. (1997). Principles of an adaptive method for measuring the slope of the psychometric function. *Vision Research*, 37(12), 1595 - 1604.
- Kline, R. (2004). *Beyond significance testing: reforming data analysis methods in behavioral research*. Washington: American Psychological Association.
- Kontsevich, L., & Tyler, C. (1999). Bayesian adaptive estimation of psychometric slope and threshold. *Vision Research*, 39(16).
- Kruschke, J. (2015). *Doing Bayesian data analysis: A tutorial with R, Jags, and Stan* (2nd ed.).
- Kujala, J. (2011). Bayesian adaptive estimation: A theoretical review. In (chap. 6).
- Kujala, J., & Lukka, T. (2006). Bayesian adaptive estimation: The next dimension. *Journal of Mathematical Psychology*, 50, 369 - 389.
- Lee, M. (2018). Bayesian methods in cognitive modeling. In J. Wixted & E.-J. Wagenmakers (Eds.), *Steven's handbook of experimental psychology* (4th ed.).
- Lesmes, L., Jeon, S.-T., Lu, Z. L., & Dosher, B. (2006). Bayesian adaptive estimation of threshold versus contrast external noise functions: The quick *tvc* method. *Vision research*, 46, 3160 - 3176.
- Lesmes, L., Lu, Z.-L., Baek, J., Tran, N., Dosher, B., & Albright, T. (2015). Developing Bayesian adaptive methods for estimating sensitivity thresholds(d') in yes-no and forced-choice tasks. *Frontiers in Psychology*, 6(1070).
- Newell, A. (1973). You can't play 20 questions with nature and win: projective comments on the papers of this symposium. In W. G. Chase (Ed.), *Visual information processing*. Academic Press, New York.

- Oron, A. P. (2007). *Up-and-down and the percentile-finding problem* (Unpublished doctoral dissertation). University of Washington.
- Pan, K. (2017). *An analytical expression for bivariate normal distribution*. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=2924071. (Accessed: 26.11.2019)
- Platt, J., & Racine, R. (1985). Effect of frequency, timbre, experience and feedback on musical tuning skills. *Perception & Psychophysics*, 38(6), 543 - 553.
- Prins, N. (2012). The psychometric function: The lapse rate revisited. *Journal of vision*, 12(6), 1-16.
- Prior choice recommendations*. (n.d.). <https://github.com/stan-dev/stan/wiki/Prior-Choice-Recommendations>. (Accessed: 9.7.2020)
- R Core Team. (2019). R: A language and environment for statistical computing [Computer software manual]. Vienna, Austria. Retrieved from <https://www.R-project.org/>
- Shen, J., Sivakumar, R., & Richards, M. (2014). Rapid estimation of high-parameter auditory filters. *Journal of the Acoustical Society of America*, 136(4), 1857 - 1868.
- Shen, Y., & Richards, V. (2013). Bayesian adaptive estimation of the auditory filter. *Journal of Acoustical Society of America*, 10(2).
- Silbert, N. (2010). *Integration of phonological information in obstruent consonant identification* (Unpublished doctoral dissertation). Indiana University.
- Silbert, N., & Thomas, R. (2013). Decisional separability, model identification, and statistical inference in the general recognition theory framework. *Psychonomic Bulletin & Review*, 20, 1 - 20.
- Silbert, N., & Thomas, R. (2016). <https://arxiv.org/pdf/1606.05598.pdf>. (Accessed: 7.6.2018)
- Silbert, N., Townsend, J., & Lentz, J. (2009). Independence and separability in the perception of complex nonspeech sounds. *Attention, Perception, & Psychophysics*, 71(8), 1900 - 1915.
- Skrondahl, A., & Rabe-Hesketh, S. (2004). *Generalized latent variable modeling: Multilevel, longitudinal, and structural equation models*. Bosa Roca, United States: Taylor & Francis Inc.
- Soto, F., Vucoovich, L., Musgrave, R., & Ashby, F. (2015). General recognition theory with individual differences: a new method for examining perceptual and decisional interactions with an application to face perception. *Psychonomic Bulletin & Review*, 22(1), 88 - 111.
- Soto, F. A., Zheng, E., Fonseca, J., & Ashby, F. (2017). Testing separability and independence of perceptual dimensions with general recognition theory: A tutorial and new R package (*grtools*). *Frontiers in Psychology*.
- Speekenbrink, M. (2016). A tutorial on particle filters. *Journal of Mathematical Psychology*, 73, 140 - 152.
- Stan Development Team. (2017). Stan modeling language: User's guide and reference manual [Computer software manual]. (This is a manual for an earlier version of Stan)
- Stan Development Team. (2019a). *RStan: the R interface to Stan*. Retrieved from <http://mc-stan.org/> (R package version 2.19.2)
- Stan Development Team. (2019b). Stan reference manual [Computer software manual]. Retrieved from https://mc-stan.org/docs/2_23/reference-manual/index.html (Version 2.23)
- Steiger, J., & Fouladi, R. (1997). Noncentrality interval estimation and the evaluation of statistical models. In L. Harlow, S. Mulaik, & J. Steiger (Eds.), *What if there were no significance tests?* (chap. 9). Mahwah, NJ: Erlbaum.

- Stigler, S. (2003). *The history of statistics: The measurement of uncertainty before 1900*. Cambridge: Belknap Press of Harvard University Press.
- Townsend, J., Hu, G., & Kadlec, H. (1988). Feature sensitivity, bias, and interdependencies as a function of energy and payoffs. *Perception & Psychophysics*, 43, 575 - 591.
- Treuwein, B. (1995). Adaptive psychophysical procedures. *Vision Research*, 35(17), 250 - 2522.
- Treuwein, B., & Strasburger, H. (1999). Fitting the psychometric function. *Perception & psychophysics*, 61(1), 87 - 106.
- Verde, M., MacMillan, N., & Rotello, C. (2006). Measures of sensitivity based on a single hit and false alarm rate: The accuracy, precision and robustness of d' , A_z , and A' . *Perception & Psychophysics*, 68(4), 643 - 654.
- Watson, A. (1983). QUEST: A bayesian adaptive psychometric method. *Perception & Psychophysics*, 32(2), 113 - 120.
- Watson, A. (2017). QUEST+: A general multidimensional Bayesian adaptive psychometric method. *Journal of Vision*, 17(3), 1 - 27.
- Wichmann, F., & Hill, N. (2001). The psychometric function I: Fitting, sampling, and goodness-of-fit. *Perception & psychophysics*, 63, 1293 - 1313.
- Wickelgren, W. (1968). Unidimensional strength theory and component analysis of noise in absolute and comparative judgments. *Journal of Mathematical Psychology*, 5, 102 - 122.
- Wickens, T. (1992). Maximum-likelihood estimation of a multivariate gaussian rating model with excluded data. *Journal of Mathematical Psychology*, 36, 213 - 234.
- Wickens, T. (2002). *Elementary signal detection theory* (3rd ed.). New York: Oxford University Press.
- Zeigenfuse, M., & Lee, M. (2010). A general latent assignment approach for modeling psychological contaminants. *Journal of Mathematical Psychology*, 54, 352 - 362.

A Algorithm for calculating ϕ_2

The algorithm is based on Boys (1989) and Pan (2017).

```
functions {
    real bivariate_cdf(real h, real k, real rho) {
        real p;
        real a1;
        real a2;
        real denominator;
        real deltacoeff;

        if(h != 0 || k != 0){
            // The point of this is to get rid of sign in
            // the zero, if there is one. I.e. change -0 to 0.

            real h_ = h == 0 ? 0 : h;
            real k_ = k == 0 ? 0 : k;

            // These ratios (between k and h) are the wrong
            // way round in the paper
            denominator = sqrt(1 - rho * rho);
            a1 = (k_ / h_ - rho) / denominator;
            a2 = (h_ / k_ - rho) / denominator;

            // The a-parameter is limited between (-1e6, 1e6),
            // so that the gradients don't "blow up":

            if(a1 > 1e6){
                a1 = 1e6;
            }
            else if(a1 < -1e6){
                a1 = -1e6;
            }

            if(a2 > 1e6){
                a2 = 1e6;
            }
            else if(a2 < -1e6){
                a1 = -1e6;
            }
        }
    }
}
```

```

if ((h_ * k_) > 0 || (h_ * k_) == 0 && (h_ + k_) >= 0){
    deltacoeff = 0.0;
} else{
    deltacoeff = 1.0;
}

p = 0.5 * (Phi(h_) + Phi(k_) - deltacoeff) - owens_t(h_, a1) - owens_t(k_, a2);
} else{
    // Form the "other" paper, eq. 10:
    p = 0.1591549 * asin(rho) + 0.25;
}

if (p < 0){
    // Since a is limited, the approximated integral will
    // sometimes be slightly negative. But not much! Don't worry!
    // print("P was lower than 0: ", p);
    p = 0;
}
return p;
}
}

```

B Code for models for the Yes/No task

B.1 Model 1

```
#include /bivariate_cdf.stan

data{
    int NTrials;
    int R[NTrials, 2];
    real S[NTrials, 2];
}

parameters{
    vector [2] sigma;
    vector [2] crit;
    vector [2] beta;
    vector [2] kappa;

    real rho;
}

model{
    real zVals [2];
    real p[NTrials];

    sigma ~ normal(log(2.5), 0.75);
    crit ~ normal(1.5, 0.30);
    beta ~ normal(log(1), 0.30);

    kappa ~ normal(0, 0.30);

    rho ~ normal(0, 0.70);

    // Trial loop:
    for(t in 1:NTrials){

        zVals [1] = -crit [1] + pow(S[t,1] / exp(sigma [1]),
            exp(beta [1])) + kappa [1] * S[t,2];
        zVals [2] = -crit [2] + pow(S[t,2] / exp(sigma [2]),
            exp(beta [2])) + kappa [2] * S[t,1];

        p[t] = 0.02 * 0.25 + 0.98 * bivariate_cdf(R[t,1] * zVals [1],
            R[t,2] * zVals [2], R[t,1] * R[t,2] * tanh(rho));
    }
}
```

```

    }
    target += sum(log(p));
}

```

B.2 Model 2

```

#include /bivariate_cdf.stan

data{
    int NTrials;
    int R[NTrials, 2];
    real S[NTrials, 2];
}

parameters{
    vector[2] sigma;
    vector[2] crit;
    vector[2] beta;
    vector[2] kappa;
    real rho;

    real<lower = 0, upper = 1> lambda;
}

model{
    real zVals[2];
    real p[NTrials];

    sigma ~ normal(log(2.5), 0.75);
    crit ~ normal(2.0, 1.00);
    beta ~ normal(log(1), 0.30);

    kappa ~ normal(0, 0.30);

    rho ~ normal(0, 0.70);

    lambda ~ beta(2, 12);

    // Trial loop:
    for(t in 1:NTrials){

        zVals[1] = -crit[1] + pow(S[t,1] / exp(sigma[1] +

```

```

    kappa[1] * S[t,2]) , exp(beta[1]));
zVals[2] = -crit[2] + pow(S[t,2] / exp(sigma[2] +
    kappa[2] * S[t,1]) , exp(beta[2]));

p[t] = lambda * 0.25 + (1 - lambda) *
    bivariate_cdf(R[t,1] * zVals[1] , R[t,2] * zVals[2] ,
    R[t,1] * R[t,2] * tanh(rho));

}
target += sum(log(p));
}

```

B.3 Model 3

```
#include /bivariate_cdf.stan

data{
    int NTrials;
    int R[NTrials, 2];
    real S[NTrials, 2];
}

parameters{
    vector [2] sigma;
    vector [2] crit;
    vector [2] beta;
    vector [2] kappa;
    real rho;

    real<lower = 0, upper = 1> lambda;
}

model{
    real zVals[2];
    real p[NTrials];
    real sigmas[2];

    sigma ~ normal(log(2.5), 0.75);
    crit ~ normal(2.0, 1.00);
    beta ~ normal(log(1), 0.30);

    kappa ~ normal(0, 1.00);
}
```

```

rho ~ normal(0, 0.70);

lambda ~ beta(2, 12);

// Trial loop:
for(t in 1:NTrials){

sigmas[1] = exp(sigma[1] + kappa[1] * S[t,2]);
sigmas[2] = exp(sigma[2] + kappa[2] * S[t,1]);

zVals[1] = -crit[1]/sigmas[1] + pow(S[t,1] /
sigmas[1], exp(beta[1]));
zVals[2] = -crit[2]/sigmas[2] + pow(S[t,2] /
sigmas[2], exp(beta[2]));

p[t] = lambda * 0.25 + (1 - lambda) *
bivariate_cdf(R[t,1] * zVals[1], R[t,2] * zVals[2],
R[t,1] * R[t,2] * tanh(rho));

}

target += sum(log(p));
}

```

B.4 Both

```

#include /bivariate_cdf.stan

data{
int NTrials;
real R[NTrials,2];
real S[NTrials,2];
}

parameters{
real sigma[2];
real crit[2];
real beta[2];
real kappa_mu[2];
real kappa_sig[2];

real<lower = -0.8, upper = 0.8> rho;
real<lower = 0, upper = 1> lambda;
}

```

```

model{
  real p[NTrials];
  real zVals[2];

  sigma ~ normal(log(2.5), 0.75);

  for(i in 1:2){
    crit[i]/exp(sigma[i]) ~ normal(2.0, 1.00);
  }

  beta ~ normal(log(1), 0.30);

  kappa_mu ~ normal(0, 0.30);
  kappa_sig ~ normal(0, 0.30);

  rho ~ uniform(-0.8, 0.8);
  lambda ~ beta(2, 48);

  // Trial loop:
  for(t in 1:NTrials){

    zVals[1] = -crit[1] / exp(sigma[1] + kappa_sig[1] * S[t,2]) +
      (pow(S[t,1] / exp(sigma[1] + kappa_sig[1] * S[t,2]),
        exp(beta[1]))) + kappa_mu[1] * S[t,2];
    zVals[2] = -crit[2] / exp(sigma[2] + kappa_sig[2] * S[t,1]) +
      (pow(S[t,2] / exp(sigma[2] + kappa_sig[2] * S[t,1]),
        exp(beta[2]))) + kappa_mu[2] * S[t,1];

    p[t] = lambda * 0.25 + (1 - lambda) *
      bivariate_cdf(R[t,1] * zVals[1], R[t,2] * zVals[2],
      (R[t,1] * R[t,2]) * rho);
  }

  target += sum(log(p));
}

```

B.5 Both (Truncated)

```
#include /bivariate_cdf.stan
```

```
data{
  int NTrials;
  real R[NTrials,2];
  real S[NTrials,2];
```

```

}

parameters{
  real sigma[2];
  real crit[2];
  real beta[2];
  real kappa_mu[2];
  real<lower = 0> kappa_sig[2];

  real<lower = -0.8, upper = 0.8> rho;
  real<lower = 0, upper = 1> lambda;
}

model{
  real p[NTrials];
  real zVals[2];

  sigma ~ normal(log(2.5), 0.75);

  for(i in 1:2){
    crit[i]/exp(sigma[i]) ~ normal(2.0, 1.00);
  }

  beta ~ normal(log(1), 0.30);

  kappa_mu ~ normal(0, 0.30);
  kappa_sig ~ normal(0, 0.30);

  rho ~ uniform(-0.8, 0.8);
  lambda ~ beta(2, 48);

  // Trial loop:
  for(t in 1:NTrials){

    zVals[1] = -crit[1] / exp(sigma[1] + kappa_sig[1] * S[t,2]) +
      (pow(S[t,1] / exp(sigma[1] + kappa_sig[1] * S[t,2]), exp(beta[1]))) + kappa_mu[1] * S[t,2];
    zVals[2] = -crit[2] / exp(sigma[2] + kappa_sig[2] * S[t,1]) +
      (pow(S[t,2] / exp(sigma[2] + kappa_sig[2] * S[t,1]), exp(beta[2]))) + kappa_mu[2] * S[t,1];

    p[t] = lambda * 0.25 + (1 - lambda) *
      bivariate_cdf(R[t,1] * zVals[1], R[t,2] * zVals[2],

```

```
(R[t,1] * R[t,2]) * rho);  
}  
  
target += sum(log(p));  
}
```

C Code for models for the 2I-4AFC task

C.1 Model 1

```
#include bivariate_cdf.stan

data{
    int  NTrials;
    int  R[NTrials , 2];
    real S[NTrials , 2];
}

parameters{
    vector [2] sigma;
    vector [2] beta;
    vector [2] kappa;
    real rho;
}

model{
    real zVals [2];
    real p[NTrials];

    sigma ~ normal(log(2.5) , 0.75);
    beta ~ normal(log(1.0) , 0.30);

    kappa ~ normal(0 , 0.30);

    rho ~ normal(0 , 0.70);

    // Trial loop:
    for(t in 1:NTrials){

        zVals [1] = pow(S[t,1] / exp(sigma[1]) , exp(beta[1])) /
            sqrt(2) + kappa[1] * S[t,2];
        zVals [2] = pow(S[t,2] / exp(sigma[2]) , exp(beta[2])) /
            sqrt(2) + kappa[2] * S[t,1];

        p[t] = 0.02 * 0.25 + 0.98 *
            bivariate_cdf(R[t,1] * zVals[1] , R[t,2] * zVals[2] ,
            R[t,1] * R[t,2] * tanh(rho));
    }
}
```

```

    }
    target += sum(log(p));
}

```

C.2 Model 2

```
#include bivariate_cdf.stan
```

```

data{
    int NTrials;
    int R[NTrials, 2];
    real S[NTrials, 2];
}

parameters{
    vector [2] sigma;
    vector [2] beta;
    vector [2] kappa;

    real rho;
    real<lower = 0, upper = 1> lambda;
}

model{
    real zVals[2];
    real p[NTrials];

    sigma ~ normal(log(2.5), 0.75);
    beta ~ normal(log(1.0), 0.30);

    kappa ~ normal(0, 0.30);

    rho ~ normal(0, 0.70);
    lambda ~ beta(2, 12);

    // Trial loop:
    for(t in 1:NTrials){

        zVals[1] = pow(S[t,1] / exp(sigma[1]), exp(beta[1])) /
            sqrt(2 + kappa[1] * S[t,2];
        zVals[2] = pow(S[t,2] / exp(sigma[2]), exp(beta[2])) /
            sqrt(2 + kappa[2] * S[t,1];
    }
}
```

```

    p[t] = lambda * 0.25 + (1.0 - lambda) *
        bivariate_cdf(R[t,1] * zVals[1], R[t,2] * zVals[2],
        R[t,1] * R[t,2] * tanh(rho));

}
target += sum(log(p));
}

```

C.3 Model 3

```

#include bivariate_cdf.stan

data{
    int NTrials;
    int R[NTrials, 2];
    real S[NTrials, 2];
}

parameters{
    vector [2] sigma;
    vector [2] beta;
    vector [2] kappa;

    real rho;
    real<lower = 0, upper = 1> lambda;
}

model{
    real zVals [2];
    real p[NTrials];

    sigma ~ normal(log(2.5), 0.75);
    beta ~ normal(log(1.0), 0.30);

    kappa ~ normal(0, 0.30);

    rho ~ normal(0, 0.70);
    lambda ~ beta(2, 12);

    // Trial loop:
    for(t in 1:NTrials){

        zVals[1] = pow(S[t,1] / exp(sigma[1] + kappa[1] * S[t,2]),

```

```

    exp(beta[1])) / sqrt(2);
zVals[2] = pow(S[t,2] / exp(sigma[2] + kappa[2] * S[t,1]), 
exp(beta[2])) / sqrt(2);

p[t] = lambda * 0.25 + (1.0 - lambda) *
bivariate_cdf(R[t,1] * zVals[1], R[t,2] * zVals[2],
R[t,1] * R[t,2] * tanh(rho));

}
target += sum(log(p));
}

```

C.4 Both

```
#include /bivariate_cdf.stan

data{
  int NTrials;
  real R[NTrials,2];
  real S[NTrials,2];
}

parameters{
  real sigma[2];
  real beta[2];
  real kappa_mu[2];
  real kappa_sig[2];

  real<lower = -0.8, upper = 0.8> rho;
  real<lower = 0, upper = 1> lambda;
}

model{
  real p[NTrials];
  real zVals[2];

  sigma ~ normal(log(2.5), 0.75);

  beta ~ normal(log(1), 0.30);

  kappa_mu ~ normal(0, 0.30);
  kappa_sig ~ normal(0, 0.30);
}
```

```

rho ~ uniform(-0.8, 0.8);
lambda ~ beta(2, 48);

// Trial loop:
for(t in 1:NTrials){

zVals[1] = (pow(S[t,1] / exp(sigma[1] + kappa_sig[1] * S[t,2]),
exp(beta[1])) + kappa_mu[1] * S[t,2]) / 1.414214;
zVals[2] = (pow(S[t,2] / exp(sigma[2] + kappa_sig[2] * S[t,1]),
exp(beta[2])) + kappa_mu[2] * S[t,1]) / 1.414214;

p[t] = lambda * 0.25 + (1 - lambda) *
bivariate_cdf(R[t,1] * zVals[1], R[t,2] * zVals[2],
(R[t,1] * R[t,2]) * rho);
}

target += sum(log(p));
}

```